

Deep Learning with Pytorch

v15



Trainer: Dr. Alfred Ang



Website: www.tertiarycourses.com.sg
Email: enquiry@tertiaryinfotech.com

About the Trainer

Dr. Alfred Ang is the founder of Tertiary Courses. He is a serial entrepreneur. He founded OSWeb2Design Singapore Pte Ltd in 1997 offering web development, e-commerce store development, graphics design, ebook publishing, mobile apps development, and digital marketing services. He established the first online gardening store in Singapore, Eco City Hydroponics Pte Ltd in 2000, offering a wide range of gardening products such as seeds, plant nutrients, hydroponics kits etc. Eco City Hydroponics has become the most popular and successful gardening store in Singapore. He founded Tertiary Infotech Pte Ltd in 2012 and transformed the business to a training platform, Tertiary Courses in 2014. Tertiary Courses offers a wide range of SkillsFuture courses for PMETs to upgrade their skills and knowledge. He also established Tertiary Courses Malaysia in 2016. He also founded Tertiary Robotics in 2015 offering Arduino, Raspberry Pi, Microbit and Robotics product

Dr. Alfred Ang earned his Ph.D. from National University of Singapore in 2000, majoring in Electrical and Electronics Engineering. He also completed an online MBA course with U21 Global based in Australia. He obtained his B.Sc (Hons) from National University of Singapore in 1992, majoring in Physics. He topped his Physics cohort for 3 consecutive years and funded his degree study with Book prizes, Study awards, bursaries and tuition. He has worked in Defence, Electronics and Semiconductor Industries. His current interests include Machine Learning, Deep Learning, Artificial Intelligence, Internet of Things, Robotics and Programming.

Dr. Alfred Ang is ACTA certified trainer and DACE certified developer. He is an associate adult educator with IAL, IBM Certified AI Practitioner and IBM Certified Design Thinking Practitioner. He was Distinguished Toastmasters (DTM) and Senior Member of IEEE. He has published more than 20 peer reviewed papers and co-inventors for more than 20 inventions.



Let's Know Each Other...

Say a bit about yourself

- Name
- What Industry you are from?
- Do you have any prior knowledge in Deep Learning or Pytorch?
- Why do you want to learn Deep Learning and Pytorch?
- What do you expect to learn from this course?

Ground Rules

- Set your mobile phone to silent mode
- Actively participate in the class. No question is stupid.
- Respect each other view
- Exit the class silently if you need to step out for phone call, toilet break

Ground Rules for Virtual Training

- Upon entering, mute your mic and turn on the video. Use a headset if you can
- Use the 'raise hand' function to indicate when you want to speak
- Participant actively. Feel free to ask questions on the chat whenever.
- Facilitators can use breakout rooms for private sessions.



Guidelines for Facilitators

1. Once all the participants are in and introduce themselves
2. Goto gallery mode, take a snapshot of the class photo - makes sure capture the date and time
3. Start the video recording (only for WSQ courses)
4. Continue the class
5. Before the class end on that day, take another snapshot of the class photo - makes sure capture the date and time
6. For NRIC verification, facilitator to create breakout room for individual participant to check (only for WSQ courses)
7. Before the assessment start, take another snapshot of the class photo - makes sure capture the date and time (only for WSQ courses)
8. For Oral Questioning assessment, facilitator to create breakout room for individual participant to OQ (only for WSQ courses)
9. End the video recording and upload to cloud (only for WSQ courses)
10. Assessor to send all the assessment records, assessment plan and photo and video to the staff (only for WSQ courses).

Prerequisite

- Basic Python
- Some knowledge of Machine Learning

Agenda

Topic 1 Overview of Deep Learning and Pytorch

- Overview of Deep Learning
- Introduction to Pytorch
- Install and Run Pytorch
- Basic Pytorch Tensor Operations
- Computation Graphs
- Compute Gradients with Autograd

Topic 2 Neural Network for Regression

- Introduction to Neural Network (NN)
- Activation Function
- Loss Function and Optimizer
- Machine Learning Methodology
- Build a NN Predictive Regression Model
- Load and Save Model

Topic 3 Neural Network for Classification

- Softmax
- Cross Entropy Loss Function
- Build a NN Classification Model

Agenda

Topic 4 Convolutional Neural Network (CNN)

- Overview of CNN
- Convolution, Max Pooling and Padding
- Build a CNN Model for Image Classification
- Overfitting Issue with Small Dataset
- Techniques to overcome Overfitting Issue

Topic 5 Transfer Learning

- Introduction to Transfer Learning
- Pre-trained Models
- Feature Extraction & Fine Tuning for Small Dataset

Topic 6 Recurrent Neural Network (RNN)

- Overview of RNN
- Long Term Dependencies
- LSTM and GRU
- Apply LSTM to Time Series Forecasting

Download Course Material

- You can download the course material from Google classroom <https://classroom.google.com>
- Enter the class code below to join the class on the top right.
- Goto Classwork>> Course material
- If you cannot access the Google Classroom, please inform trainer or staff.

nn7x4ou

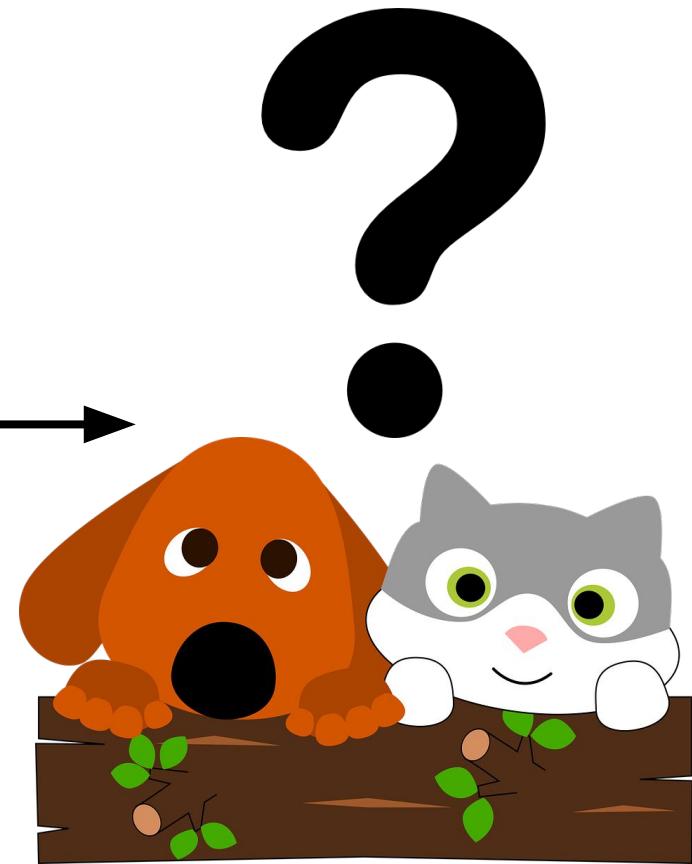
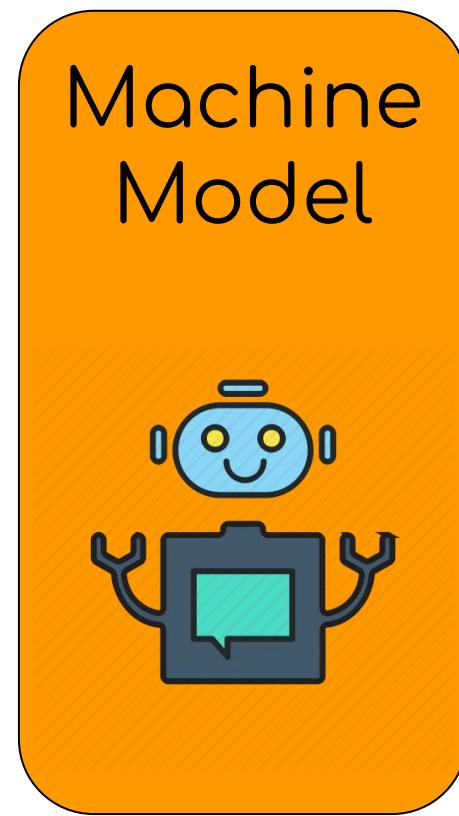
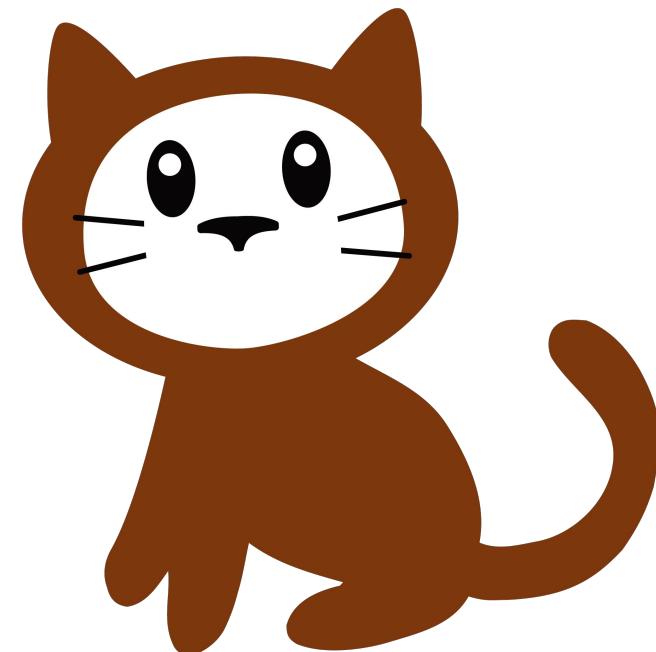
Topic 1

Overview of Deep Learning and Pytorch

**So, do you
see a crow
...or a
cat?**



After you trained the machine to recognize dog or cat, then machine is able to tell your the answer when you give it an image



Why Machine Learning?

Google AlphaGO (A Deep NN) beat World GO Champion Ke Jia on May 2017



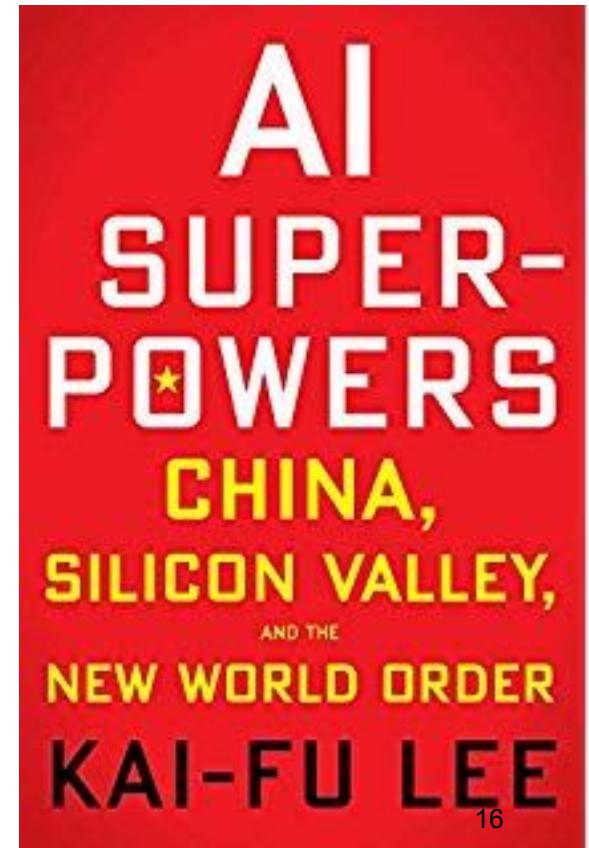


**“AI is the new electricity.
It has the potential to
transform every industry
and to create huge
economic value”**

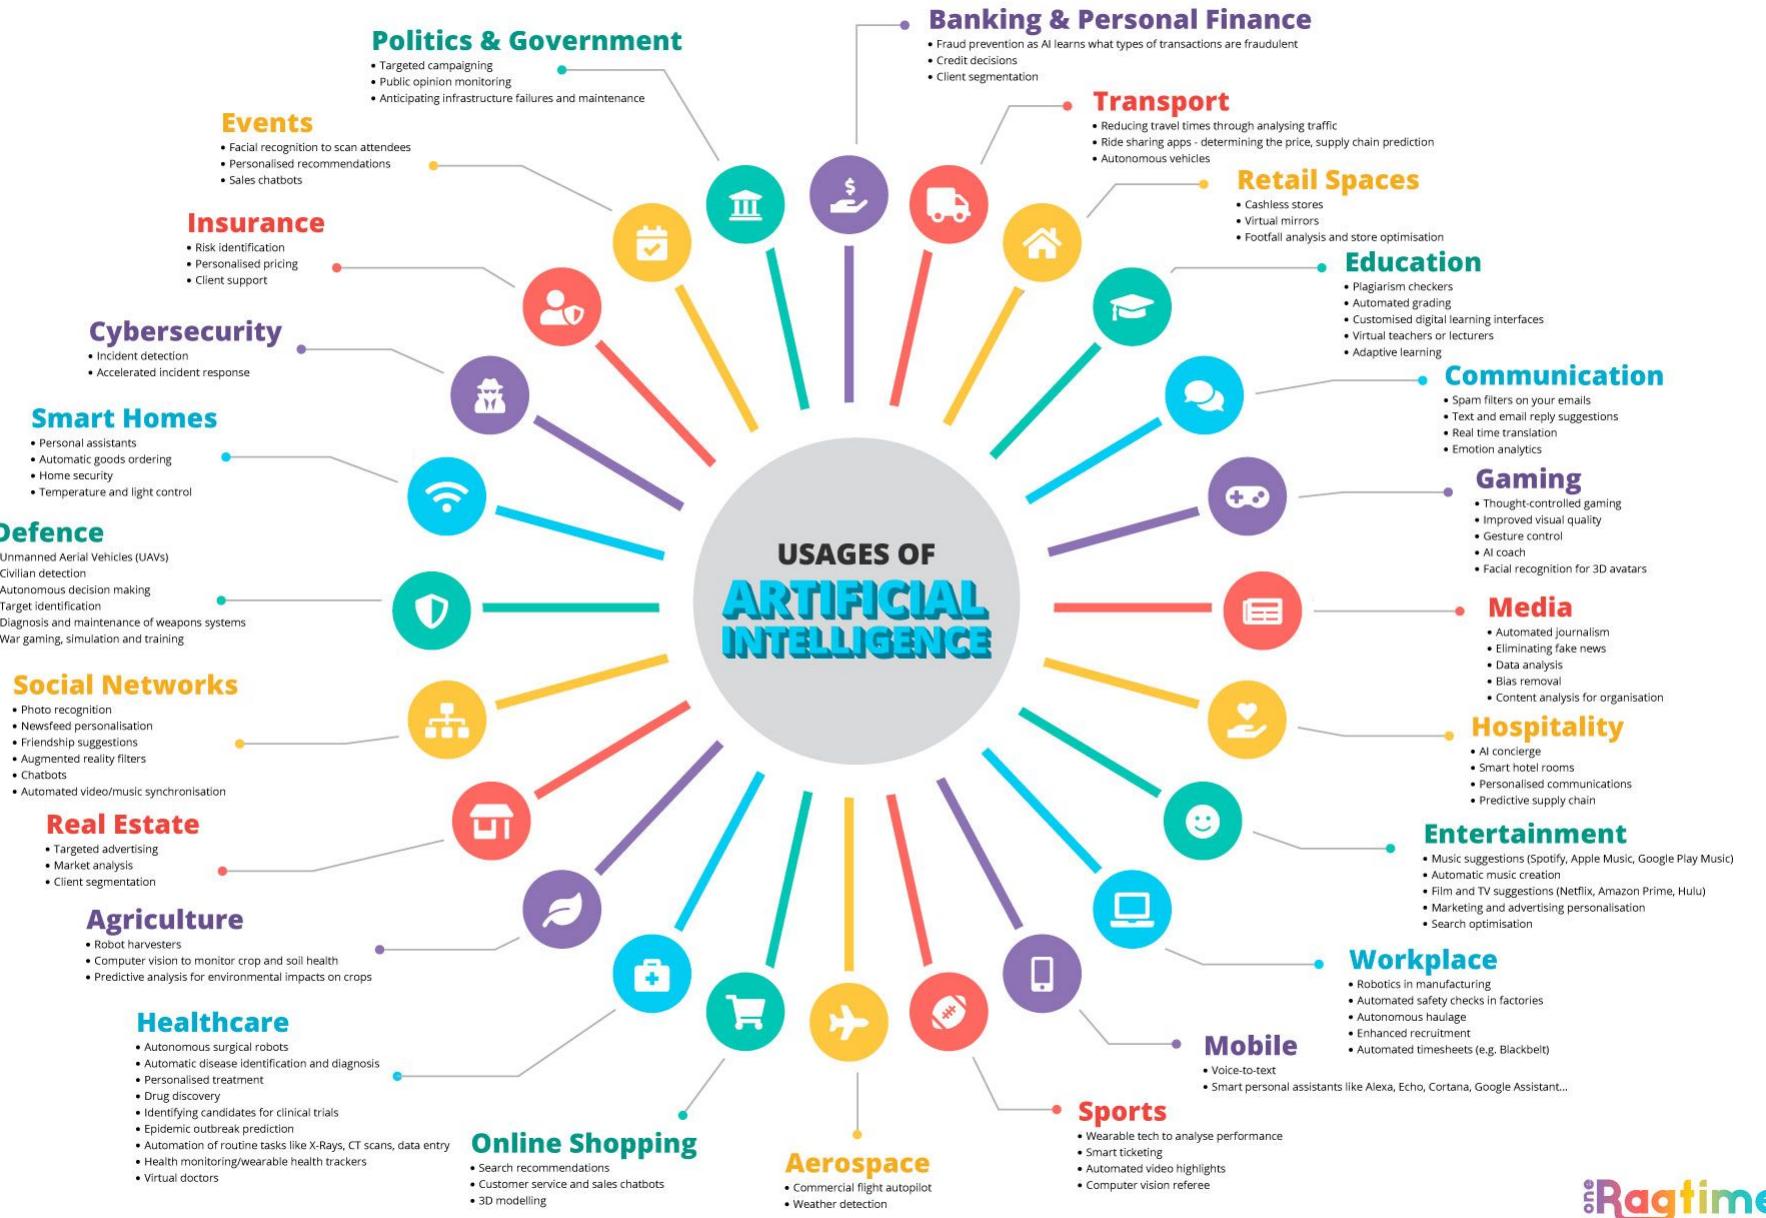
--Andrew Ng

Will AI take over Human Jobs?

"Artificial intelligence will replace half of all jobs in the next decade" - Lee Kai Fu

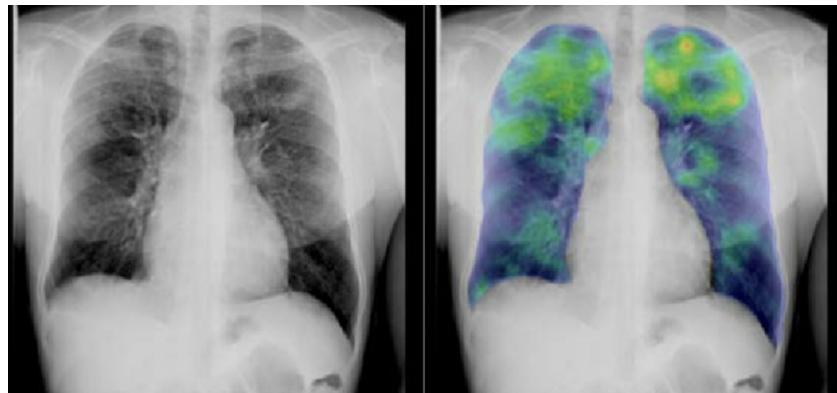


Machine Learning Applications

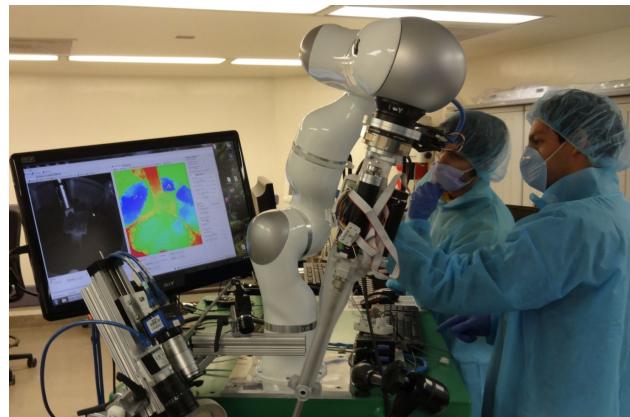


ML Applications on Health Care

Identifying tuberculosis



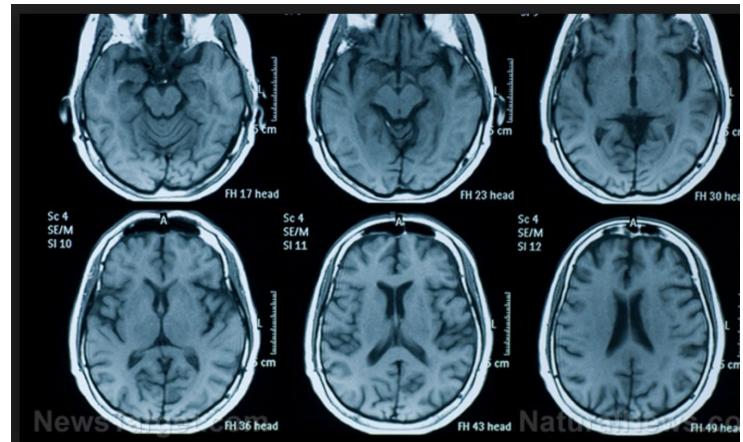
Robotics-assisted surgery



Detecting brain bleeds



Detecting Alzheimer's disease



ML Applications on Retail/Logistics

Amazon Go



Frontdesk Robots



Alibaba Unmanned store 阿里巴巴无人超市



Chatbot for retail/services

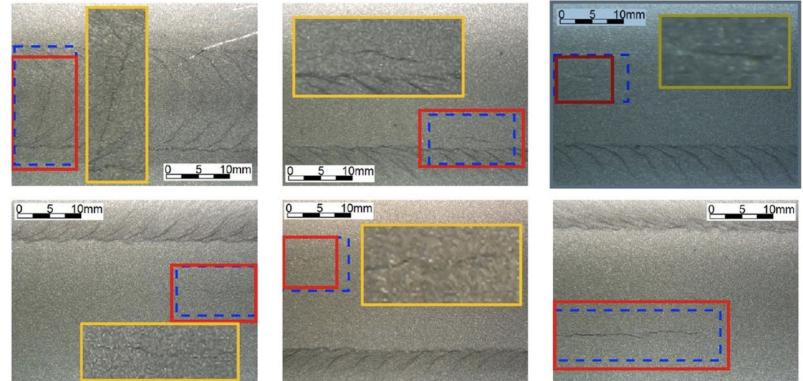


ML Applications on Security/Mfg

Face recognition
identifying criminals



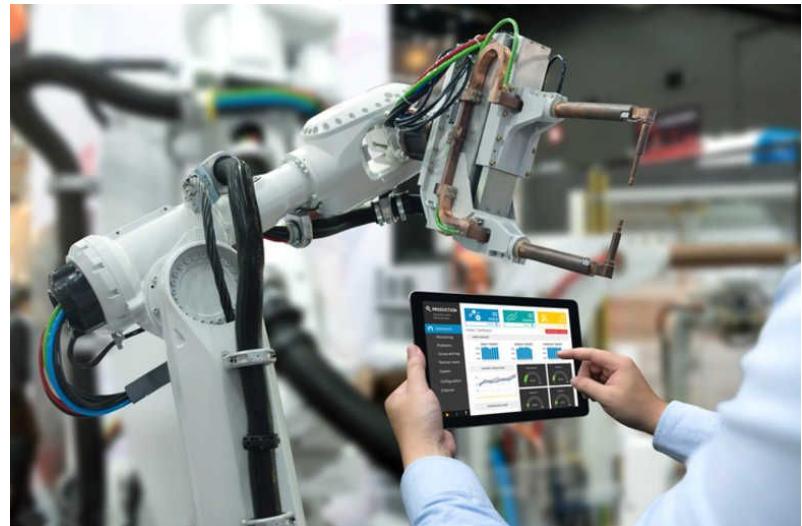
Defect detection



Parking Spotter



Robots in Mfg

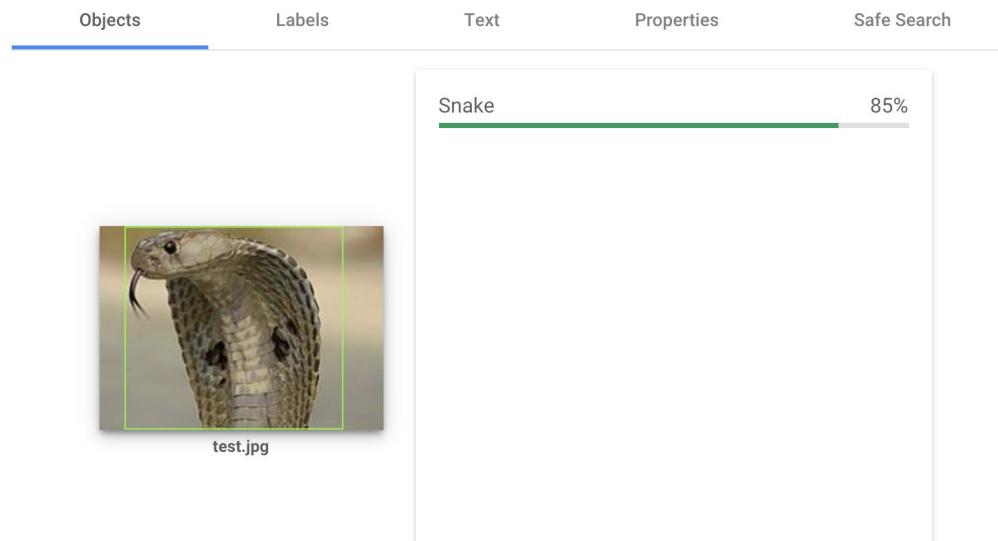


Activity : Vision

- Upload some images from internet
- Use the following Vision tools to classify the images
 - Google Vision Demo:
<https://cloud.google.com/vision>
 - IBM Vision Demo:
<https://watson-visual-recognition-duo-dev.ng.bluemix.net/>

Objects Labels Text Properties Safe Search

Snake 85%



AI vs ML vs DL

ARTIFICIAL INTELLIGENCE

Early artificial intelligence stirs excitement.



MACHINE LEARNING

Machine learning begins to flourish.



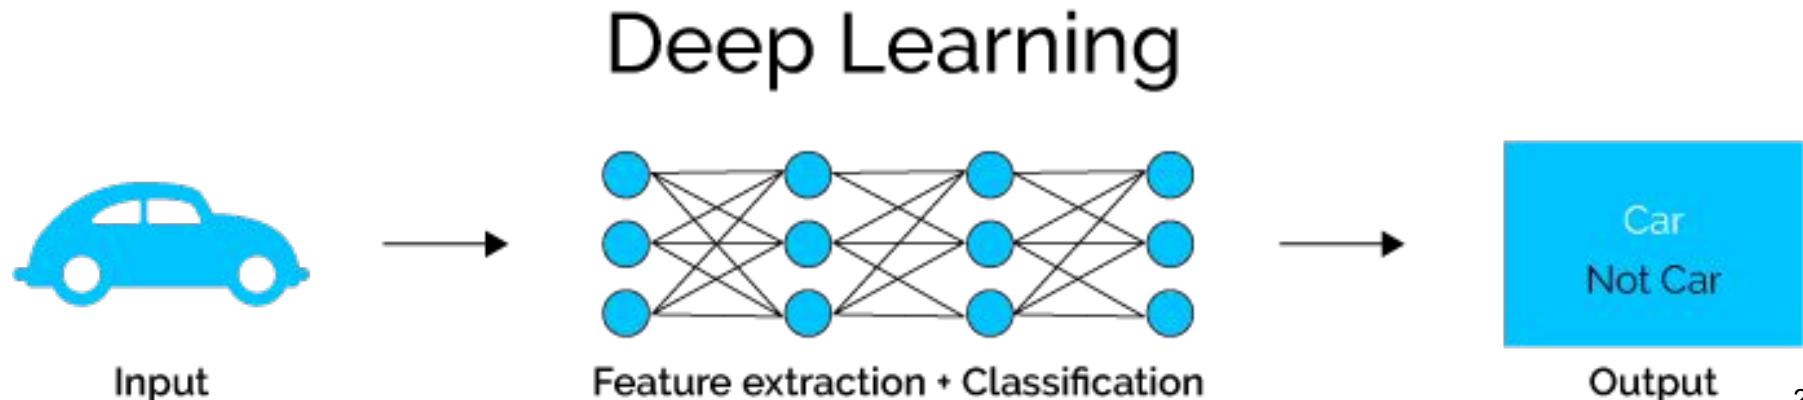
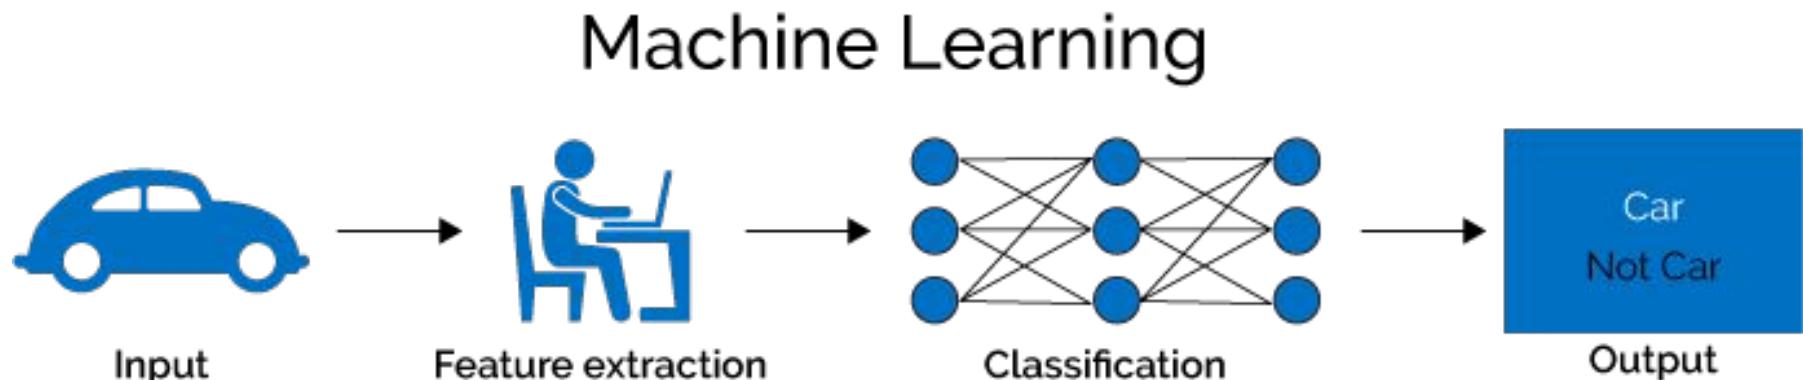
DEEP LEARNING

Deep learning breakthroughs drive AI boom.



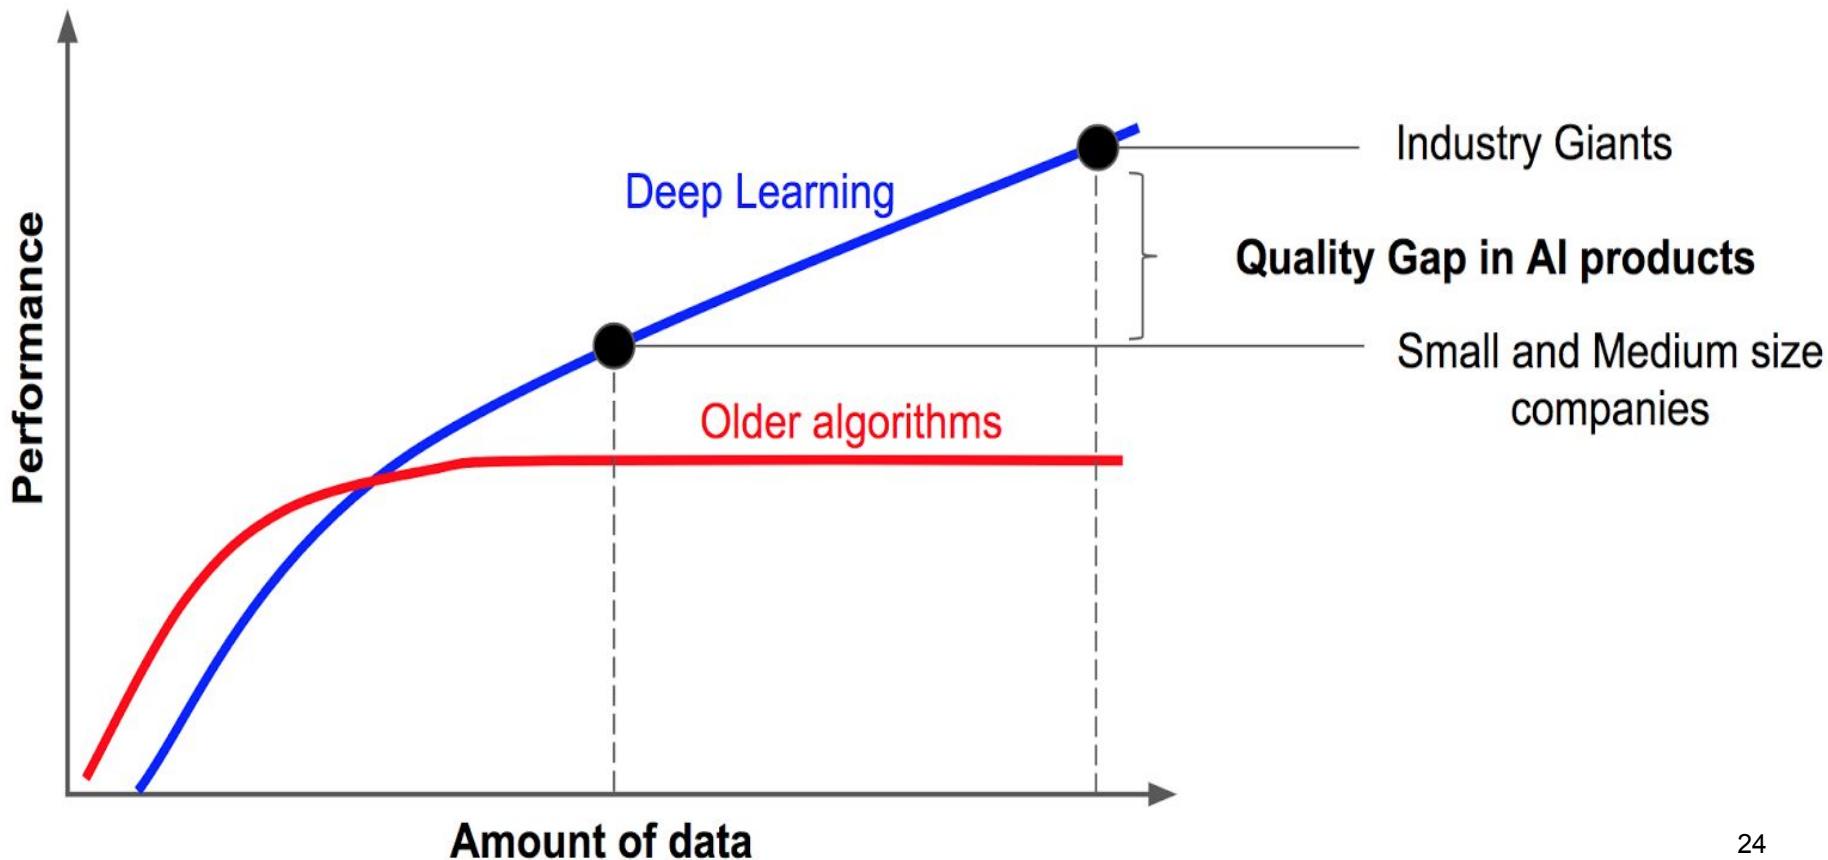
Machine Learning vs Deep Learning

Deep Learning are neural network based and consists of millions of parameters to train.



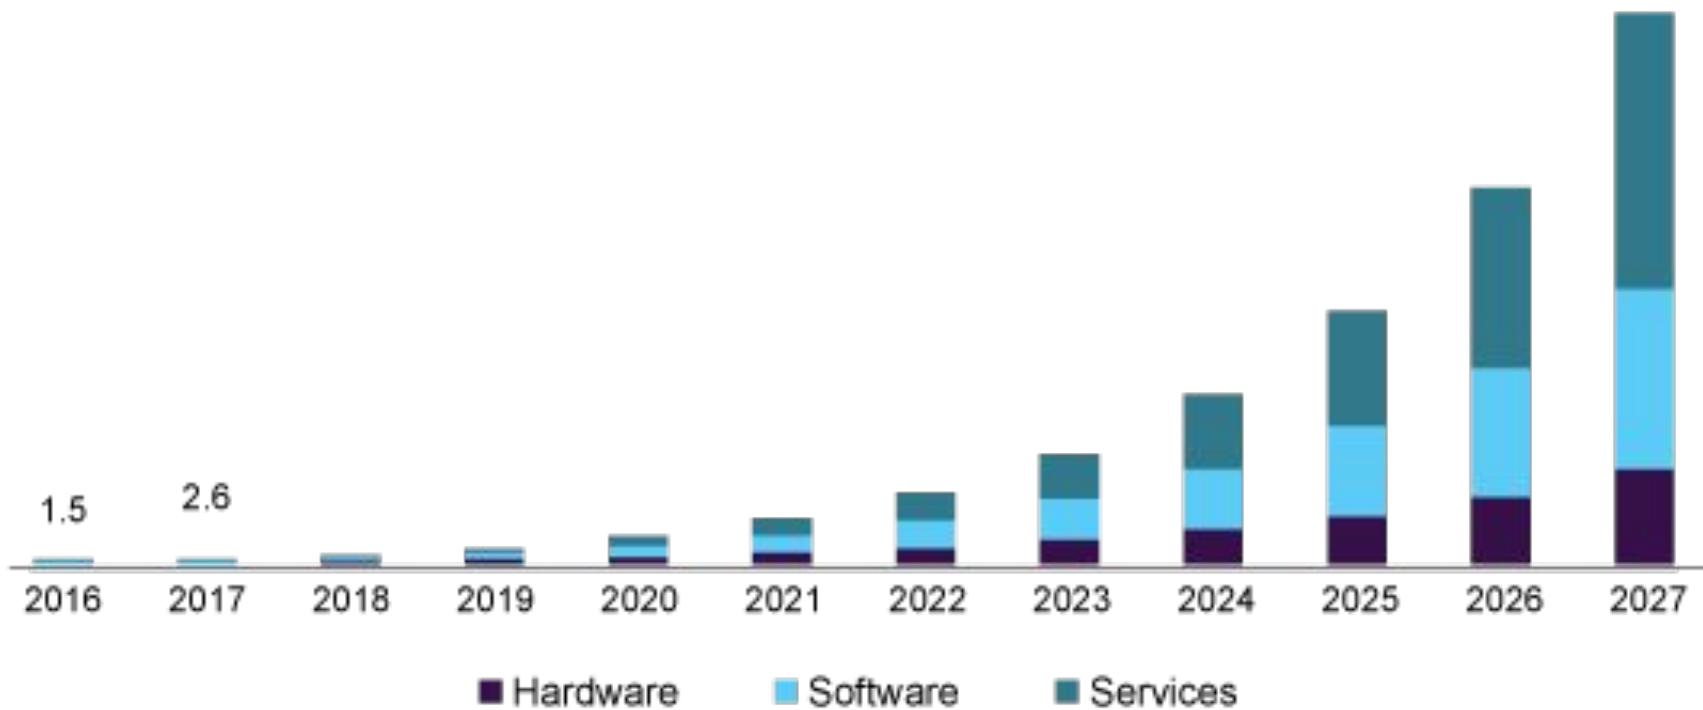
Why Deep Learning?

The performance of Deep Learning surpasses classical machine learning with millions and millions of data.



Projected AI Market Growth

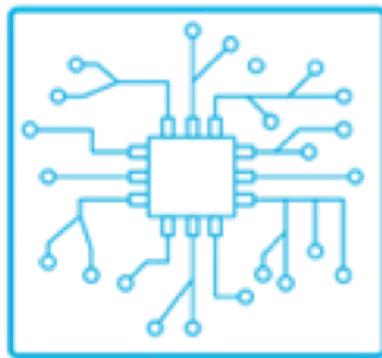
Projected Asia Pacific AI market size growth from 2016 to 2027 (USD Billion)



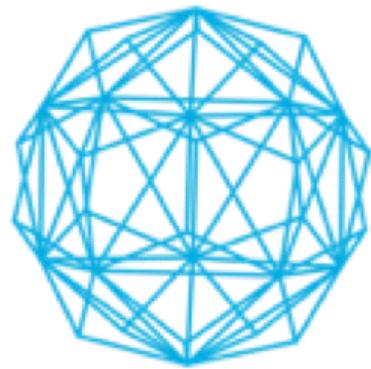
Source: www.grandviewresearch.com

Driving Forces of AI Growth

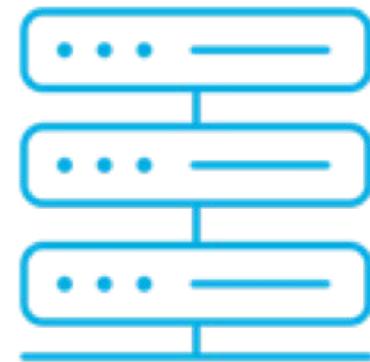
The three driving forces of AI growth is increasing computing power, algorithm and data availability



Computing Power



Algorithm Power



Data Availability

Types of Machine Learning

GAN Generative Models

Recommended Systems

Targetted Marketing

Customer Segmentation

UNSUPERVISED LEARNING

Structure Discovery

Meaningful compression

Big data Visualisation

CLUSTERING

DIMENSIONALLY REDUCTION

Feature Elicitation

Fraud Detection

Image Classification

CLASSIFICATION

Customer Retention

Diagnostics

SUPERVISED LEARNING

REGRESSION

Forecasting

Predictions

Process Optimization

New Insights

MACHINE LEARNING

REINFORCEMENT LEARNING

Real-Time Decisions

Game AI

Learning Tasks

REINFORCEMENT LEARNING

Robot Navigation

Skill Aquisition

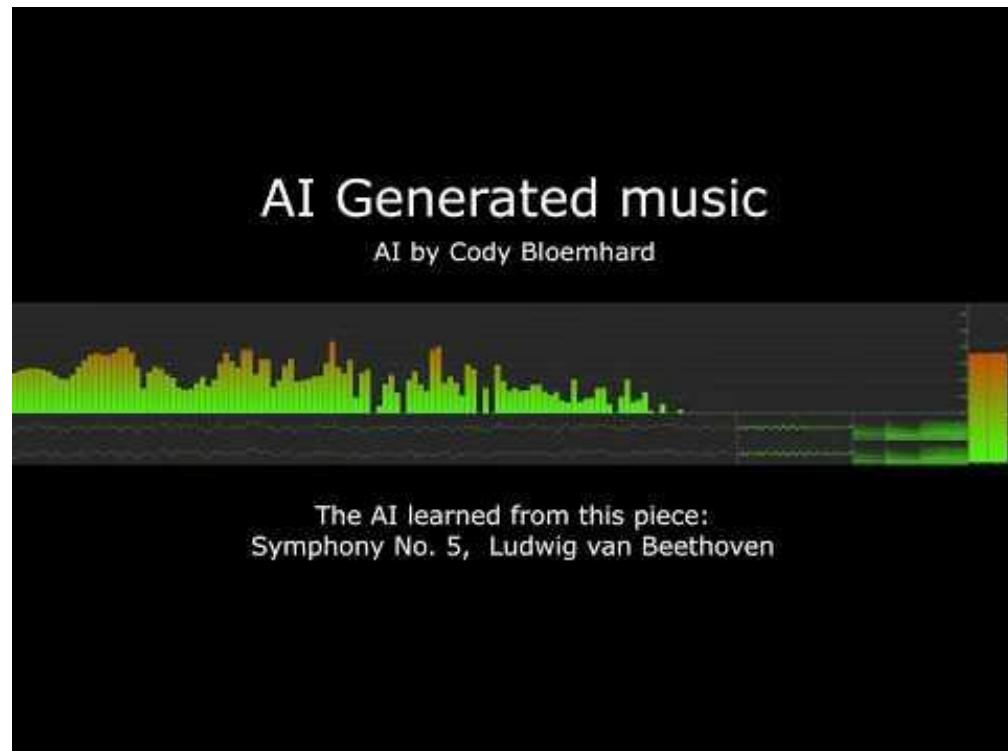
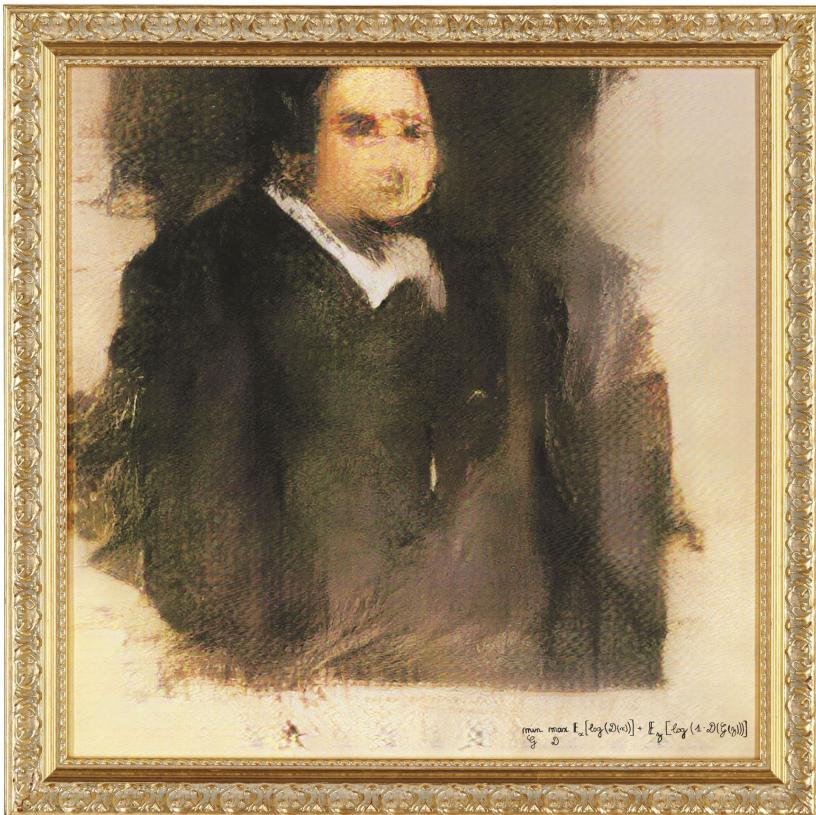
Semi Supervised Learning

Meta Learning

AI Generated Content - Big Business

AI-Generated Portrait Sells for \$432,500

<https://www.cbsnews.com/news/ai-generated-portrait-sells-for-stunning-432500-portrait-of-edmond-de-belamy/>



AI Generated Fake News



Machine Learning Frameworks

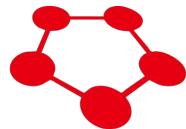
PyTorch

TensorFlow



scikit
learn

mxnet GLUON



Chainer
Caffe

torch

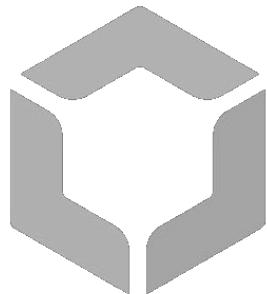
PaddlePaddle

theano

Machine Learning on Cloud



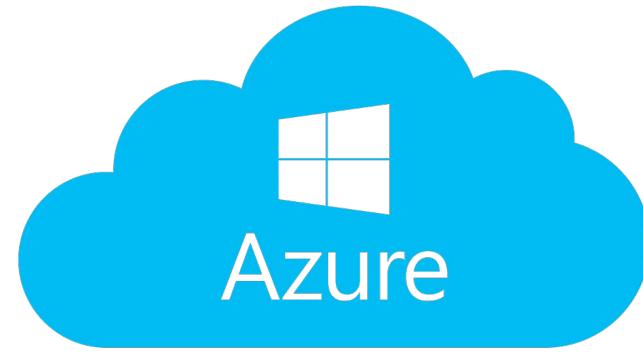
Google Cloud



百度云
cloud.baidu.com



IBM Cloud



 Alibaba Cloud

The Alibaba Cloud logo is a dark gray, stylized 'G' shape with a horizontal bar through the middle, followed by the text "Alibaba Cloud" in a dark gray, lowercase, sans-serif font.

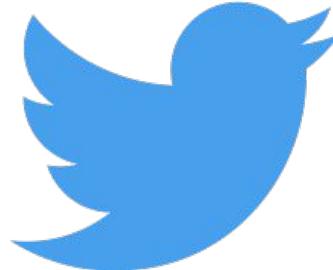
What is Pytorch?

- PyTorch (<http://pytorch.org/>) is a python package developed by Facebook
- The key features of Pytorch:
 - Tensor computation with strong GPU acceleration
 - Deep Neural Networks built on a reverse-mode automatic differentiation
- Pytorch is a rival DL platform against Tensorflow developed by Google



Who use Pytorch?

facebook®



nVIDIA®



Carnegie
Mellon
University



NYU

S Stanford
University

UNIVERSITE
PIERRE & MARIE CURIE
LA SCIENCE A PARIS

Install Pytorch on Colab

- Google Colaboratory is a free Jupyter notebook environment that requires no setup and runs entirely in the cloud.
- You can also access computing resources such as GPU and TPU for free
- You can get started with Google Colab from <https://colab.research.google.com/>
- Pytorch is pre-installed in Colab. You can re-install the latest version of using the command below in Colab

```
!pip3 install torch torchvision torchaudio
```

Setup GPU on Colab

To use GPU, Goto Edit->Notebook Setting, set hardware accelerator to GPU

Notebook settings

Hardware accelerator

GPU  

To get the most out of Colab, avoid using
a GPU unless you need one. [Learn more](#)

Omit code cell output when saving this notebook

CANCEL

SAVE

Check Pytorch Version

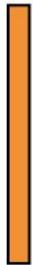
```
import torch  
print(torch.__version__)
```

Check GPU Resource

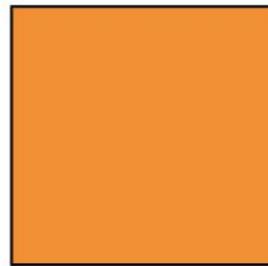
```
print("Cuda Current Device: ", torch.cuda.current_device())
print("Cuda Device Count: ", torch.cuda.device_count())
print("Cuda Device Name: ", torch.cuda.get_device_name(0))
print("Cuda Device Available : ", torch.cuda.is_available())
```

What is a Tensor?

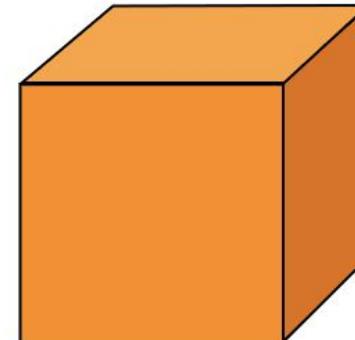
Tensor is a multi-dimensional array (data)



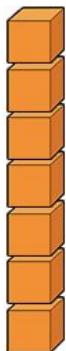
1d-Tensor



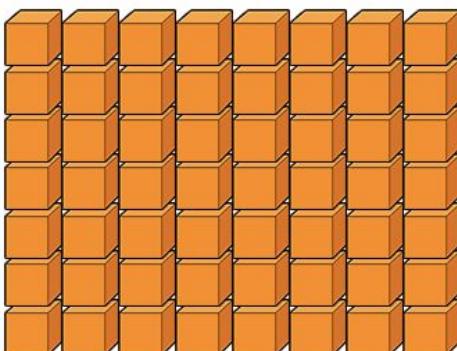
2d-Tensor



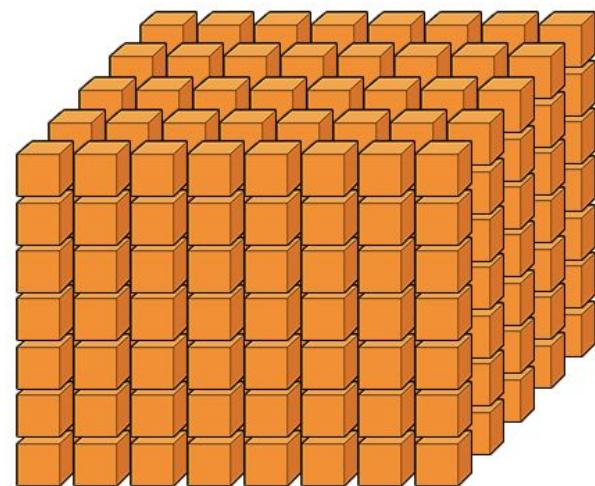
3d-Tensor



4d-Tensor



5d-Tensor



6d-Tensor

Create Torch Tensor

- Torch tensors are quite similar to Numpy arrays
- However, unlike Numpy arrays, Torch tensors are required for computing gradient (backpropagation) and GPU.
- There are two ways to create Torch tensor
 - `torch.Tensor(x)`
 - `torch.tensor(x)`
- `torch.tensor` is recommended as it can infer the data types from `x` (input) even for complex numbers.

Create Tensors from Python List

Tensors can be created from Python lists with the `torch.tensor()` function

Create a Vector

`a = [1, 2, 3]`

`b = torch.Tensor(a)` -> always convert to float data type

`b = torch.tensor(a)` -> keep to same data type (integer)

Create a Matrix

`a = [[1, 2, 3], [4, 5, 6]]`

`b = torch.Tensor(a)`

Torch & Numpy Conversion

- To convert torch tensor to numpy array
`b = a.numpy()`
- Convert numpy array to torch tensor:
`b = torch.from_numpy(a)`

Torch vs Numpy Functions

Numpy:

```
a = [1,2,3,4]  
b = np.array(a)  
c = np.sum(b)  
d = np.mean(b)  
e = np.max(b)
```

Torch:

```
a = [1.,2.,3.,4.]  
b = torch.tensor(a)  
c = torch.sum(b)  
d = torch.mean(b)  
e = torch.max(b)
```

Torch Math Operations

- `torch.abs(a)`
- `torch.pow(a, 2)`
- `torch.ceil(a)`
- `torch.floor(a)`
- `torch.add(a,b)`
- `torch.mul(a, b)`
- `torch.sum()`
- `torch.sqrt(a)`
- `torch.mean(a)`
- `torch.median(a)`
- `torch.eq(a,b)`
- `torch.gt(a,b)`
- `torch.le(a,b)`
- `torch.max(a)`
- `torch.min(a)`
- `torch.sort(a)`

Compute with GPU

```
device = torch.device('cuda' if torch.cuda.is_available()  
else 'cpu')
```

```
a = torch.tensor([1,1]).to(device)  
b = torch.tensor([2,2]).to(device)  
c = a+b
```

- To revert to cpu
c.cpu()

To convert to numpy, need to revert to cpu first
c.cpu().numpy()

Activity: Compute with GPU

Compute the following operations with GPU and convert to numpy array

a = 3

b = 4

c = 5

d = (a*b)+c

Ans: 17

Matrix Multiplication

- `torch.mm` performs a matrix multiplication of the matrices

Eg

```
mat1 = torch.randn(2, 3)
mat2 = torch.randn(3, 3)
torch.mm(mat1, mat2)
```

Activity: Matrix Multiplication

Given

$$x = [[1, 1]]$$

$$w = [[1, 2], [3, 4]]$$

$$b = [[2, 2]]$$

Compute the outcome of x^*w+b

Ans: [6,8]

Generate Special Torch Tensors

```
torch.zeros(2, 3) # zero matrix  
torch.eye(3) # diagonal matrix  
torch.linspace(1,5,10) # sequence equal spaced  
torch.range(0,10,2) # sequence with specific spacing  
torch.rand(5, 3) # random matrix from 0 to 1  
torch.randn(5, 3) # normal distributed random matrix
```

Torch Max

- `torch.max(a,dim)` returns the maximum value of each row of the input Tensor in the given dimension dim.
- The second return value is the index location of each maximum value found (argmax).

```
a = torch.Tensor([[1,0,0],[1,0,0],[0,1,0],[0,0,1]])  
print(torch.max(a,1))
```

Reshape Tensor

You can use `view` or `reshape` functions to reshape a tensor

Examples

```
a = torch.linspace(1,10,10).view(2,5)  
a = torch.linspace(1,10,10).reshape(2,5)
```

```
a = torch.linspace(1,10,10).view(-1,2)  
a = torch.linspace(1,10,10).reshape(-1,2)
```

Squeeze and Unsqueeze Dim

```
x = torch.linspace(0, 5, 5)
print(x)
x = torch.unsqueeze(x, dim=0) # row
print(x)
```

```
x = torch.linspace(0, 5, 5)
print(x)
x = torch.unsqueeze(x, dim=1) # column
print(x)
```

Concatenate

```
x = torch.tensor([1,2,3])  
y = torch.cat((x,x,x))
```

Transpose

```
x = torch.tensor([[1,2],[3,4]])  
y = torch.t(x)
```

Activity: Torch Operations

Perform the following torch operations in Pytorch

$x = [1,1]$ => convert to 1x2 matrix using unsqueeze

$w = [[1,2],[3,4]]$

$b = [[2],[2]]$ => transpose the vector

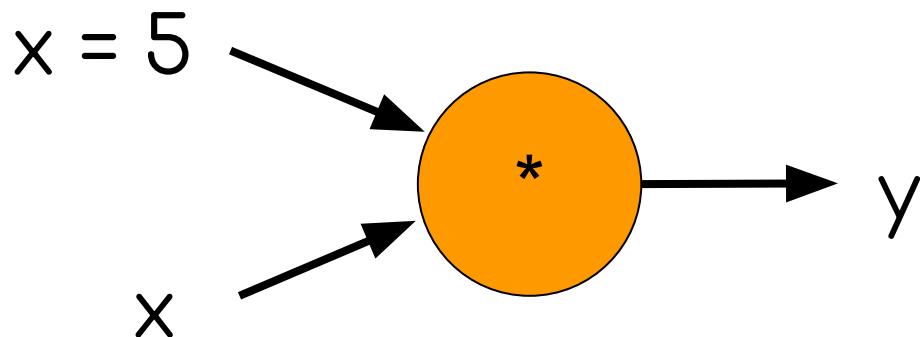
$y = x^*w+b$

Ans: [6, 8]

Gradient

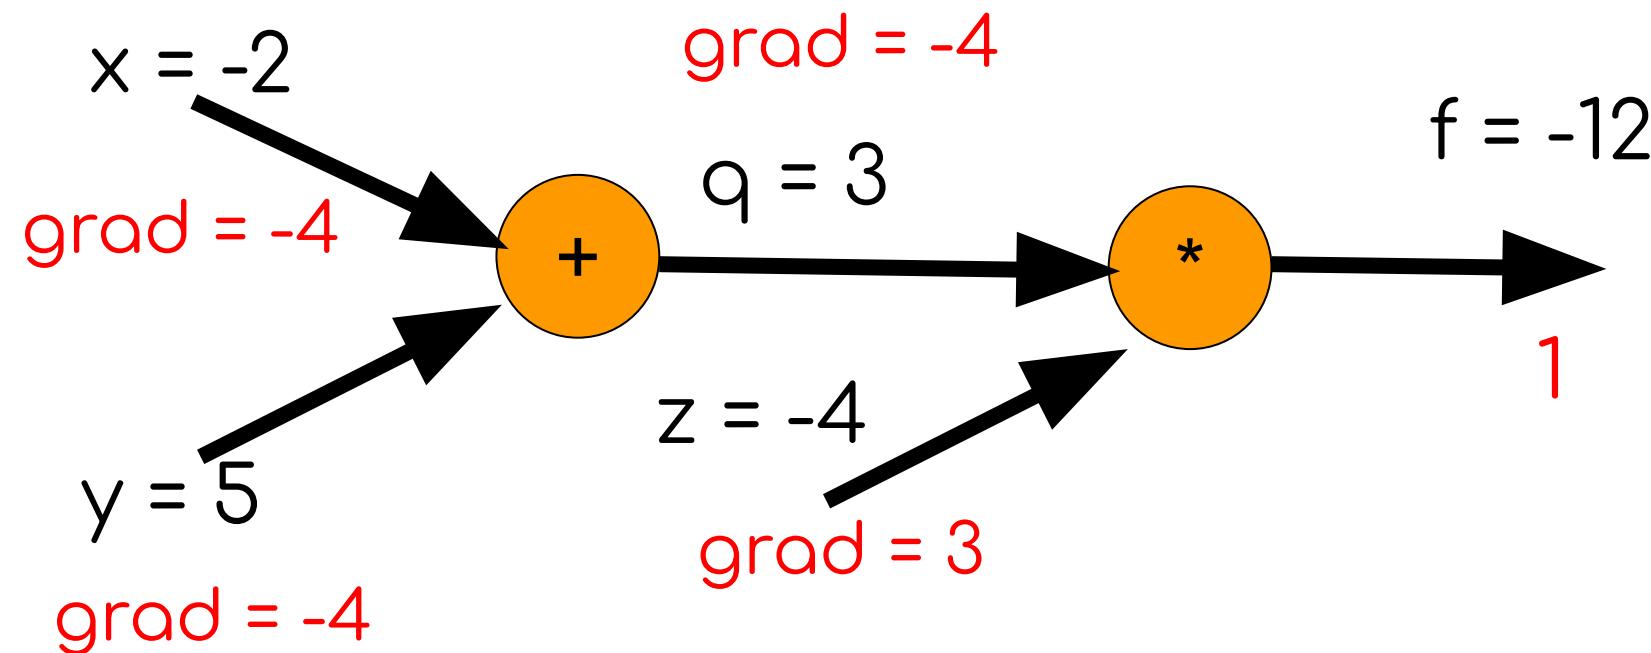
```
x = torch.tensor([5], requires_grad=True)  
y = x*x  
y.backward()  
x.grad.item()
```

$$x.\text{grad} = 2*x = 2*5=10$$



Compute Gradients

Compute gradients from computation graph
 $f = (x+y)*z = q*z$



Compute Gradients

```
x = torch.tensor([-2.], requires_grad=True)
y = torch.tensor([5.], requires_grad=True)
z = torch.tensor([-4.], requires_grad=True)
f = (x+y)*z

f.backward()
```

```
print('x gradient = ',x.grad.item())
print('y gradient = ',y.grad.item())
print('z gradient = ',z.grad.item())
```

Difference between .item and .data

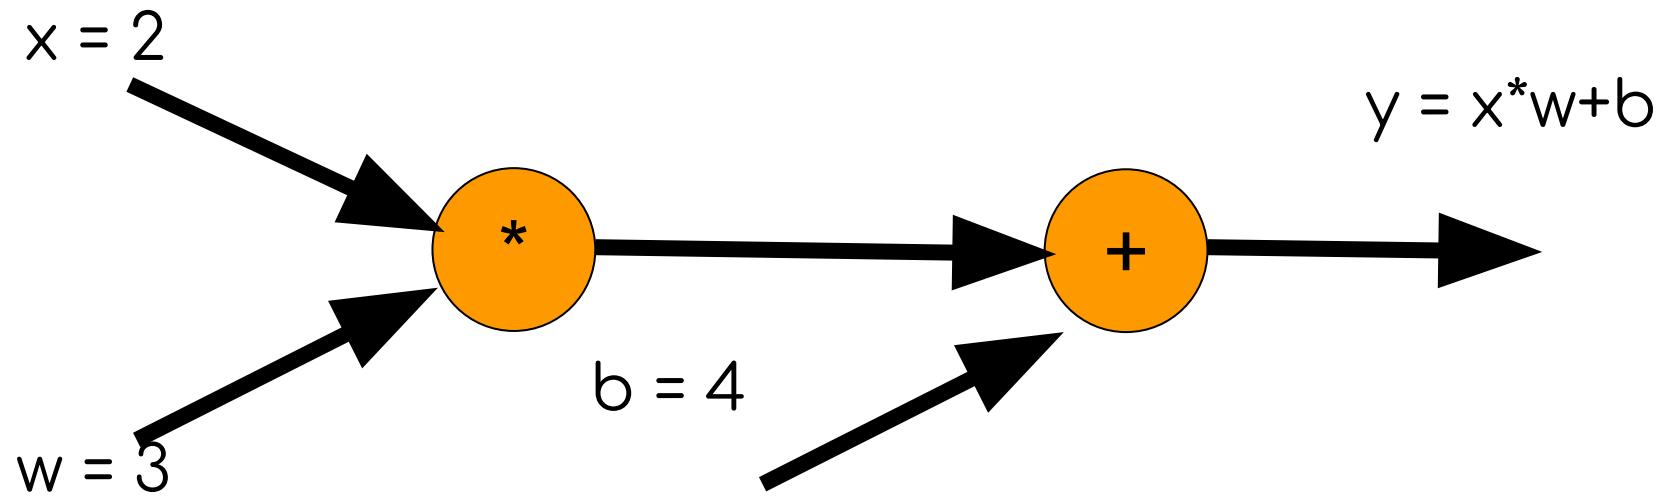
```
a = torch.randn(1)
print(a.item())
print(a.data)
print(a)
```

Output

```
-1.1979358196258545
tensor([-1.1979])
tensor([-1.1979])
```

Activity: Compute Gradients

Compute gradients at w and b for $y = x^*w+b$



Answer:

$$x \text{ gradient} = 3.0$$

$$w \text{ gradient} = 2.0$$

$$b \text{ gradient} = 1.0$$

Activity: Compute Gradients

Compute gradients for the following functions for $x=1$

$$z = 2y^2$$

$$y = x^3$$

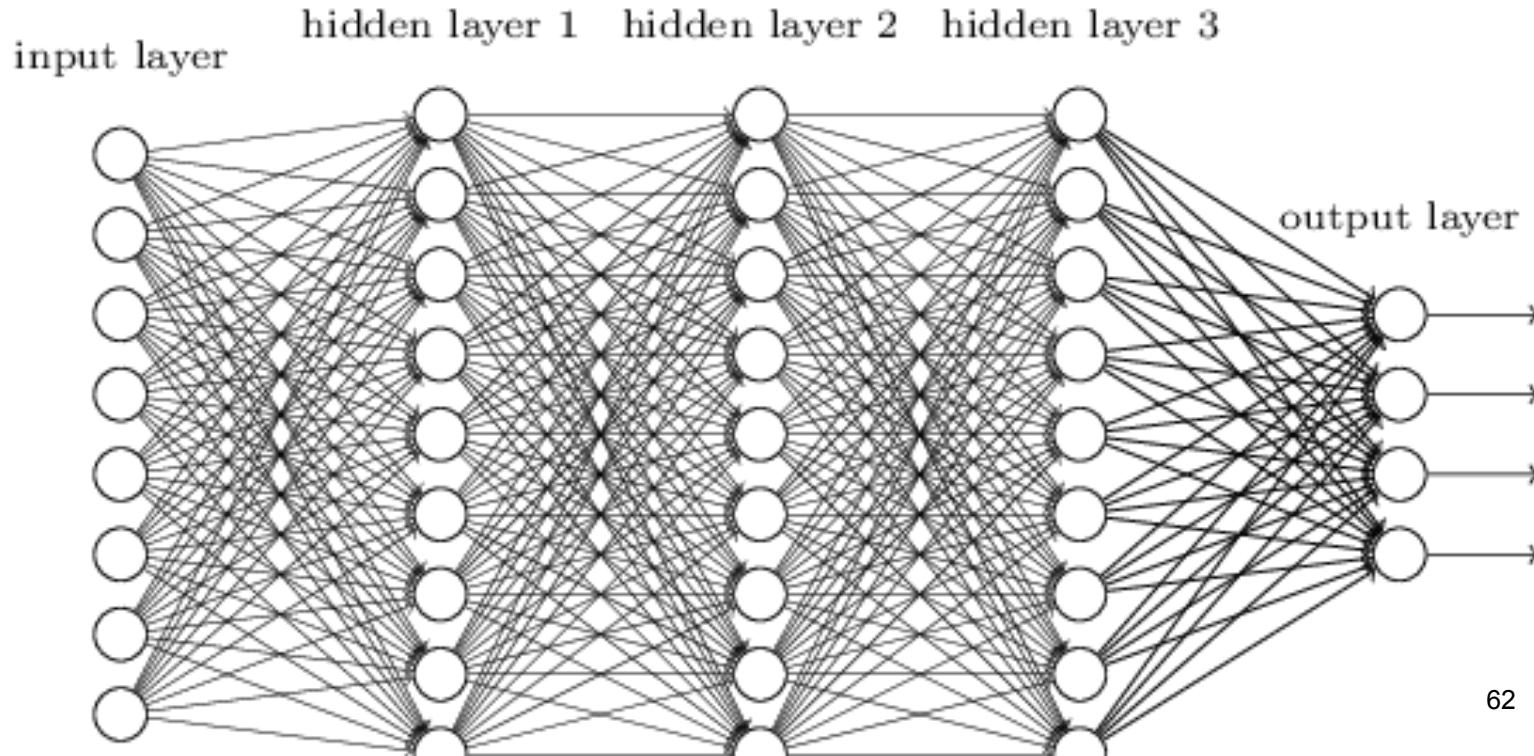
Answer: 48

Topic 2

Neural Network for Regression

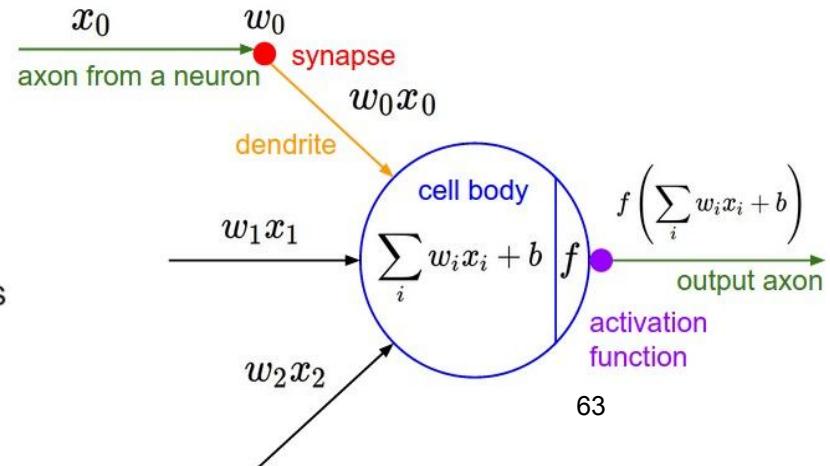
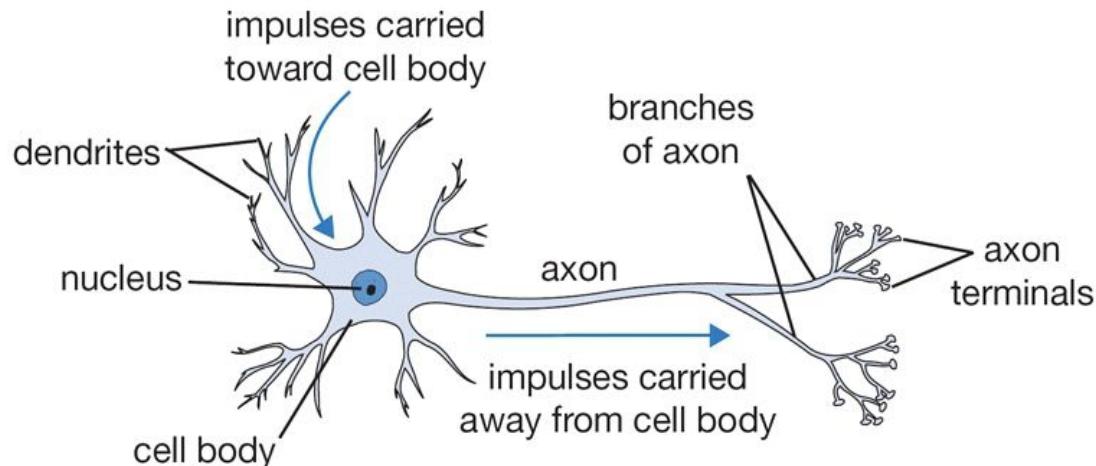
Neural Network (NN)

- Neural Networks are a set of algorithms, modeled loosely after the human brain, that are designed to recognize patterns
- NN is made up of neurons. It consists of input layer, multiple hidden layers and output layers. The simplest NN is one hidden layer.



What is a Neuron?

- A biological neuron has dendrites to receive signals, a cell body to process them, and an axon to send signals out to other neurons
- An artificial neuron has a number of input channels, a processing stage, and one output that can fan out to multiple other artificial neurons.
- The weighted sum of inputs + bias is feed through a **non-linear activation function** to output a value to next neurons.

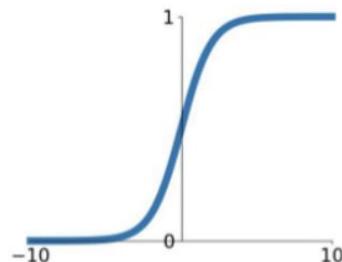


Activation Functions

Activation functions allow Neural Networks to learn complicated and Non-linear complex functional mappings between the inputs and outputs.

Sigmoid

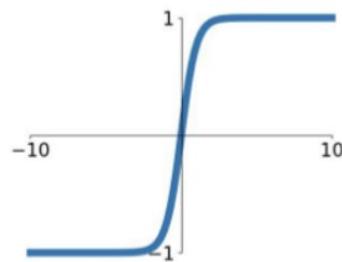
$$\sigma(x) = \frac{1}{1+e^{-x}}$$



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

tanh

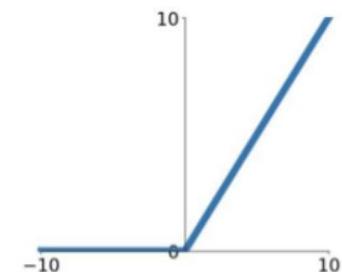
$$\tanh(x)$$



$$\sigma(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

ReLU

$$\max(0, x)$$



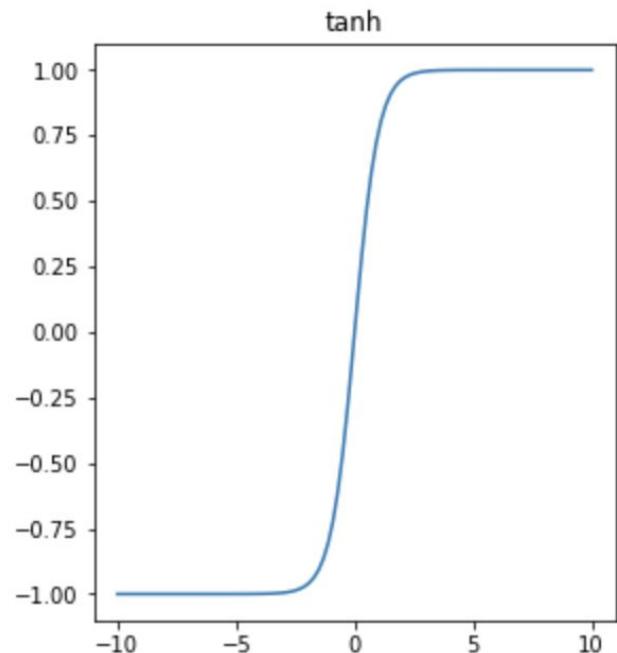
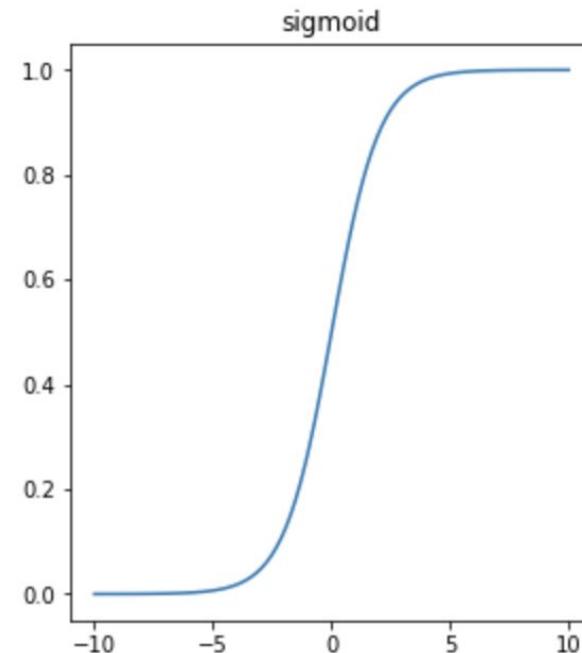
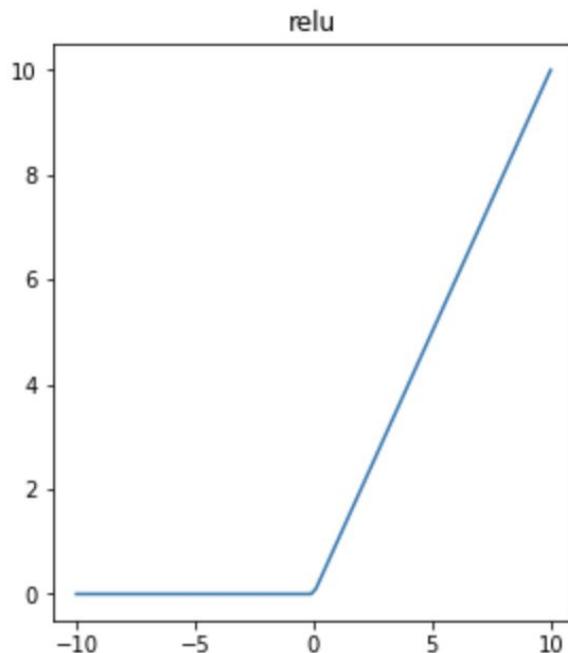
$$\sigma(z) = \max(0, z)$$

Activation Functions

`torch.relu(x)`

`torch.sigmoid(x)`

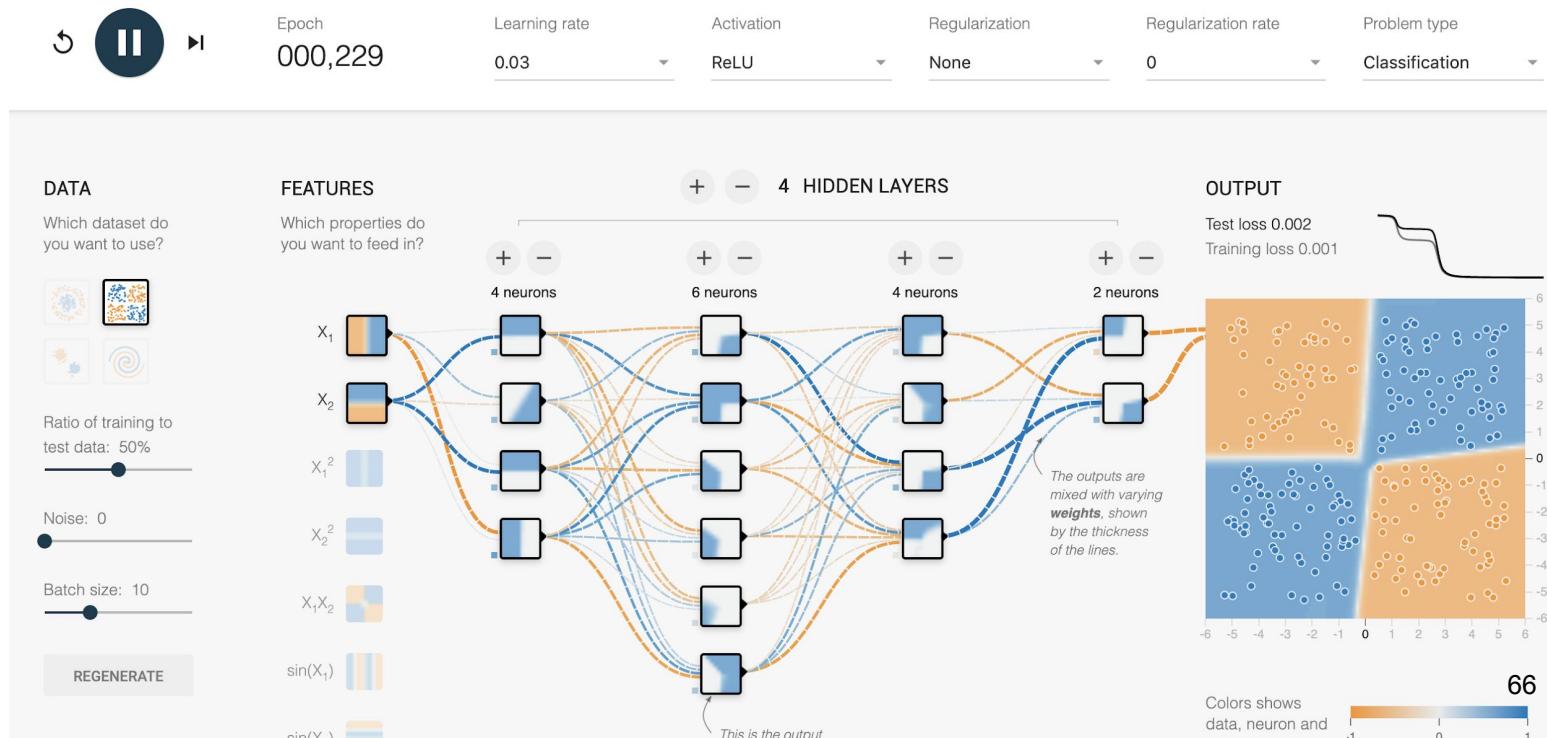
`torch.tanh(x)`



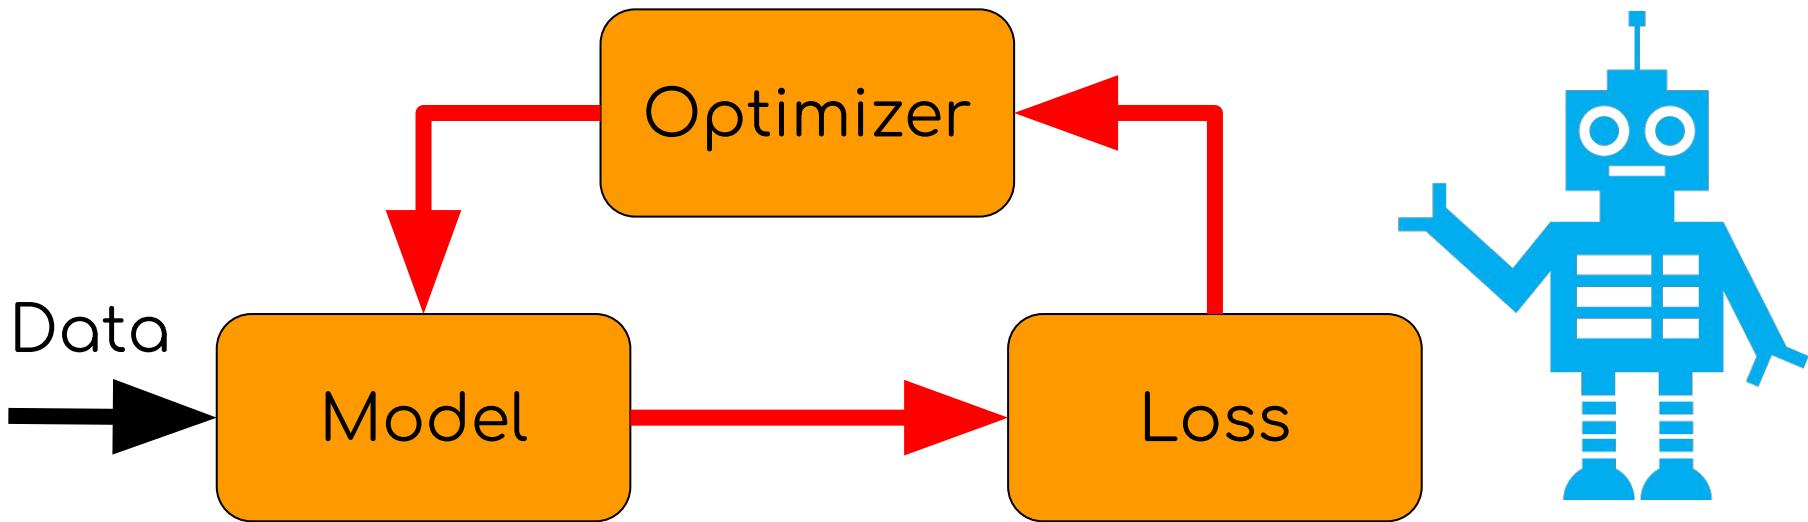
Activity: Tensorflow Playground

- Try different activation functions (sigmoid vs relu) on Tensorflow playground
<https://playground.tensorflow.org/>
- Hover over the connections to check the weights

Tinker With a **Neural Network** Right Here in Your Browser.
Don't Worry, You Can't Break It. We Promise.



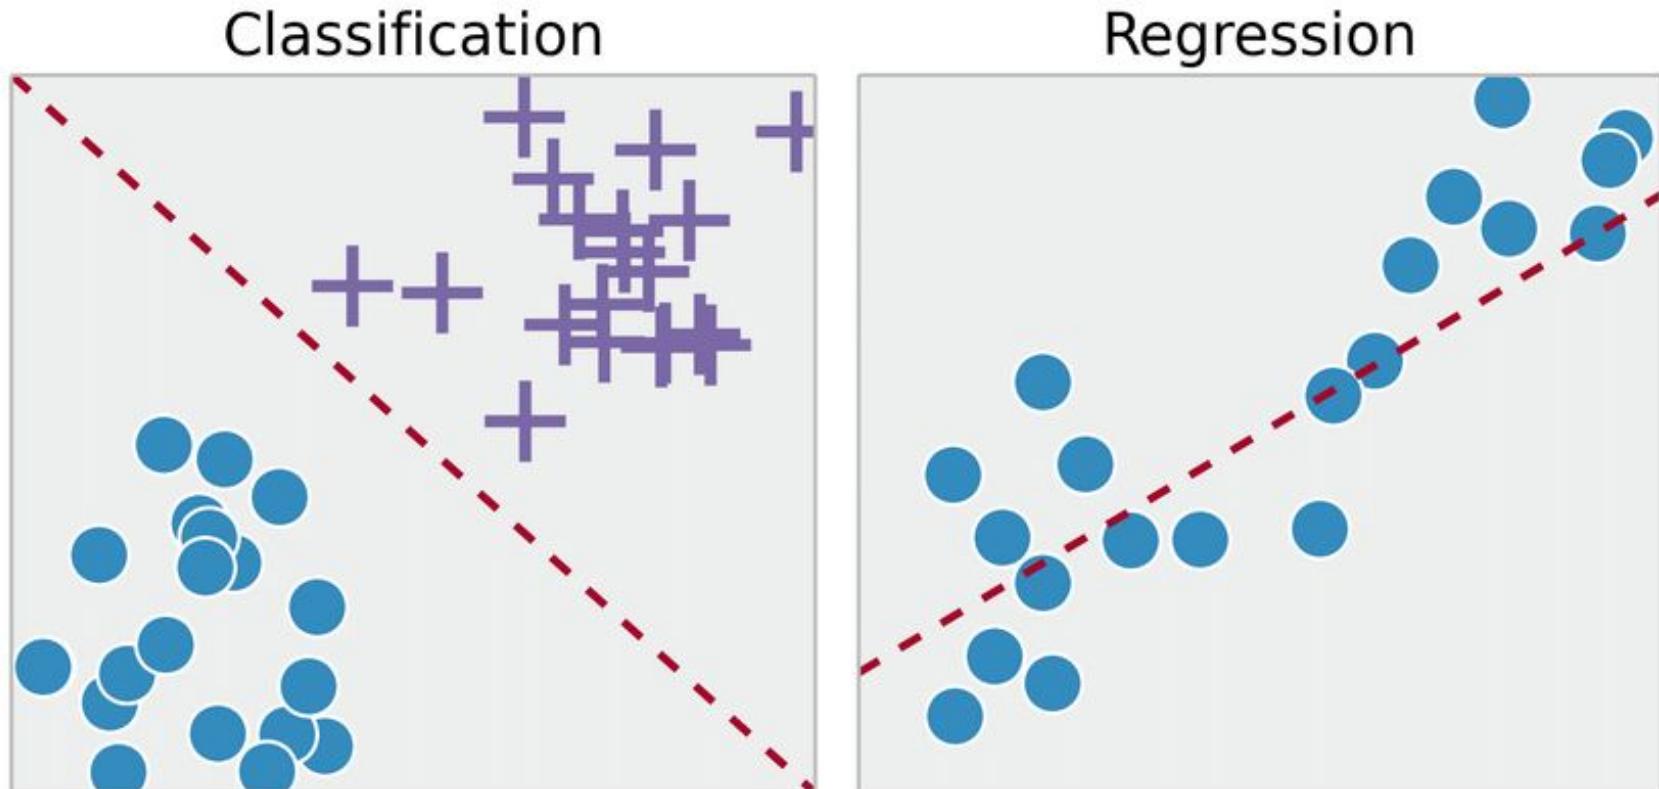
Machine Learning (Trial & Error) Approach



- Deep learning is a trial & error, brute force and blackbox approach to learn the patterns and make predictions
- The approach requires 3 components: NN Model, Loss Function, and Optimizer
- The learning (training) is driven by data to reduce the loss in successive iterations.
- The optimizer is to minimize the loss function of a model by adjusting the weights (parameters). ⁶⁷

What is Regression?

- Regression is predicting a continuous-valued attribute associated with an object.
- Regression is a supervised learning method. During training, actual values are provided.



Machine Learning Steps

1. Preprocessing Data
2. Define Model
3. Select Optimizer
4. Define Loss
5. Training Loop

Step 1: Preprocessing Data

```
import torch
```

```
X = torch.tensor([1.,2.,3,4,5])  
y = torch.tensor([0,-1.1,-1.8,-3.1,-4.5])
```

```
W = torch.rand(1,requires_grad=True)  
b = torch.rand(1,requires_grad=True)
```

Step 2: Select Optimizer

```
learning_rate = 0.001
```

```
optimizer = torch.optim.SGD([W,b], lr=learning_rate)
```

Step 3 to 4 Define Model and Loss

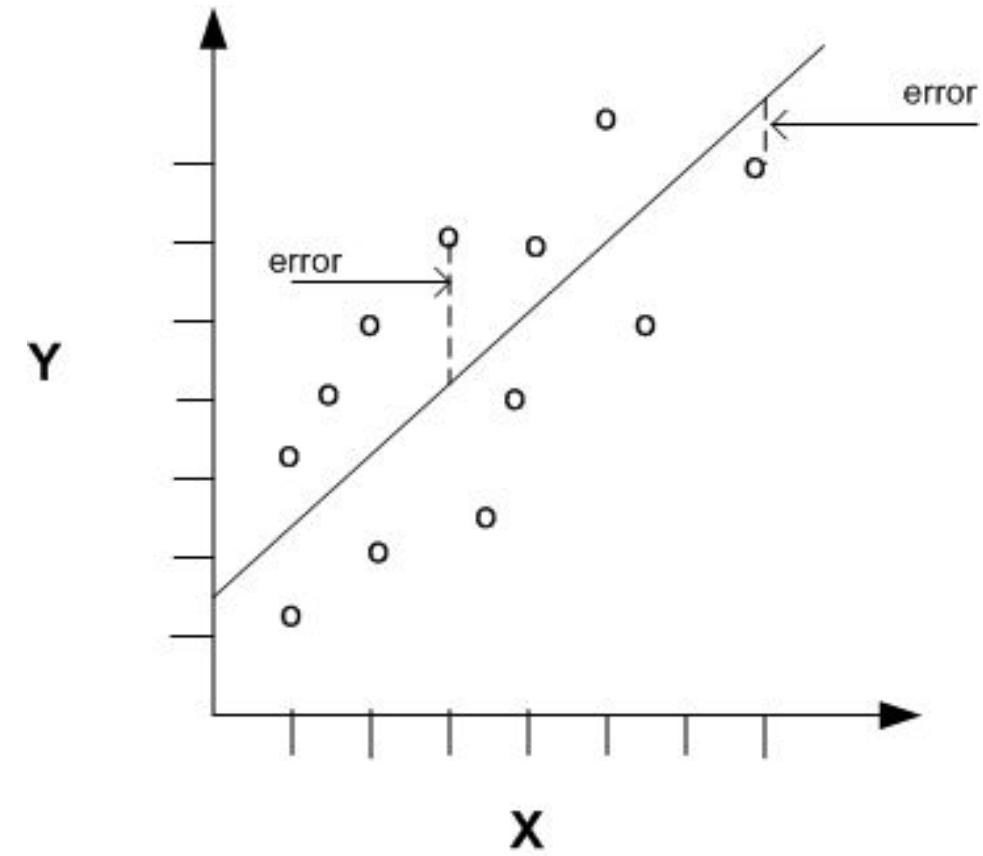
```
yhat = X*W+b # Linear Model
```

```
loss = (yhat-y).pow(2).sum() #Mean Square Error
```

Alternatively, you can use the Pytorch MSELoss function as follows:

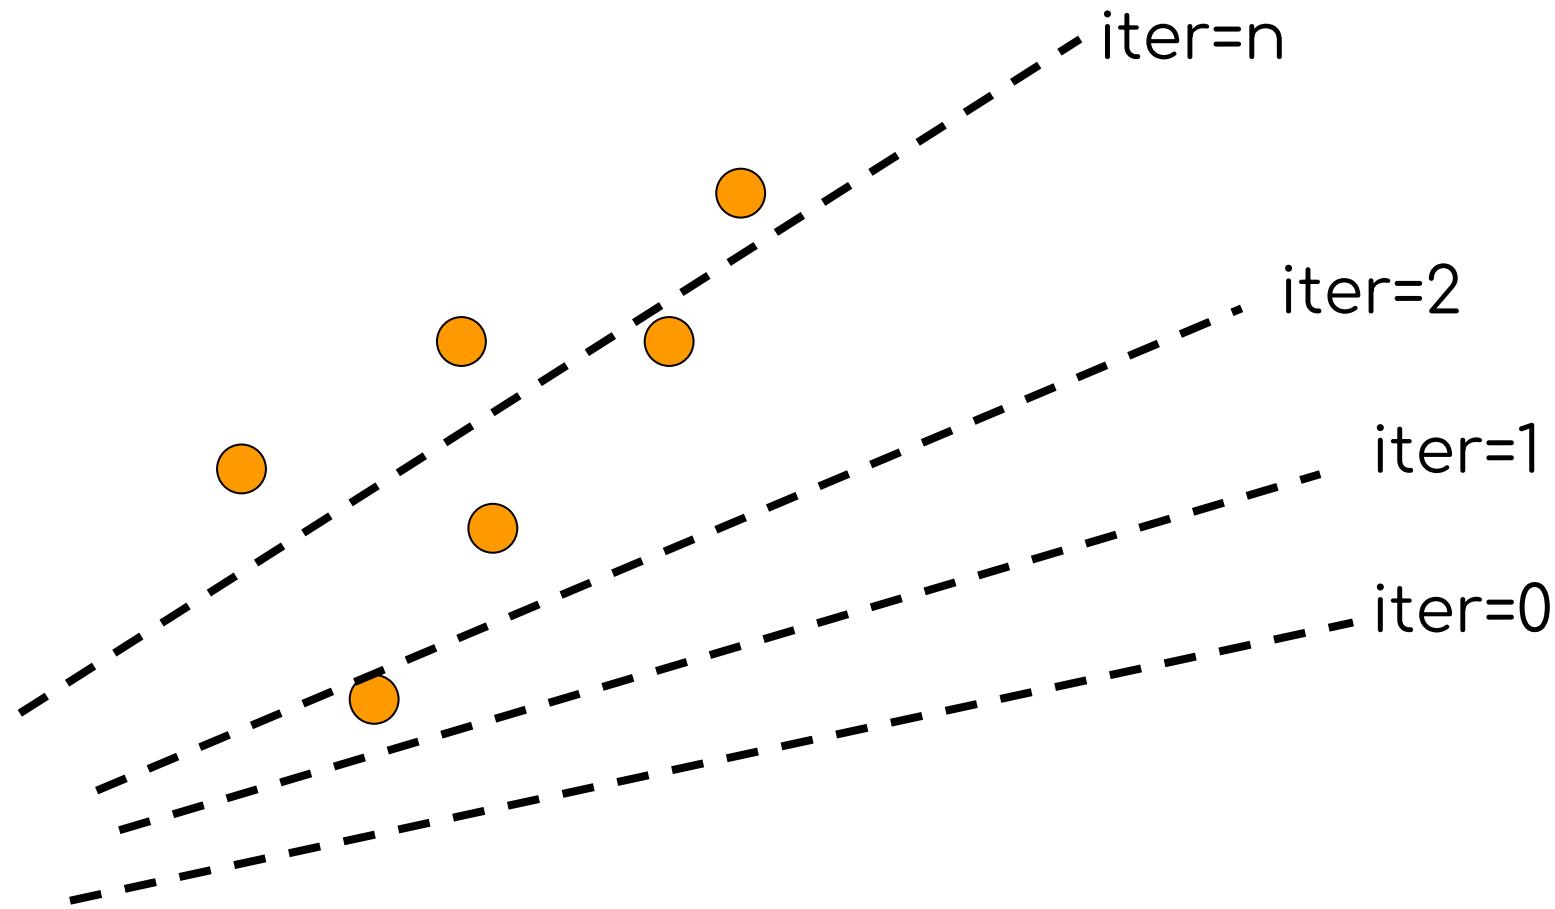
```
criterion = torch.nn.MSELoss()  
loss = criterion(yhat,y)
```

Mean Square Error



$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

ML Approach to Linear Regression



ML determine the the best fit curve by reducing the MSE to minimum

Step 5: Training Loop

```
for i in range(1000):
```

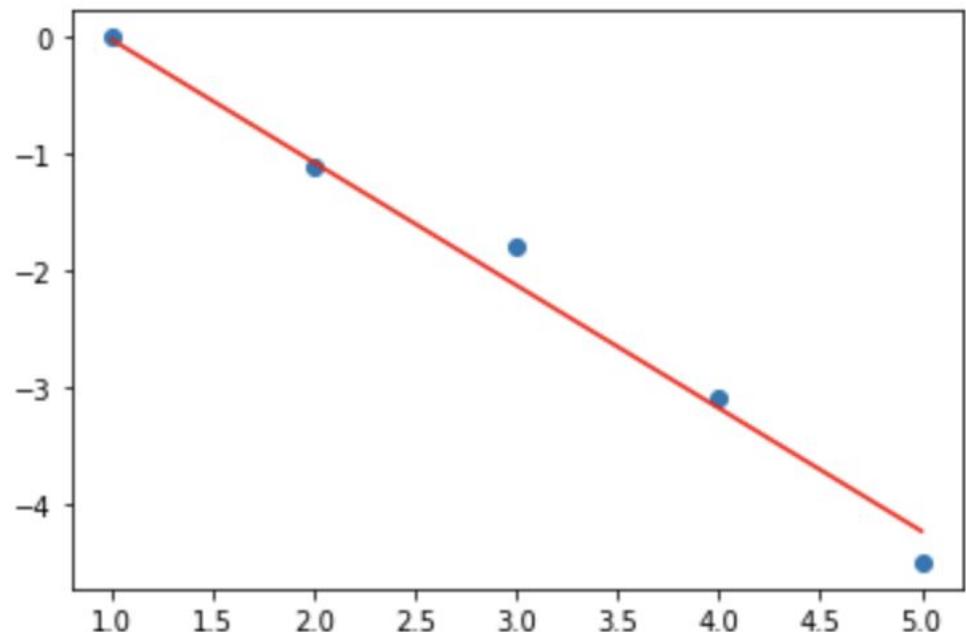
```
    yhat = X*W+b
```

```
    loss = (yhat-y).pow(2).sum()
```

```
    optimizer.zero_grad()
```

```
    loss.backward()
```

```
    optimizer.step()
```



Zero Grad Explanation

- In PyTorch, we need to set the gradients to zero before starting to do backpropagation because PyTorch accumulates the gradients on subsequent backward passes.
- This is convenient while training RNNs. Because of this, when you start your training loop, ideally you should zero out the gradients so that you do the parameter update correctly.
- Else the gradient would point in some other direction than the intended direction towards the minimum

Explicit Zero Grad Demo

```
import torch  
learning_rate = 0.001
```

```
X = torch.tensor([1.,2.,3,4,5])  
y = torch.tensor([0,-1.1,-1.8,-3.1,-4.5])  
W = torch.randn(1, requires_grad=True)  
b = torch.randn(1, requires_grad=True)
```

```
for i in range(1000):  
    yhat = X*W+b  
    loss = (yhat - y).pow(2).sum()  
    loss.backward()
```

```
W.data -= learning_rate * W.grad.item()  
b.data -= learning_rate * b.grad.item()
```

```
W.grad.data.zero_()  
b.grad.data.zero_()
```

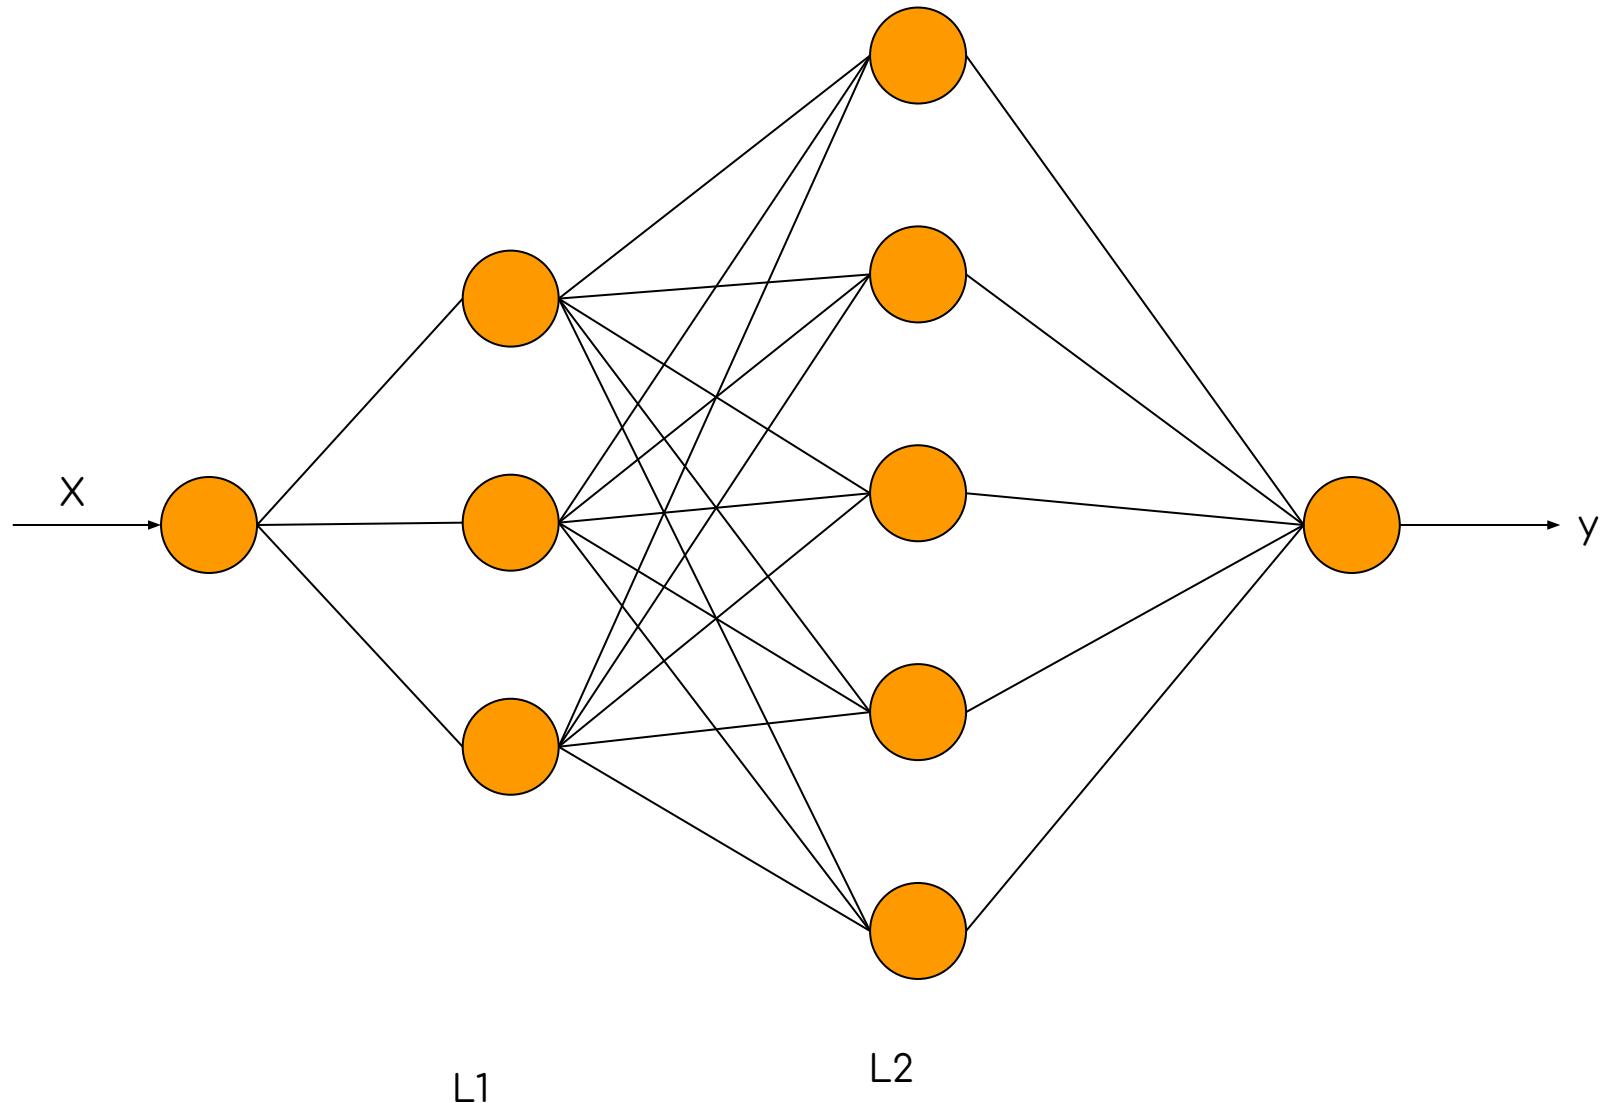
Similar to `optimizer.step()`

Similar to `zero_grad()`

```
print(f'i:{i}, W:{W.item()}, b:{b.item()}, loss:{loss.item()}')
```

```
)
```

Linear Regression NN Model



Step 1: Preprocessing Data (NN)

```
import torch
```

```
X = torch.tensor([1.,2.,3,4,5])  
y = torch.tensor([0,-1.1,-1.8,-3.1,-4.5])
```

```
X = torch.unsqueeze(X, dim=1)  
y = torch.unsqueeze(y, dim=1)
```

Step 2 Define Model (NN)

```
import torch
import torch.nn as nn
import torch.nn.functional as F

L1 = 3
L2 = 5
class Model(nn.Module):
    def __init__(self):
        super(Model,self).__init__()
        self.fc1 = nn.Linear(1,L1)
        self.fc2 = nn.Linear(L1,L2)
        self.fc3 = nn.Linear(L2,1)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

model = Model()
print(model)
```

Alternate Model Definition

Step 2: Define Model

```
import torch
import torch.nn as nn
import torch.nn.functional as F
```

```
L1 = 3
L2 = 5
```

```
model = nn.Sequential(nn.Linear(1, L1),
                      nn.ReLU(),
                      nn.Linear(L1, L2),
                      nn.ReLU(),
                      nn.Linear(L2, 1))

print(model)
```

Step 3 Select Optimizer

```
learning_rate = 1e-2
```

```
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
```

Step 4 Training

```
for i in range(1000):
```

```
    yhat = model(X)
```

```
    loss = (yhat - y).pow(2).sum()
```

```
    optimizer.zero_grad()
```

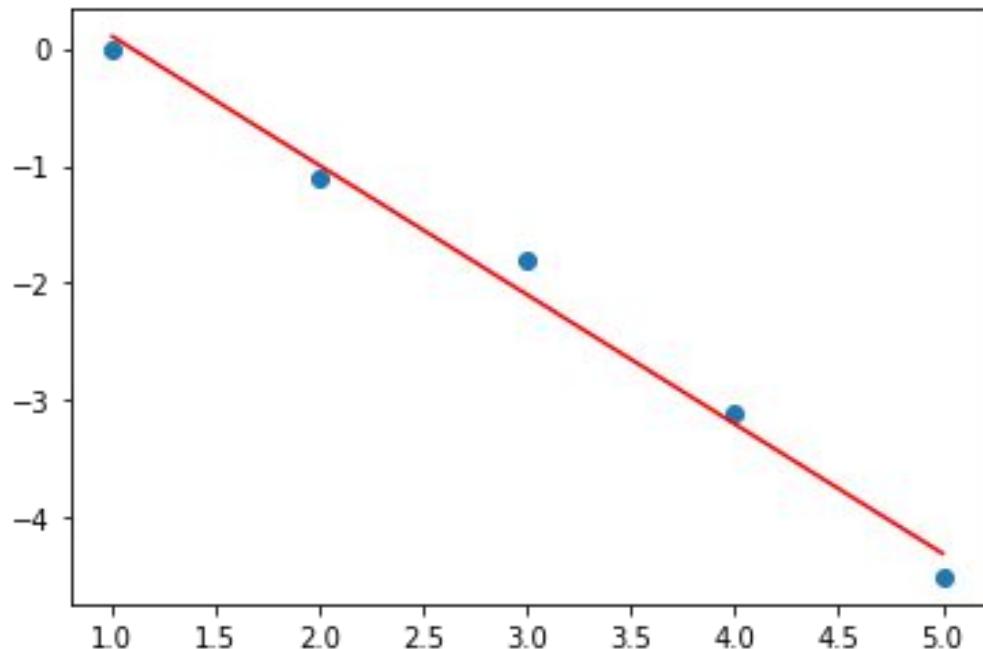
```
    loss.backward()
```

```
    optimizer.step()
```

Step 5: Evaluate Model

```
import matplotlib.pyplot as plt
```

```
plt.plot(X,y,'o')  
plt.plot(X,model(X).data,'r')  
plt.show()
```



Demo: Predictive Model for Housing Price

This is a demo of how to build a predictive regression model for housing price. The steps are as follows:

- Prepare the data
 - Import raw data
 - Clean data eg remove missing data
 - Normalize the data
 - Split raw data to train and test dataset
 - Create input and output
- Build the model
 - Create a neural network model
 - Compile the model (optimizer, loss function)
 - Inspect the model
- Train the model
- Evaluate the model
- Make prediction

Import the Data

We will get some raw data from github Eg boston hosting price

```
import pandas as pd  
dataset_path =  
"https://raw.githubusercontent.com/tertiarycourses/  
datasets/master/boston.csv"
```

Prepare the Data for Training

- Split the data to train and test dataset

```
x_train = dataset.sample(frac=0.7,random_state=0)  
x_test = dataset.drop(x_train.index)
```

- Create input and output

```
y_train = x_train.pop('medv')  
y_test = x_test.pop('medv')
```

Clean the Data

- Remove missing data

```
dataset = dataset.dropna()
```

- Normalize the data

```
x_train = (x_train - x_train.mean()) / x_train.std()
```

```
x_test = (x_test - x_test.mean()) / x_test.std()
```

Convert Data to Tensors

```
import torch
```

```
x_train = torch.Tensor(x_train.values)
y_train = torch.Tensor(y_train.values)
```

```
x_test = torch.Tensor(x_test.values)
y_test = torch.Tensor(y_test.values)
```

Increase Tensor Dimension

```
y_train = torch.unsqueeze(y_train, dim=1)  
y_test = torch.unsqueeze(y_test, dim=1)
```

Build the Model

```
class Model(nn.Module):  
  
    def __init__(self):  
        super(Model,self).__init__()  
        self.fc1 = nn.Linear(13,L1)  
        self.fc2 = nn.Linear(L1,L2)  
        self.fc3 = nn.Linear(L2,1)  
  
    def forward(self, x):  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return x  
  
model = Model()
```

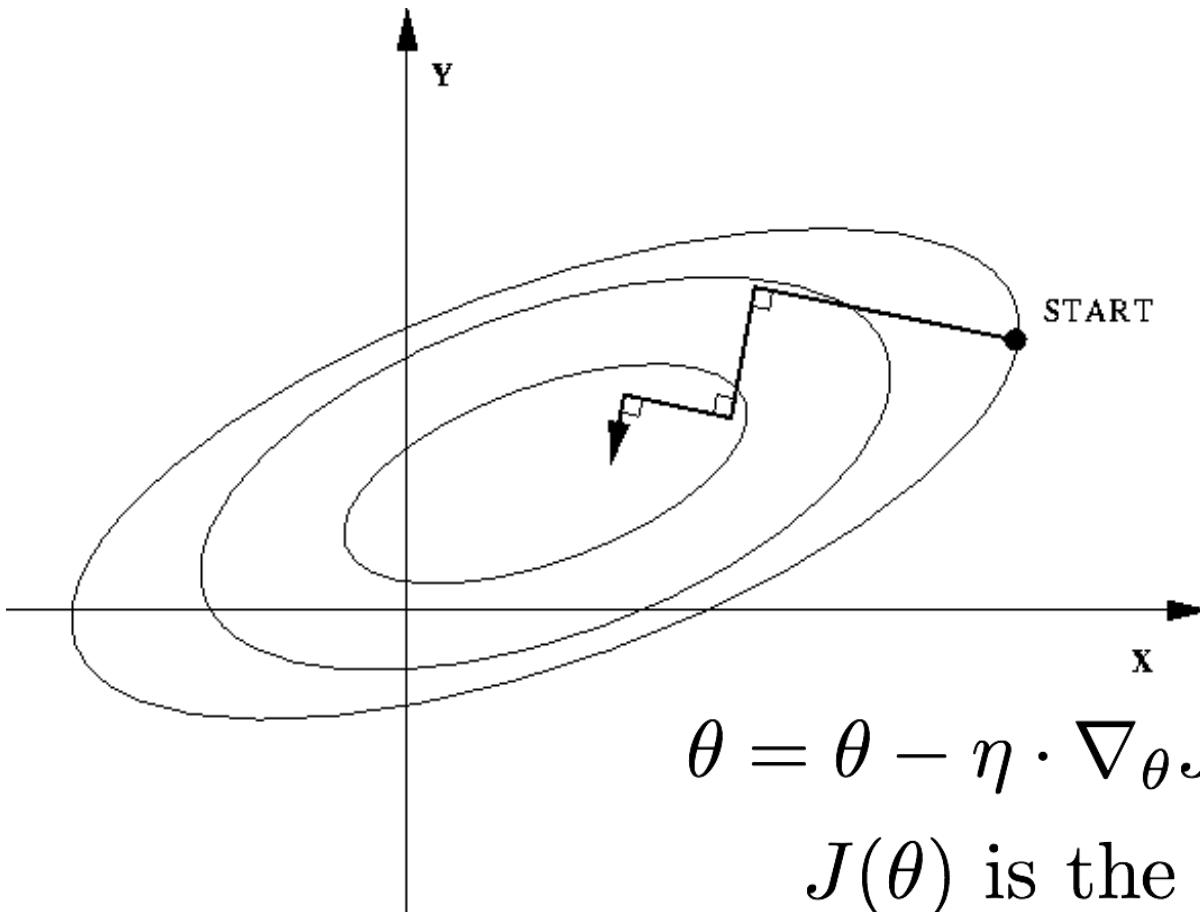
Define the Loss and Optimizer

```
criterion = nn.MSELoss()
```

```
learning_rate = 0.001
```

```
optimizer = torch.optim.SGD(model.parameters(),  
lr=learning_rate)
```

GradientDescentOptimizer (SGD)



$$\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta; x^{(i)}; y^{(i)})$$

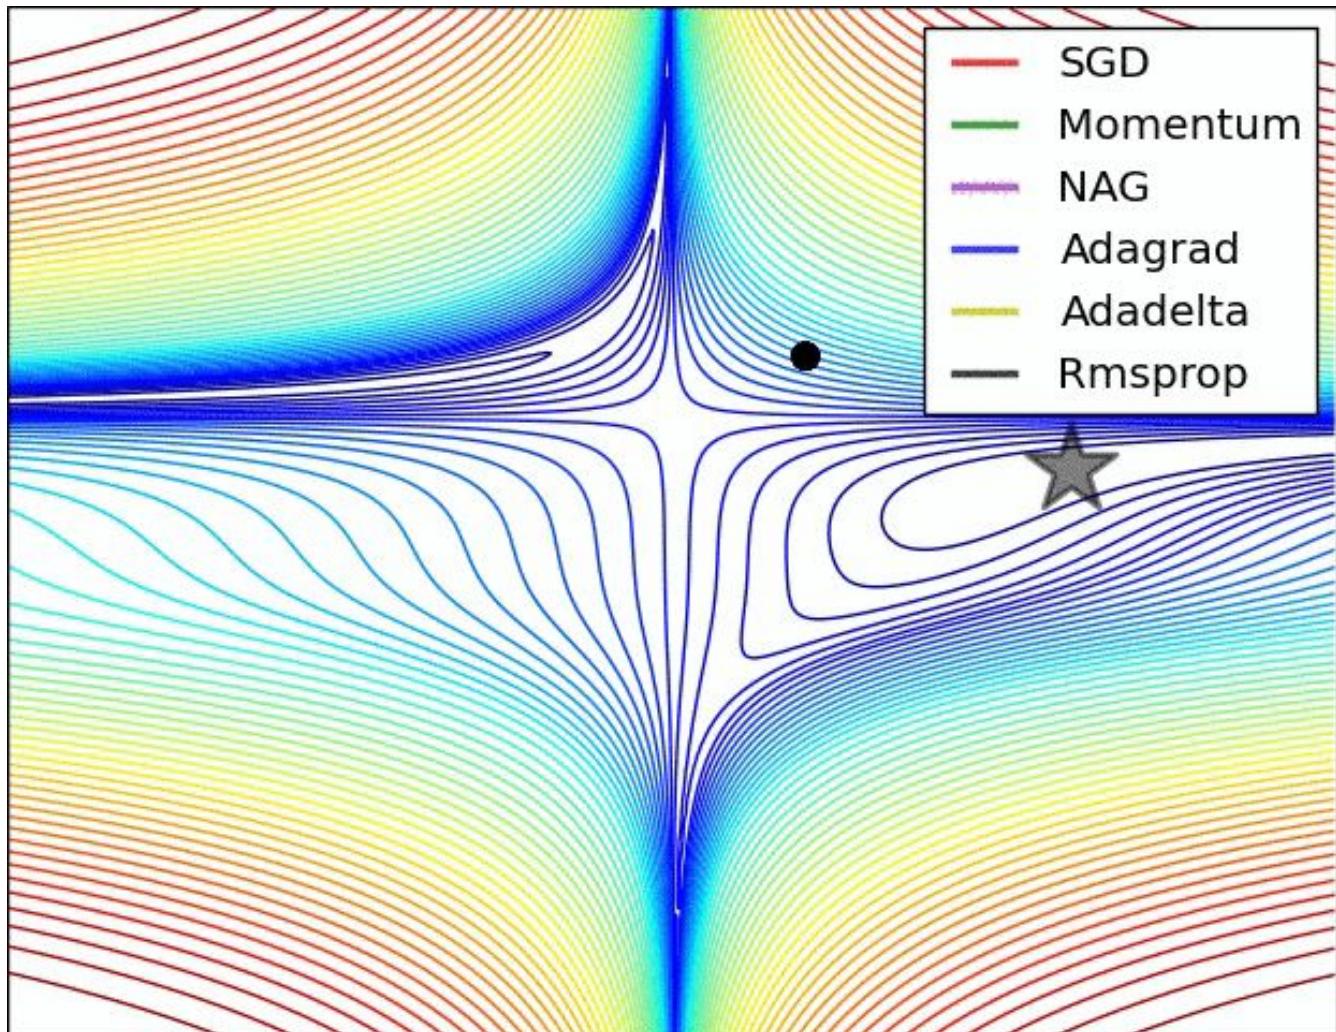
$J(\theta)$ is the loss function

η is the learning rate

θ are the model parameters

Other Popular Optimizers

- Adam
- RMSProp
- Adagrad
- Adadelta
- Momentum
- NAG



Train the Model

```
for i in range(1000):
```

```
# Model prediction  
yhat = model(x_train)
```

```
# Compute loss  
loss = criterion(yhat, y_train)
```

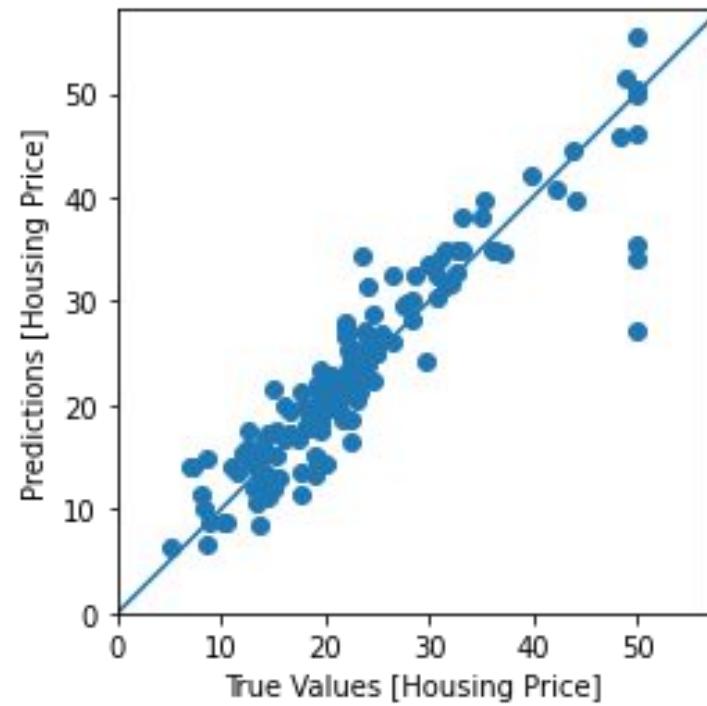
```
# Compute gradients and update parameters  
optimizer.zero_grad()  
loss.backward()  
optimizer.step()
```

we need to manually set the gradients to zero using `optimizer.zero_grad()` since PyTorch by default accumulates gradients.

Evaluate the Model

```
import matplotlib.pyplot as plt
```

```
plt.scatter(y_test, model(x_test).data)
plt.xlabel('True Values [Housing Price]')
plt.ylabel('Predictions [Housing Price]')
plt.axis('equal')
plt.axis('square')
plt.xlim([0,plt.xlim()[1]])
plt.ylim([0,plt.ylim()[1]])
plt.plot([0, 100], [0, 100])
```



Save and Load the Model

```
# Save the model
```

```
torch.save(model,'./regression.pkl')
```

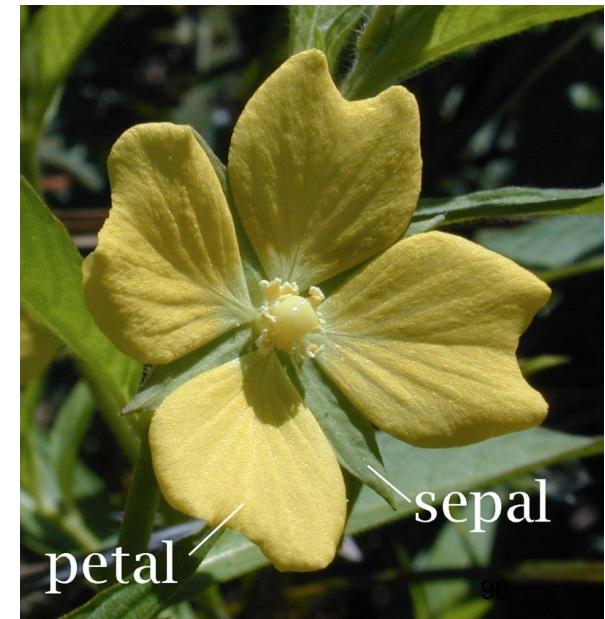
```
# Load the model
```

```
new_model=torch.load('./regression.pkl')
```

Activity: Predictive Regression Model

- Import the iris.dataset from github
<https://raw.githubusercontent.com/tertiarycourses/datasets/master/iris.csv>
- Create a NN model is with 3 hidden layers (64,32,32) to predict the sepal width from sepal length, petal length and petal width
- Perform all the steps as shown in the demo
 - Prepare the data
 - Build the model
 - Train the model
 - Make prediction
 - Save the model

Time: 30 mins



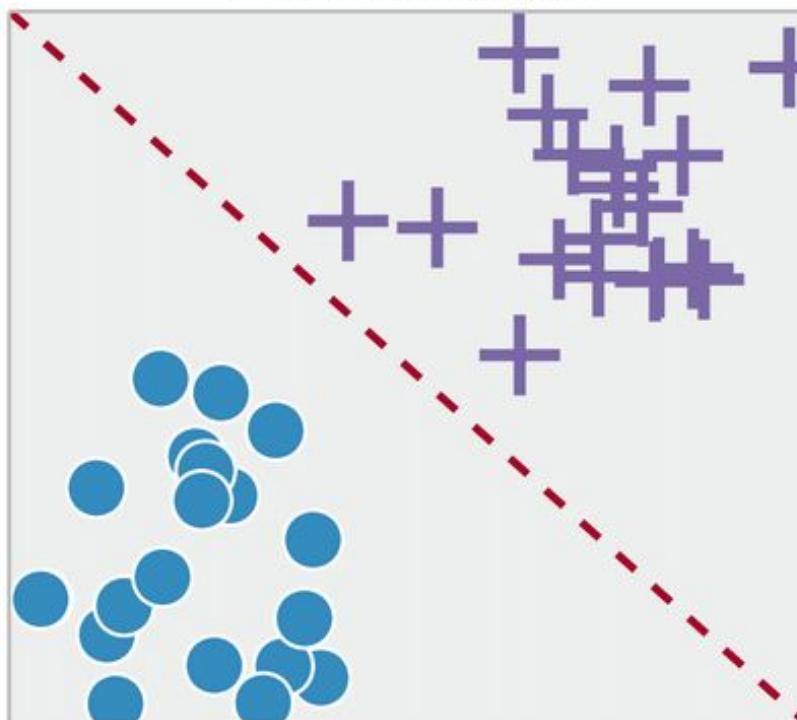
Topic 3

Neural Network for Classification

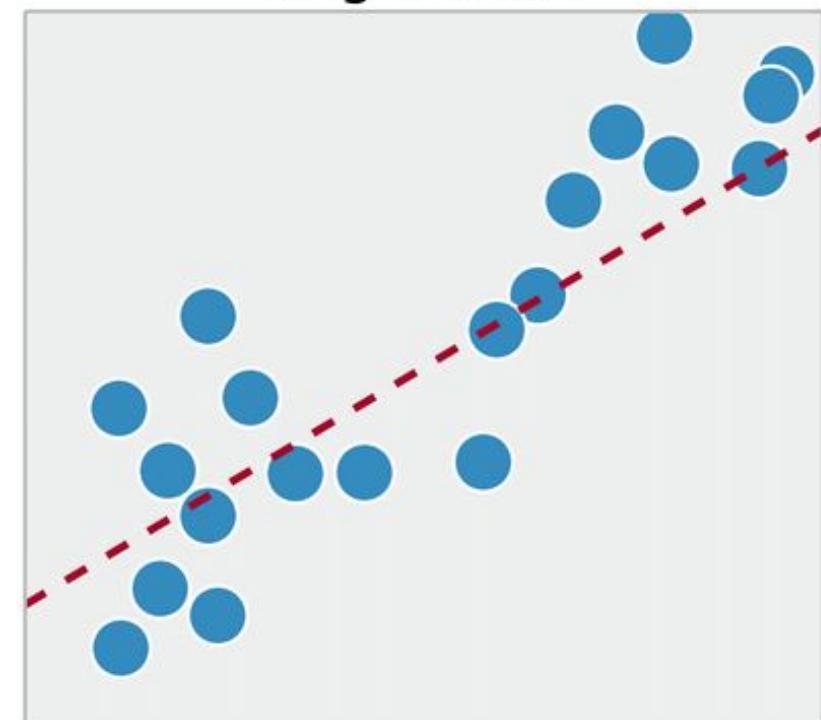
What is Classification?

- Classification is identifying to which category an object belongs to.
- Classification is a supervised learning method.
During training, labels/classes are provided.

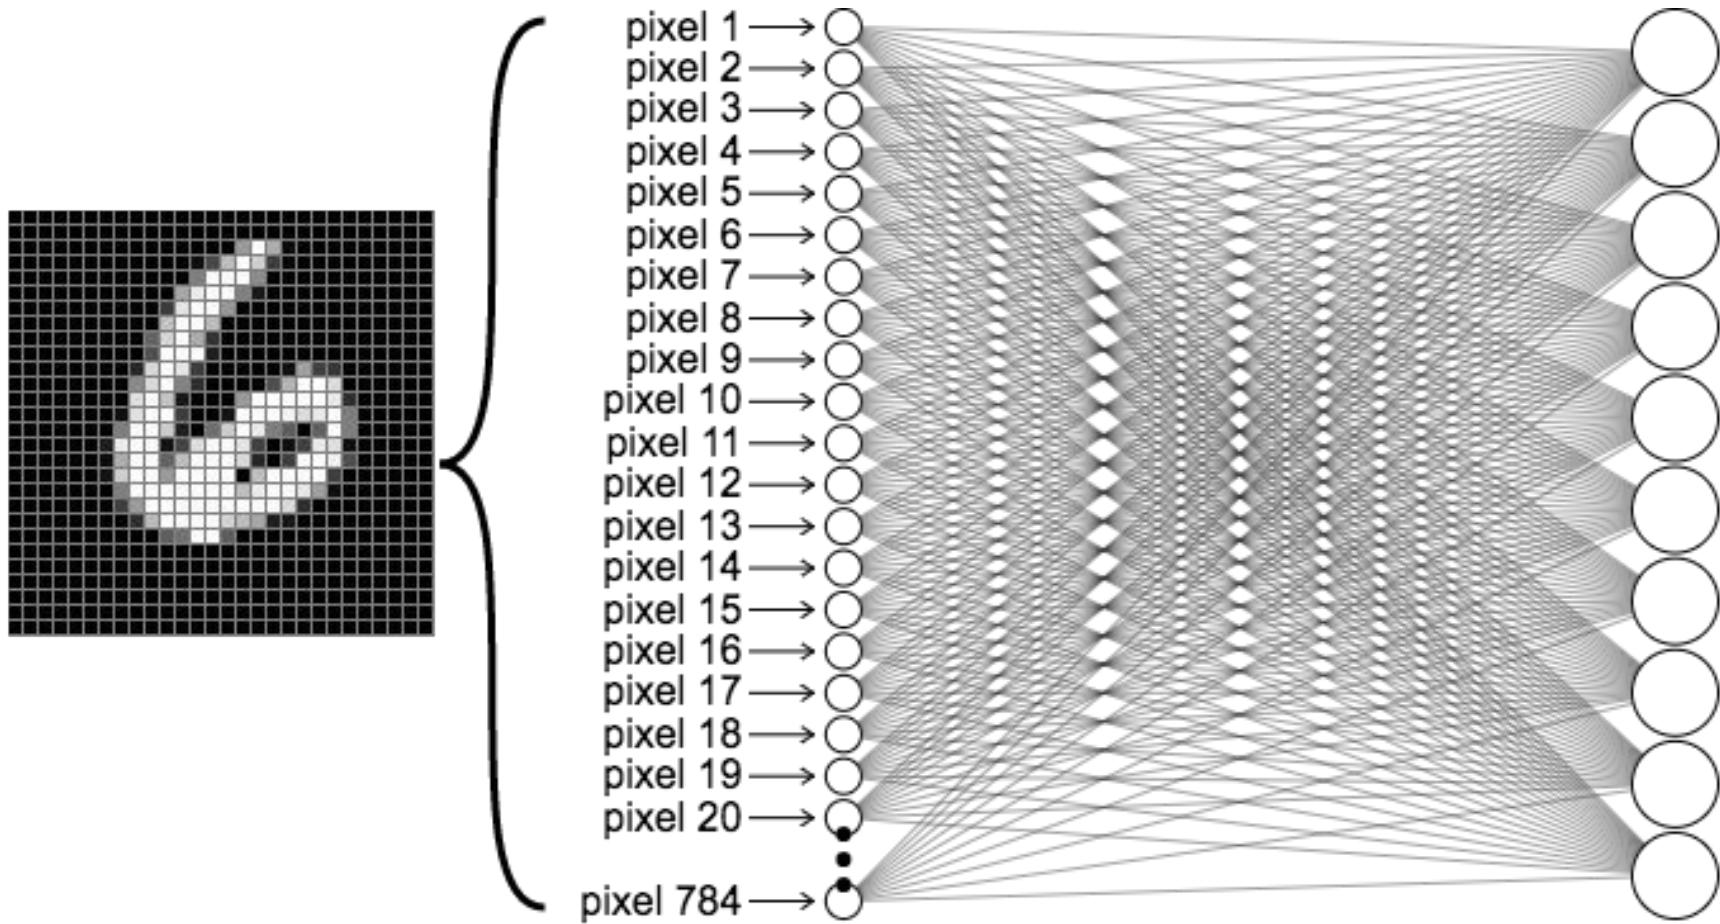
Classification



Regression

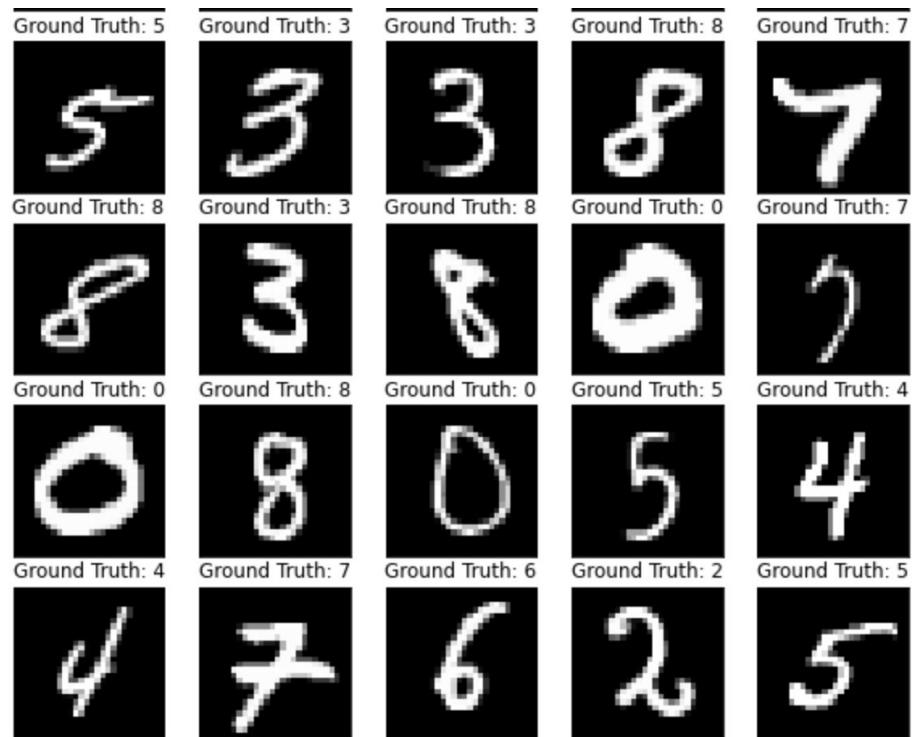


NN for HandWritten Digit Classification



MNIST Handwritten Digits

- The MNIST data is split into three parts:
 - 55,000 training data
 - 10,000 test data
 - 5,000 validation data
- Each image is 28 pixels by 28 pixels and serialize to a 1D vector of 784 numbers
- The label is one hot encoded.



Download MNIST Datasets

```
mean, std = (0.5,), (0.5,
```

```
# Create a transform and normalise data
# Download MNIST test dataset and load test data
transform = transforms.Compose([transforms.ToTensor(),
                               transforms.Normalize(mean, std)
                           ])
```

```
trainset = datasets.MNIST('~/.pytorch/MNIST/', download=True, train=True,
                         transform=transform)
```

Use datasets to:
Download dataset
Transform data

```
trainloader = torch.utils.data.DataLoader(trainset, batch_size=64, shuffle=True)
```

Use DataLoader
function for
Mini-batch Training

```
# Download MNIST test dataset and load test data
testset = datasets.MNIST('~/.pytorch/MNIST/', download=True, train=False,
                        transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=64, shuffle=False)
```

One Hot Encoding for Digits

0	[1 0 0 0 0 0 0 0 0]
1	[0 1 0 0 0 0 0 0 0]
2	[0 0 1 0 0 0 0 0 0]
3	[0 0 0 1 0 0 0 0 0]
4	[0 0 0 0 1 0 0 0 0]
5	[0 0 0 0 0 1 0 0 0]
6	[0 0 0 0 0 0 1 0 0]
7	[0 0 0 0 0 0 0 1 0]
8	[0 0 0 0 0 0 0 0 1]
9	[0 0 0 0 0 0 0 0 1]

Probabilistic Output

- Classification neural network model always gives a probabilistic output, for example



Probability

0.8



0.15

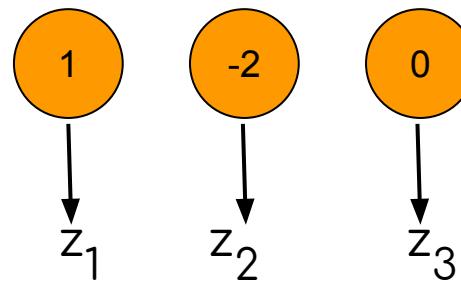


0.05

Softmax

- Softmax extends logistics regression (produce 0 to 1 for a single output) into a multi-class problem.
- Softmax assigns probabilities to each class in a multi-class problem using the formula below

$$p_j = \sigma(z)_j = \frac{e^{z_j}}{\sum_{i=1}^k e^{z_i}}$$



$$[1, -2, 0] \rightarrow e^1, e^{-2}, e^0 = [2.71, 0.14, 1] \rightarrow 0.7, 0.04, 0.26$$

Define the NN Model

L1 = 32

L2 = 64

```
class Model(nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.fc1 = nn.Linear(784, L1)
        self.fc2 = nn.Linear(L1, L2)
        self.fc3 = nn.Linear(L2, 10)
```

```
    def forward(self, x):
        x = x.view(x.shape[0], -1)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

```
model = Model()
```

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model.to(device)
```

Alternative NN Model Definition

```
import torch
import torch.nn as nn
import torch.nn.functional as F

L1 = 32
L2 = 64
model = nn.Sequential(nn.Flatten(),
                      nn.Linear(784, L1),
                      nn.ReLU(),
                      nn.Linear(L1, L2),
                      nn.ReLU(),
                      nn.Linear(L2, 10))
```

Add Model to GPU

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')  
model.to(device)
```

Loss Function and Optimizer

```
# Loss Function
```

```
criterion = nn.CrossEntropyLoss()
```

```
# Optimizer
```

```
optimizer = torch.optim.Adam(model.parameters(),  
lr=0.01)
```

Cross Entropy

Cross entropy is used to measure the closeness of two probability distribution

$$H_{y'}(y) = - \sum_i y'_i \log(y_i)$$

Example of Cross Entropy

Cross entropy = $1 * \log(0.6) + 0 * \log(0.3) + 0 * \log(0.1) = 0.51$

Compute Cross Entropy Example

Softmax Prediction	Case 1	Case 2
0.73	0	1
0.02	0	0
0.25	1	0

$$\text{Loss for Case 1} = -\log(0.25) = 1.46$$

$$\text{Loss for Case 2} = -\log(0.73) = 0.31$$

Activity: Cross Entropy

Compute the Cross Entropy for the following 2 cases

Correct Label (y)	Predicted Softmax (yhat) case 1	Predicted Softmax (yhat) case 2
0	0.8	0.1
1	0.1	0.8
0	0.1	0.1

Ans: Case 1 = 1, Case 2 = 0.1

Cross Entropy Loss

`torch.nn.CrossEntropyLoss()`

- Combines LogSoftMax and NLLLoss
- useful when training a classification problem with n classes
- The input is expected to contain scores for each class
- input has to be a 2D Tensor of size (minibatch,n).

Train the Model

```
for i in range(num_epochs):
    train_loss = 0

    for batch, (X, y) in enumerate(trainloader):
        X = X.to(device)
        y = y.to(device)

        yhat = model(X)

        optimizer.zero_grad()
        loss = criterion(yhat, y)
        loss.backward()
        optimizer.step()

        train_loss += loss.item()

    train_tracker.append(train_loss/len(trainloader))
    print(f"Epoch{[i+1]}/{[num_epochs]} | Training loss: {[train_loss/len(trainloader)]} | ",end="")

Epoch(1/10) | Training loss: 0.493915855884552 | Test loss: 0.36642441153526306 | Accuracy : 0.8849999904632568
Epoch(2/10) | Training loss: 0.3544181287288666 | Test loss: 0.27184629440307617 | Accuracy : 0.9181999564170837
Epoch(3/10) | Training loss: 0.3391174077987671 | Test loss: 0.38928401470184326 | Accuracy : 0.8816999793052673
Epoch(4/10) | Training loss: 0.3254077732563019 | Test loss: 0.2915478050708771 | Accuracy : 0.9121999740600586
Epoch(5/10) | Training loss: 0.31116983294487 | Test loss: 0.29876407980918884 | Accuracy : 0.911300003528595
Epoch(6/10) | Training loss: 0.32121512293815613 | Test loss: 0.3007839322090149 | Accuracy : 0.9145999550819397
Epoch(7/10) | Training loss: 0.3109281361103058 | Test loss: 0.2833910584449768 | Accuracy : 0.9146999716758728
Epoch(8/10) | Training loss: 0.3096688389778137 | Test loss: 0.33763203024864197 | Accuracy : 0.8985999822616577
Epoch(9/10) | Training loss: 0.3109353184700012 | Test loss: 0.39147722721099854 | Accuracy : 0.8889999985694885
Epoch(10/10) | Training loss: 0.31191393733024597 | Test loss: 0.31013721227645874 | Accuracy : 0.910099983215332

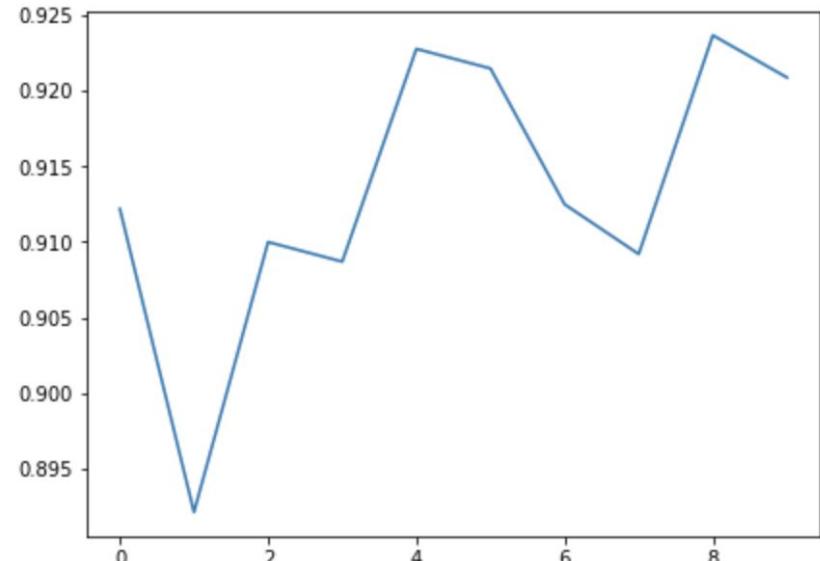
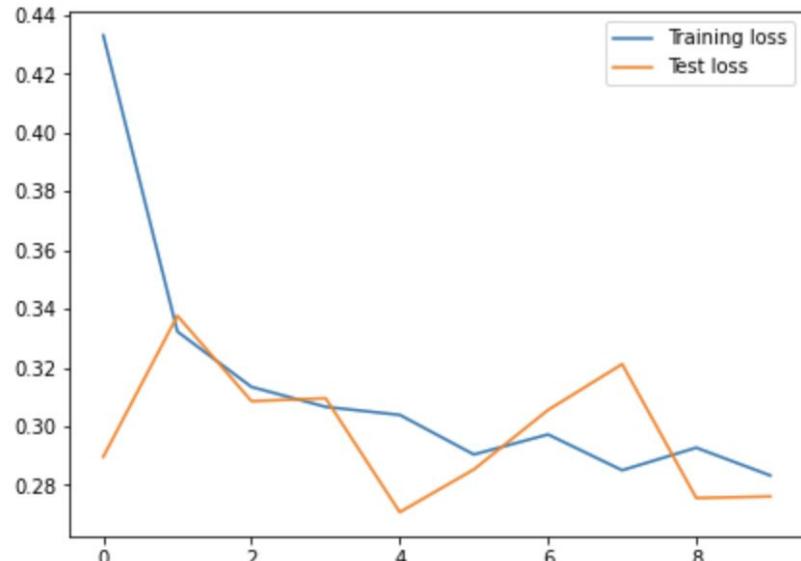
Number correct : 9101, Total : 10000
Accuracy of the model after 30 epochs on the 10000 test images: 91.00999450683594%
```

Evaluate the Model

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.plot(train_tracker, label='Training loss')
plt.plot(test_tracker, label='Test loss')
plt.legend()
```

```
plt.subplot(1,2,2)
plt.plot(accuracy_tracker, label='Test accuracy')
plt.show()
```



Fashion MNIST

- Similar to Handwritten Digits, the fashion MNIST data is split into three parts:
 - 55,000 training data
 - 10,000 test data
 - 5,000 validation data
- Each image is 28 pixels by 28 pixels and serialize to a 1D vector of 784 numbers
- The label is one hot encoded.



Fashion MNIST Classes

T-shirt/top	[1 0 0 0 0 0 0 0 0 0]
Trouser	[0 1 0 0 0 0 0 0 0 0]
Pullover	[0 0 1 0 0 0 0 0 0 0]
Dress	[0 0 0 1 0 0 0 0 0 0]
Coat	[0 0 0 0 1 0 0 0 0 0]
Sandal	[0 0 0 0 0 1 0 0 0 0]
Shirt	[0 0 0 0 0 0 1 0 0 0]
Sneaker	[0 0 0 0 0 0 0 1 0 0]
Bag	[0 0 0 0 0 0 0 0 1 0]
Ankle boot	[0 0 0 0 0 0 0 0 0 1]

Activity: Fashion MNIST Classification

- Import the Fashion MNIST data
- Create a NN model is with 3 hidden layers (128,64,32) to classify the fashion
- Perform all the steps as shown in the demo
 - Prepare the data
 - Build the model
 - Train the model
 - Evaluate the model

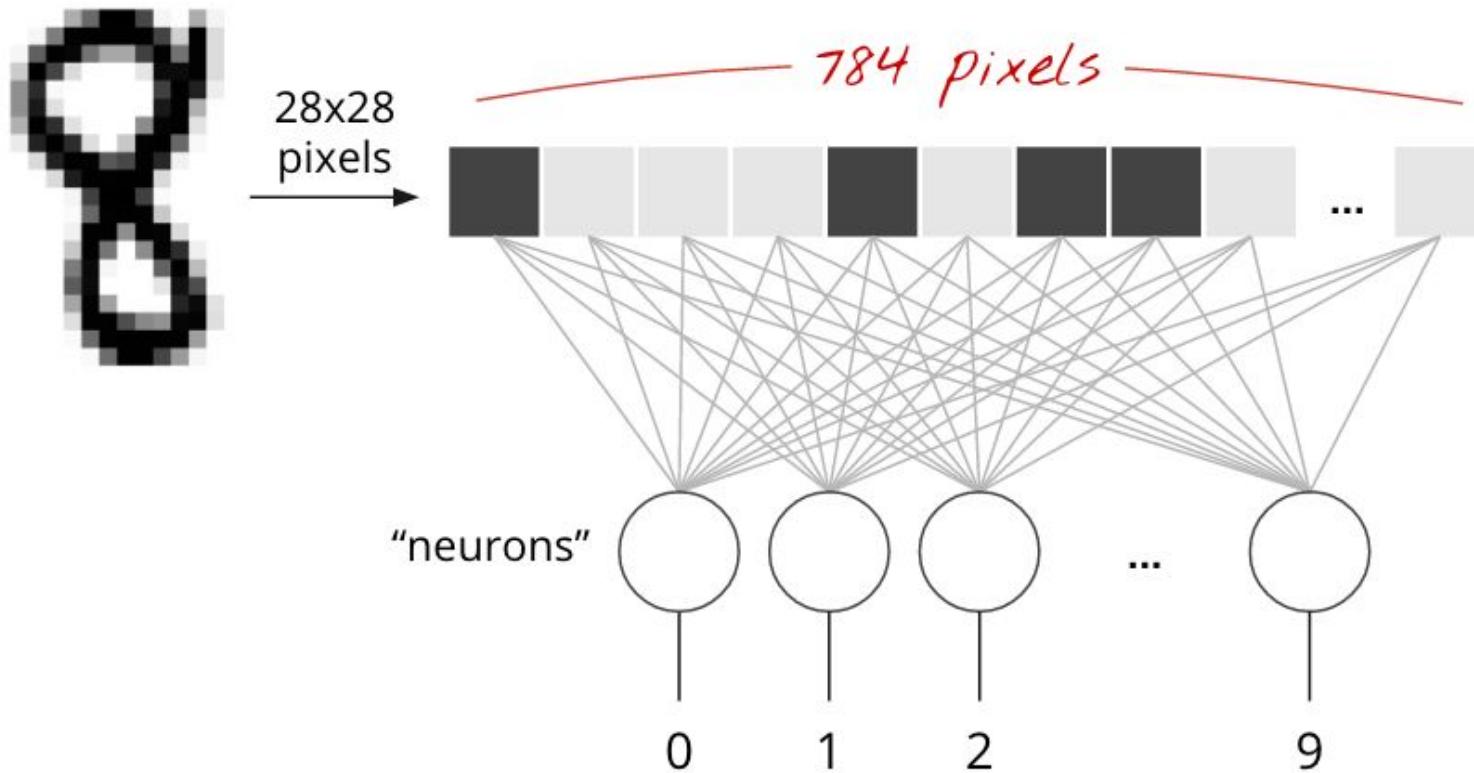
Time: 30 mins

Topic 4

Convolutional Neural Network (CNN)

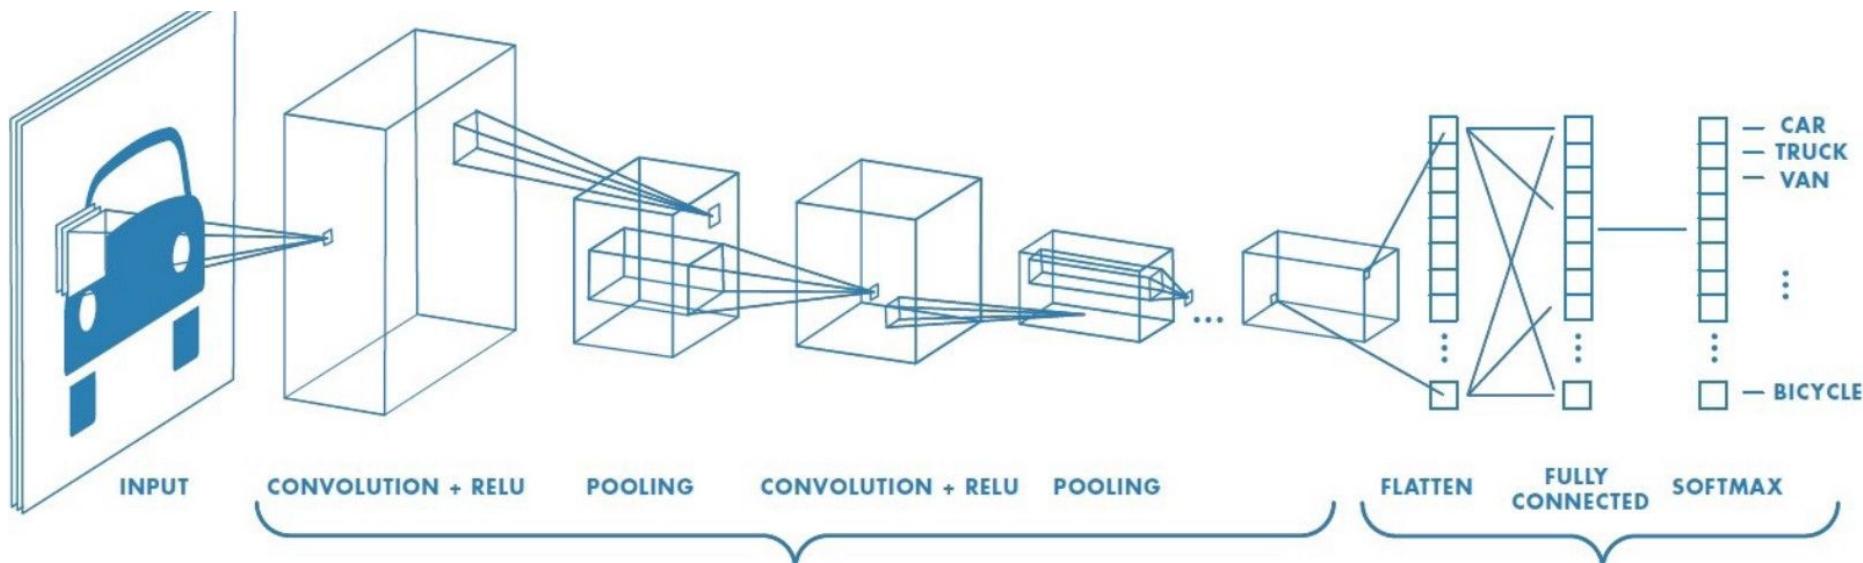
Limitation of DNN

When we vectorize a 28×28 mnist image to a 784×1 vector, we lose the spatial information of the image, and thus reduce the accuracy of the prediction



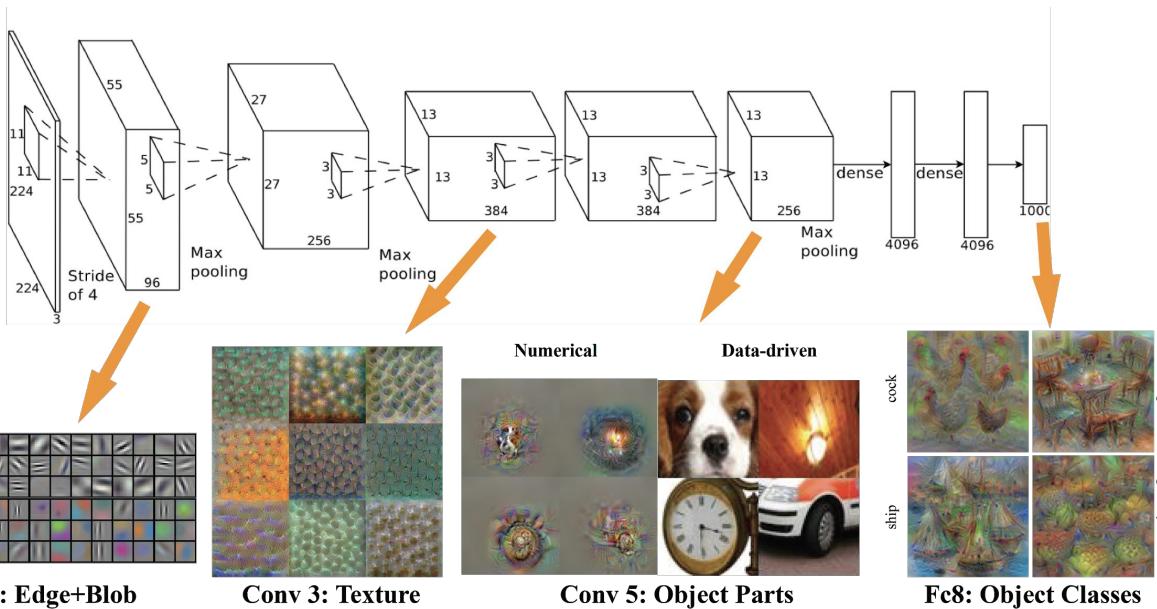
Convolutional Neural Network (CNN)

- The key technology behind computer vision is convolutional neural network (CNN)
- Convolution networks are generally composed of Convolution layers, pooling layers and a last component which is the fully connected that will be used for an appropriate task like classification or detection.



Why CNN?

- CNN preserves the spatial relationship between pixels by learning localized image features using small filters (kernels).
- The first conv layer extracts general features like edges, and subsequent conv layers extract more detailed features



AlexNet
8 layers:
5 conv layers +
3 FC layers

Brief History of CNN

- 1959 Hubel & Wiesel found two types of cells in the visual primary cortex, and proposed a cascading model of these two types of cells for use in pattern recognition task
- 1980 - Inspired by Hubel and Wiesel work, Kunihiko Fukushima introduced neocognition neural network, similar to CNN concept (not using BP yet)
- 1982-1985 Geoff Hinton, Yann LeCun and others successfully used BP on NN with sigmoid activation function.
- 1990 - Yann LeCun invented CNN with BP for handwritten digits recognition
- 2012 - Alex Krizhevsky won the ImageNet LSVRC competition with AlexNet (deep CNN +ReLU+Data Augmentation+Dropout+GPU)

Inventor of CNN

- CNN was invented by Yann LeCun, currently the Director of AI Research, Facebook (<http://yann.lecun.com/>)
- LeNet-5 is the latest convolutional network designed for handwritten and machine-printed character recognition (<http://yann.lecun.com/exdb/lenet/index.html>)



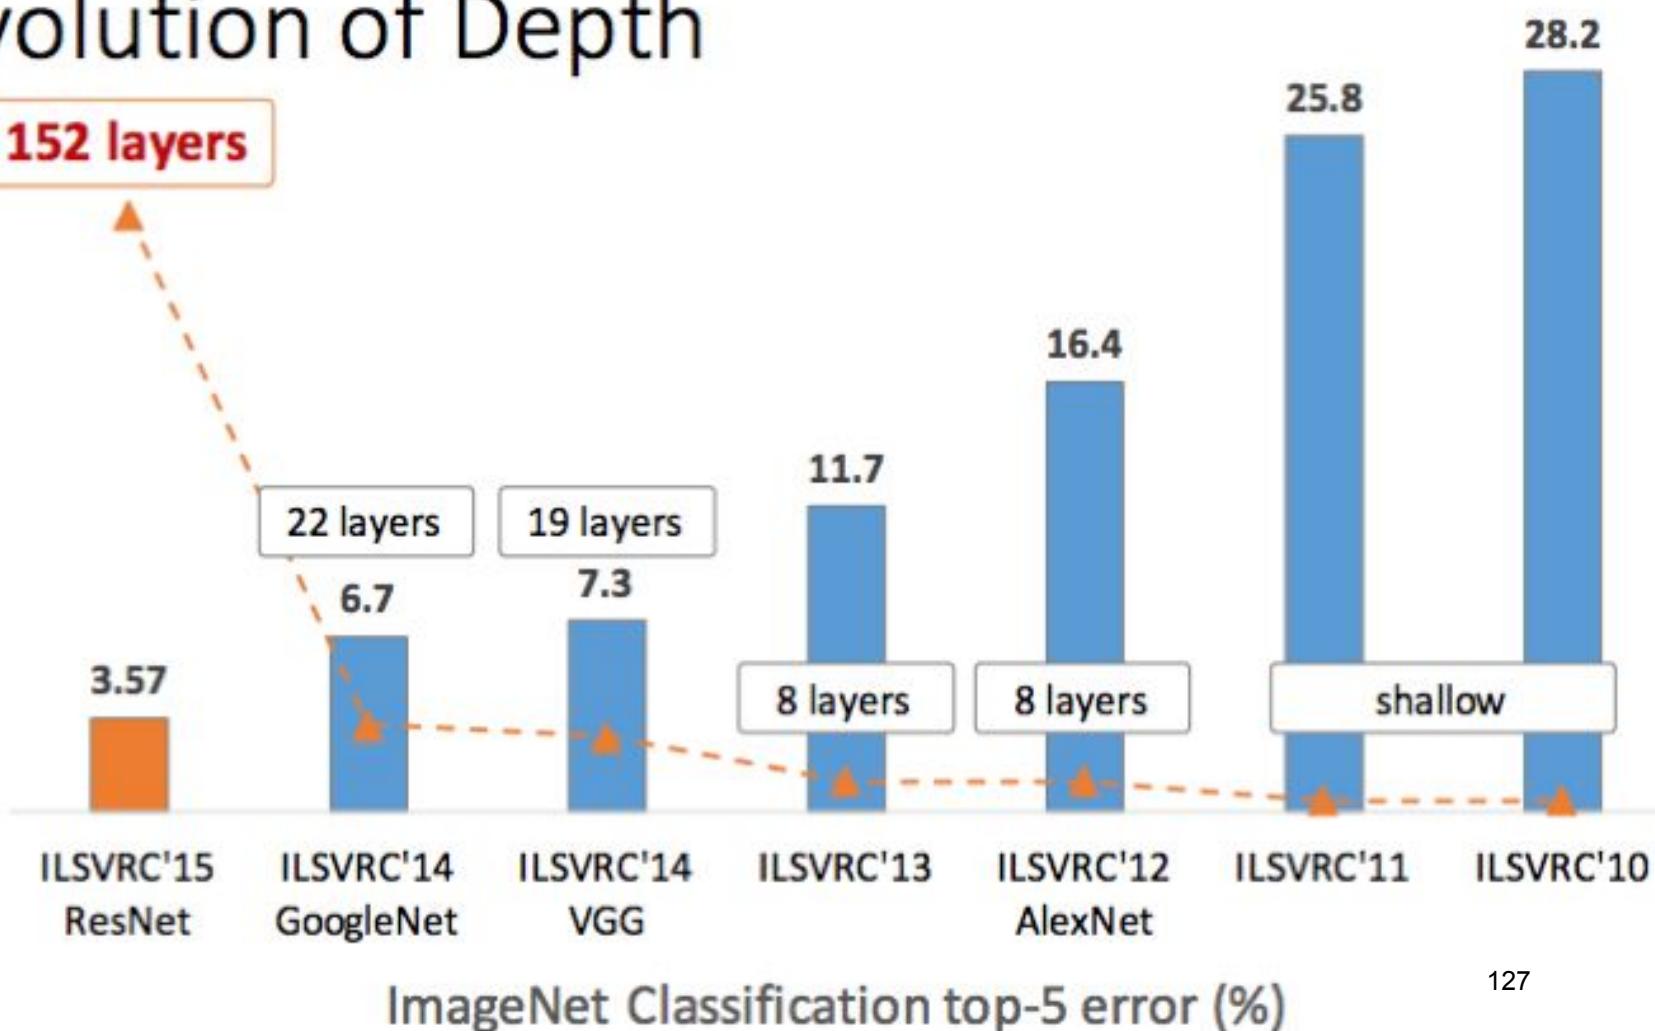
ImageNET ILSVRC

- ImageNet (image-net.org) was initiated by Li Fei Fei in 2007 with more than 15 Million images
- Held ILSVRC (International Large Scale Visual Recognition Challenge) since 2010



ILSVRC Winners

Revolution of Depth



Heroes of Machine Learning

There are many heroes in Machine Learning. The Turing Award 2019 is award to

- Geoff Hinton
- Yann Lecun
- Yoshua Bengio



*backpropagation,
boltzmann machines*



Geoff Hinton
Google

convolution



Yann Lecun
Facebook

*stacked auto-
encoders*



Yoshua Bengio
U. of Montreal

GPU utilization



Andrew Ng
Baidu

dropout



Alex Krizhevsky
Google

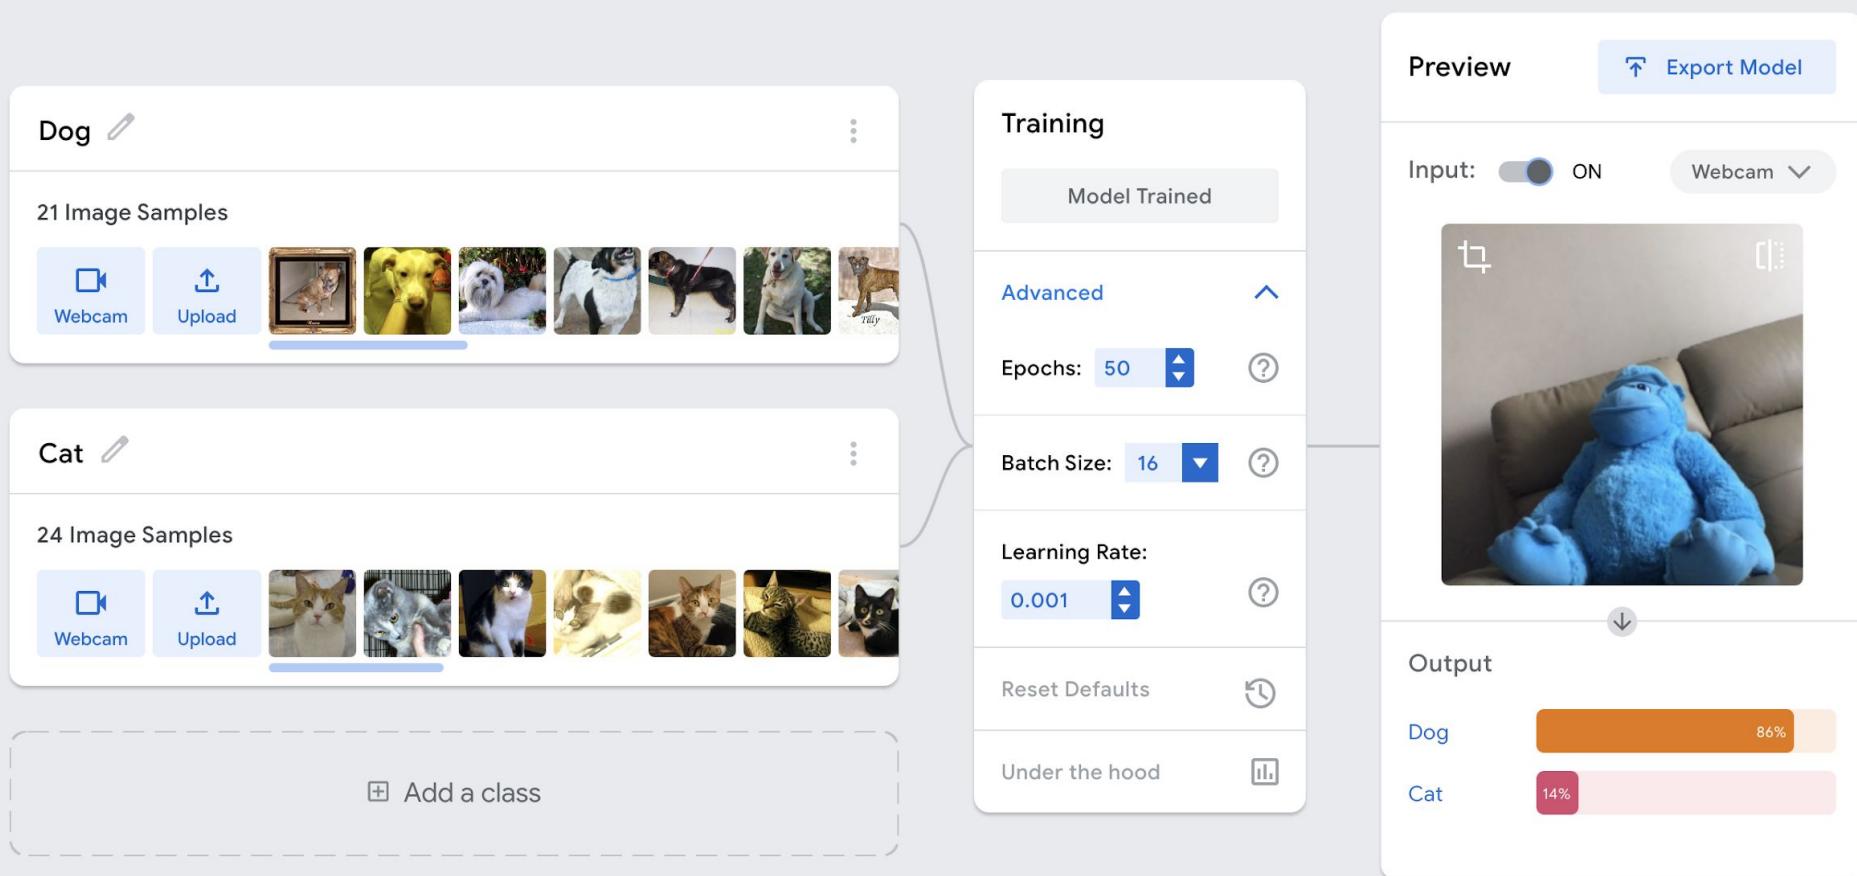
GAN



Ian Goodfellow
OpenAI

Activity : Image Classification

- Create an image classification model and training
<https://teachablemachine.withgoogle.com/train/image>
- You can use the cat and dog datasets here
- https://storage.googleapis.com/mledu-datasets/cats_and_dogs_filtered.zip



The screenshot shows the Teachable Machine interface for image classification. On the left, there are two sections: 'Dog' and 'Cat'. The 'Dog' section has 21 image samples, and the 'Cat' section has 24 image samples. Both sections include 'Webcam' and 'Upload' buttons. A large curved arrow points from the 'Dog' and 'Cat' sections towards the central 'Training' and 'Output' panels. The 'Training' panel on the right shows settings for a trained model: Epochs: 50, Batch Size: 16, and Learning Rate: 0.001. The 'Output' panel shows the classification results for a blue stuffed animal: Dog (86%) and Cat (14%).

Dog

21 Image Samples

Webcam Upload

Cat

24 Image Samples

Webcam Upload

Add a class

Training

Model Trained

Advanced

Epochs: 50

Batch Size: 16

Learning Rate: 0.001

Reset Defaults

Under the hood

Preview Export Model

Input: ON Webcam

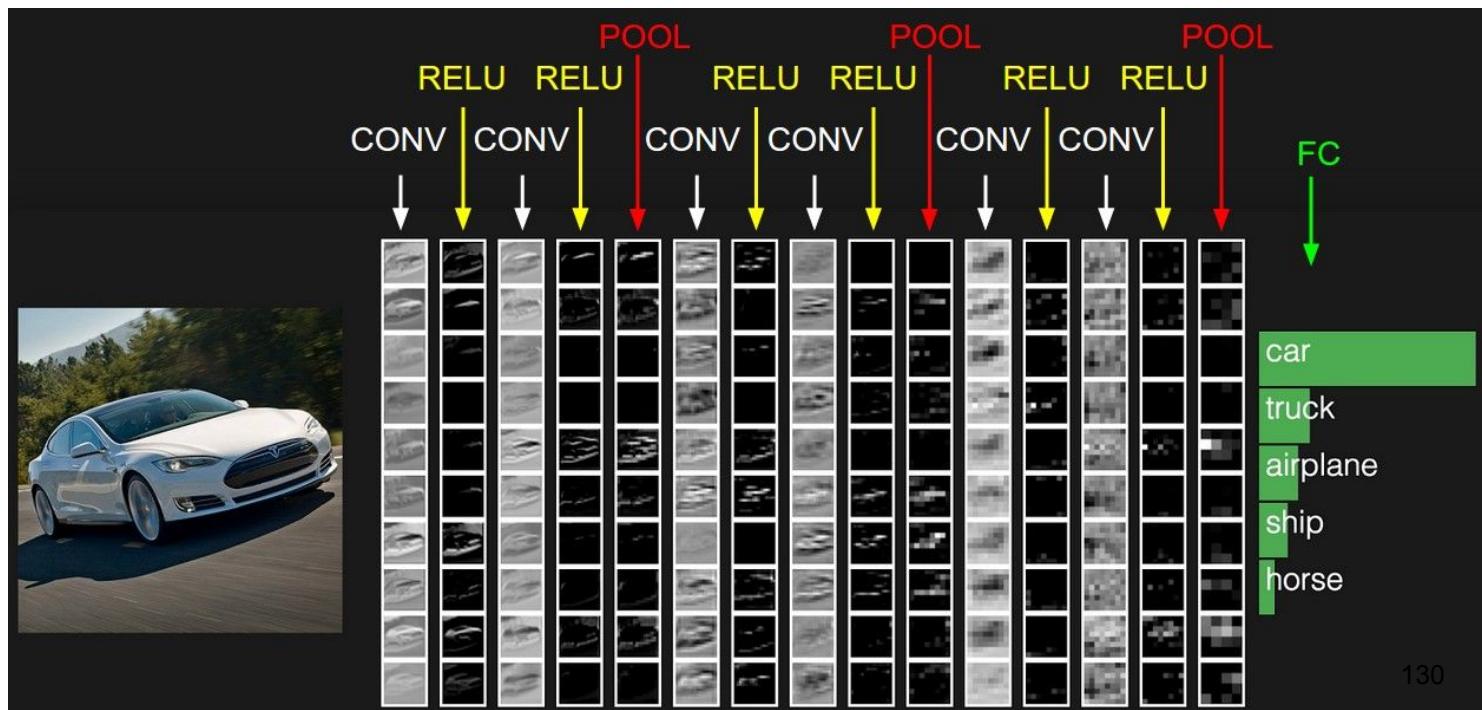
Dog 86%

Cat 14%

ConvNET Architecture

ConvNET architecture consists the follow layers

- Convolutional layer
- Activation layer
- Max Pooling layer
- Dropout layer
- Fully Connected layer



Convolution

- Convolution is the process of sliding a filter across the image
- The filter act as a feature detector

1 <small>$\times 1$</small>	1 <small>$\times 0$</small>	1 <small>$\times 1$</small>	0	0
0 <small>$\times 0$</small>	1 <small>$\times 1$</small>	1 <small>$\times 0$</small>	1	0
0 <small>$\times 1$</small>	0 <small>$\times 0$</small>	1 <small>$\times 1$</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Output Image Size

The output image size can be computed by this formula

$$(N-F)/\text{stride} + 1$$

N: input size

F: filter size

S: stride

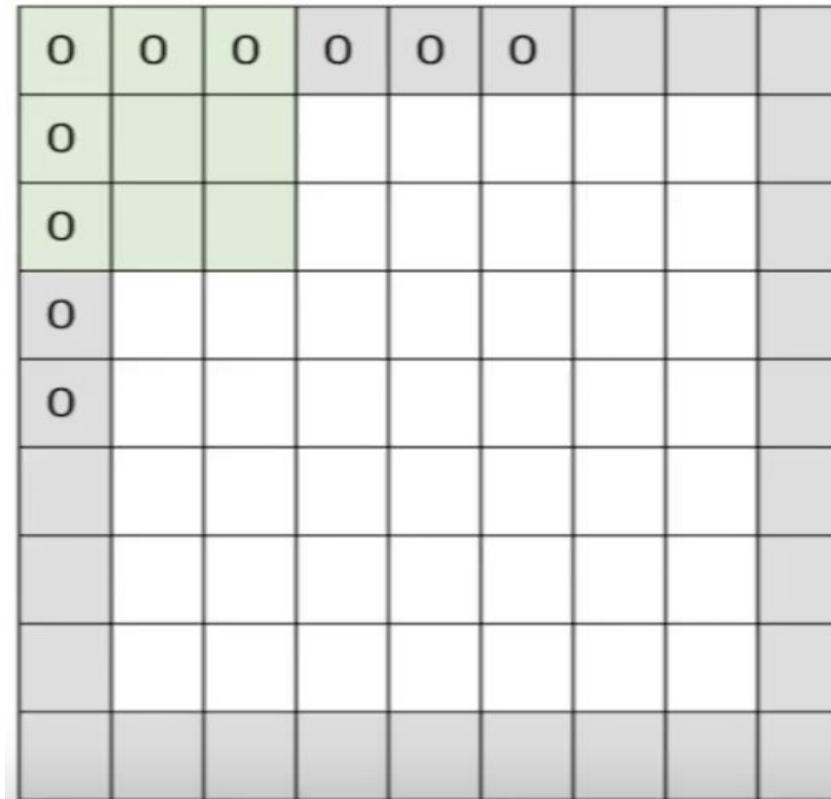
Eg $(5-3)/1+1 = 3$

$(5-3)/2+1 = 2$

Stride is the step taken when sliding the filter

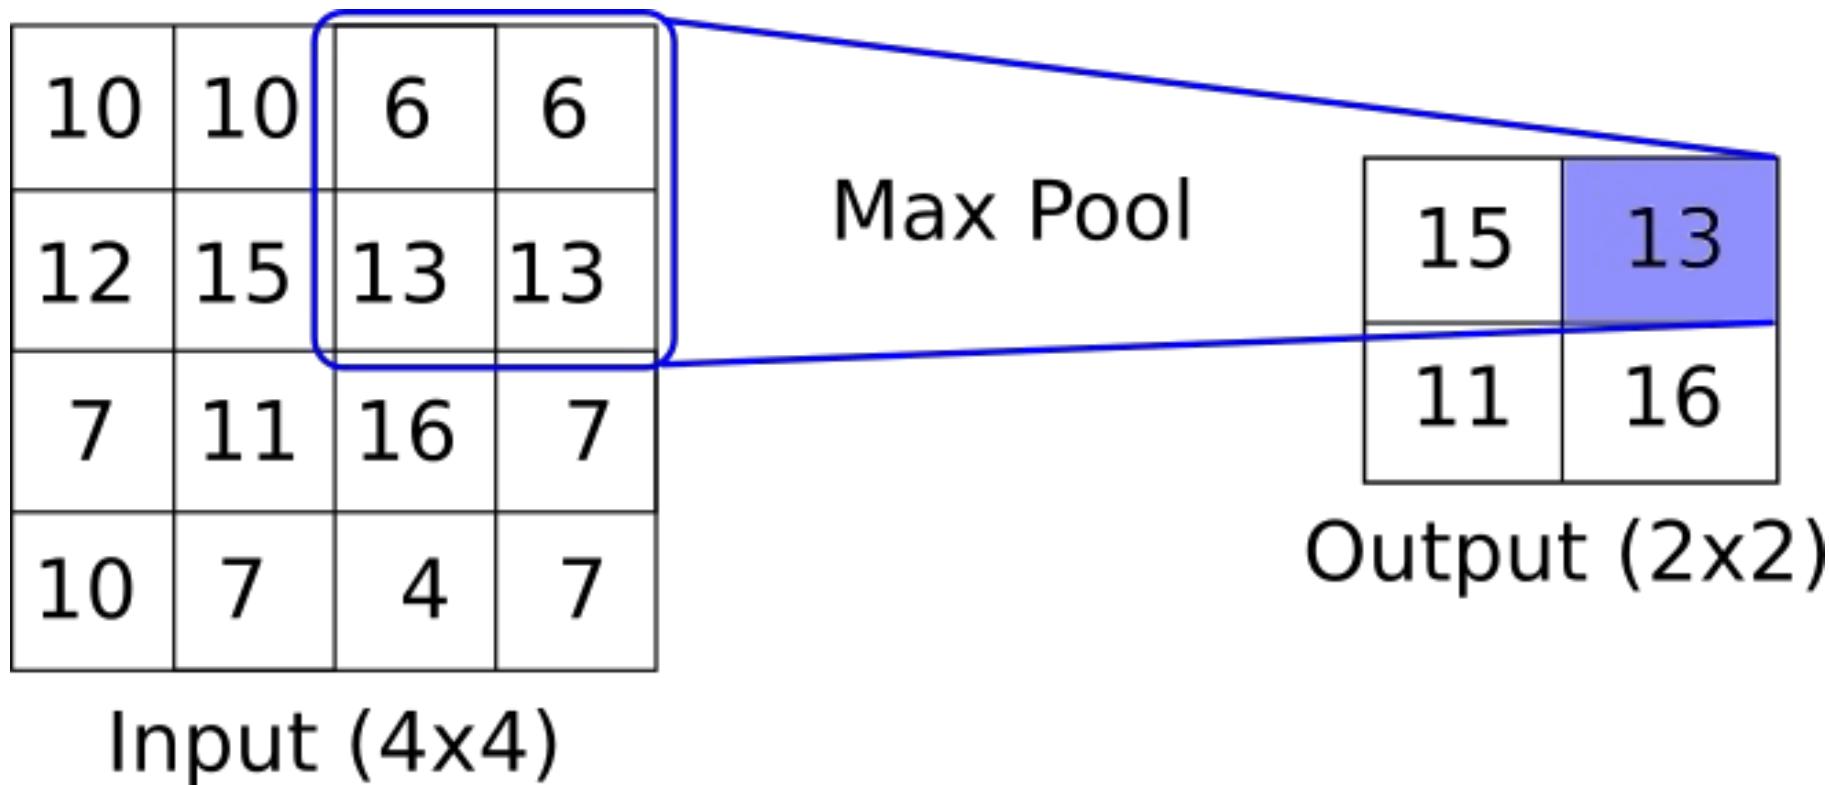
Padding

- Zero padding is used to preserve the output image size and reduce edge information loss.
- There are two types of padding
 - Valid padding - no padding
 - Same padding - keep the same output image size



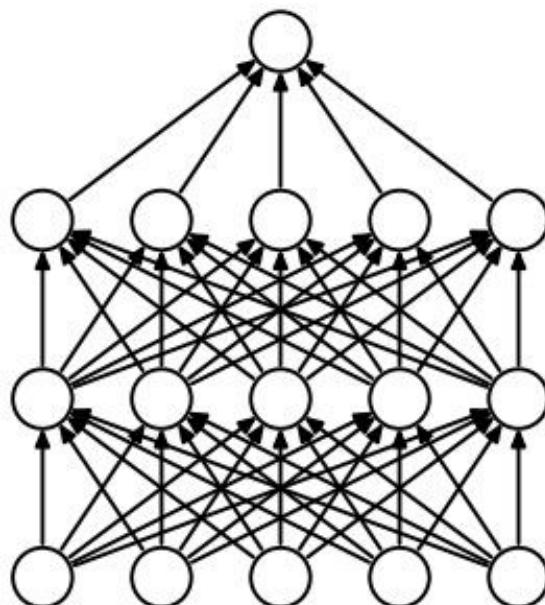
Pooling

- Similar to stride, the purpose of a pooling layer is to downsample the output size. and extract the salient features.
- Maximum or average pooling layers

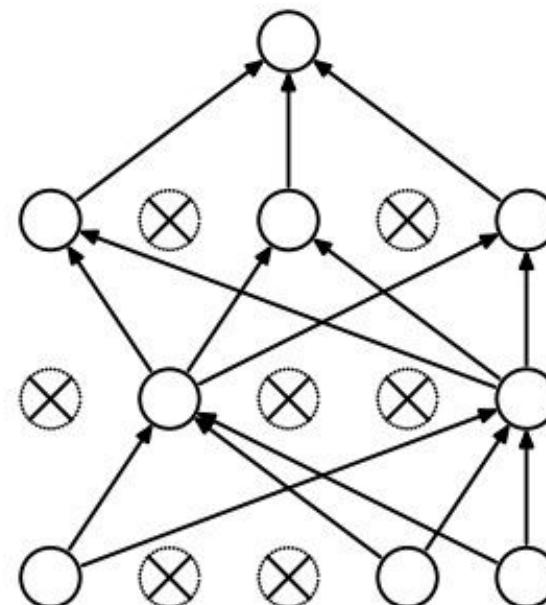


Dropout

- Dropout is a way regularize the parameters, thus reduce over-fitting and general fit better.
- During training, we randomly select activation functions and make the output equal to zero.
- It is generally apply to the fully connected layer.

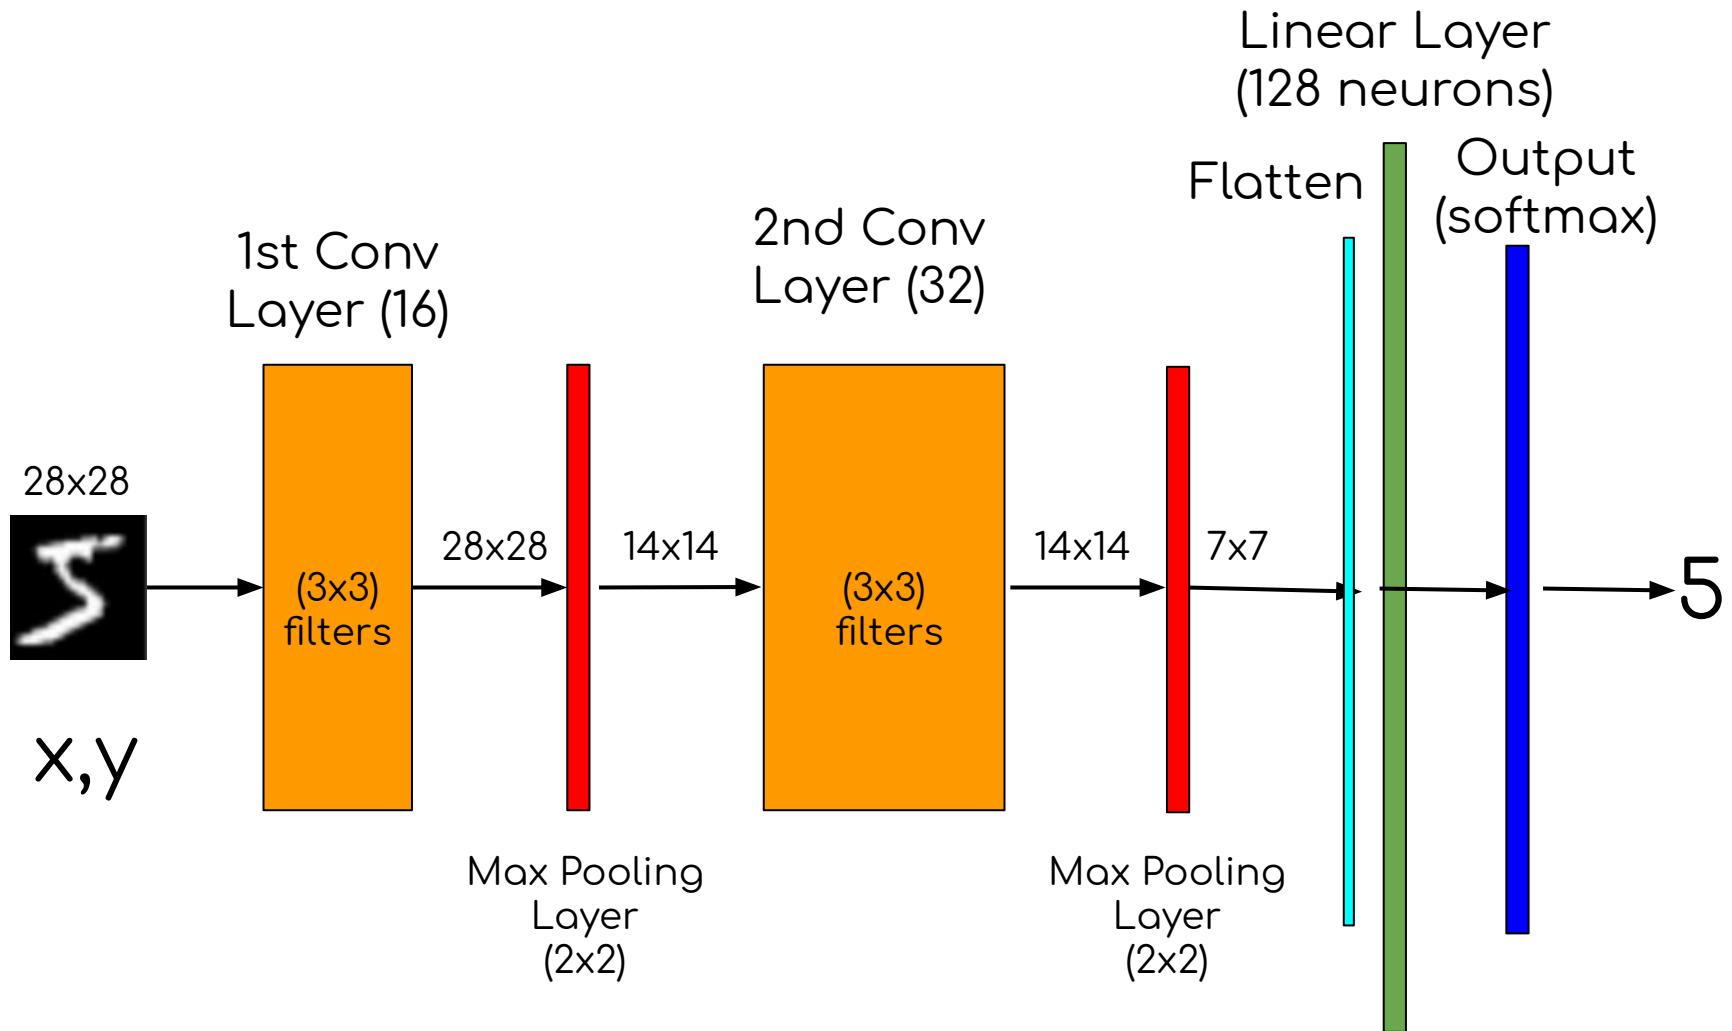


(a) Standard Neural Net



(b) After applying dropout.

CNN Demo for MNIST



Define the CNN Model

L1 = 16

L2 = 32

L3 = 128

```
class Model(nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.conv1 = nn.Conv2d(1, L1, 3, 1, 1)
        self.conv2 = nn.Conv2d(L1, L2, 3, 1, 1)
        self.fc1 = nn.Linear(L2 * 7 * 7, L3)
        self.fc2 = nn.Linear(L3, 10)

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = torch.flatten(x, 1)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x

model = Model()
```

Alternative CNN Model Definition

```
import torch  
import torch.nn as nn  
import torch.nn.functional as F
```

```
L1 = 16  
L2 = 32  
L3 = 128
```

```
model = nn.Sequential(nn.Conv2d(1,L1,3,1,1),  
                      nn.ReLU(),  
                      nn.MaxPool2d(2),  
                      nn.Conv2d(L1,L2,3,1,1),  
                      nn.ReLU(),  
                      nn.MaxPool2d(2),  
                      nn.Flatten(),  
                      nn.Linear(L2*7*7,L3),  
                      nn.Linear(L3, 10))
```

Add Model to GPU

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')  
model.to(device)
```

Train the Model

```
for i in range(num_epochs):
    train_loss = 0

    for batch, (X, y) in enumerate(trainloader):
        X = X.to(device)
        y = y.to(device)

        yhat = model(X)

        optimizer.zero_grad()
        loss = criterion(yhat, y)
        loss.backward()
        optimizer.step()

        train_loss += loss.item()

train_tracker.append(train_loss/len(trainloader))
print(f'Epoch{[i+1]}/{[num_epochs]} | Training loss: {[train_loss/len(trainloader)]} | ',end=)
```

```
Epoch(1/10) | Training loss: 0.1815527715781795 | Test loss: 0.07517748720911302 | Accuracy : 0.9764999747276306
Epoch(2/10) | Training loss: 0.08686092886530153 | Test loss: 0.06301685600368392 | Accuracy : 0.9801999926567078
Epoch(3/10) | Training loss: 0.07878129867931653 | Test loss: 0.07113557542655892 | Accuracy : 0.9785999655723572
Epoch(4/10) | Training loss: 0.07240879847404066 | Test loss: 0.07097163051218507 | Accuracy : 0.9775999784469604
Epoch(5/10) | Training loss: 0.06776289350617164 | Test loss: 0.06310828684380197 | Accuracy : 0.9818999767303467
Epoch(6/10) | Training loss: 0.06799703942527295 | Test loss: 0.0705306206628807 | Accuracy : 0.9801999926567078
Epoch(7/10) | Training loss: 0.06663249518403581 | Test loss: 0.08165361432873594 | Accuracy : 0.9777999520301819
Epoch(8/10) | Training loss: 0.06236315957801568 | Test loss: 0.06675530408520106 | Accuracy : 0.983299970626831
Epoch(9/10) | Training loss: 0.06283533013741194 | Test loss: 0.0810895188391511 | Accuracy : 0.9786999821662903
Epoch(10/10) | Training loss: 0.06212432895201757 | Test loss: 0.07680490346007038 | Accuracy : 0.979699969291687
```

Number correct : 9797, Total : 10000

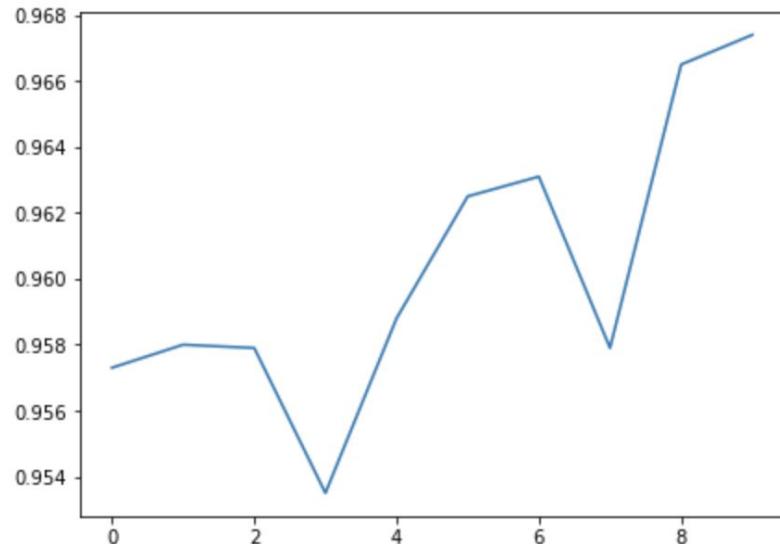
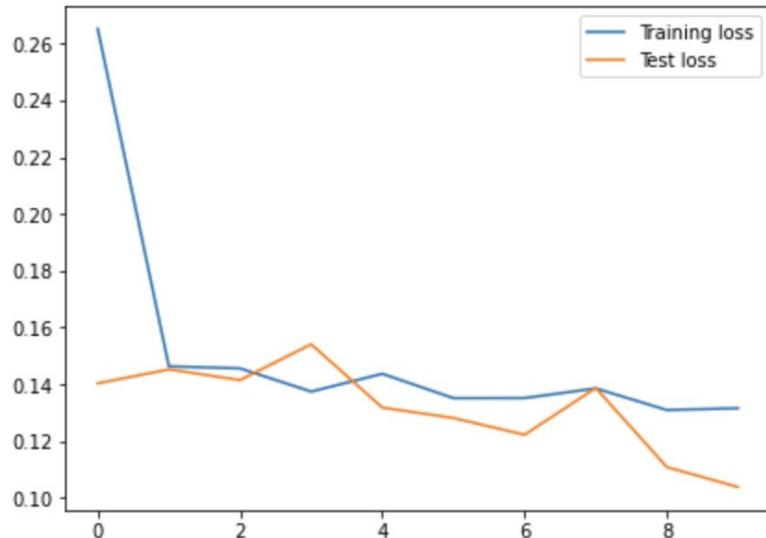
Accuracy of the model after 30 epochs on the 10000 test images: 97.97000122070312%

Evaluate the Model

```
import matplotlib.pyplot as plt
```

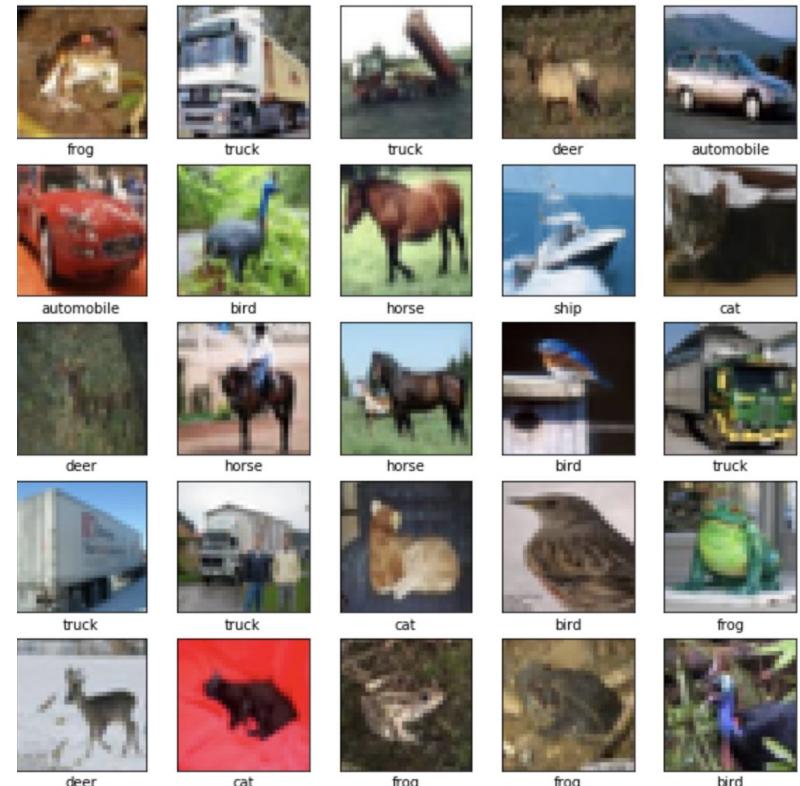
```
plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.plot(train_tracker, label='Training loss')
plt.plot(test_tracker, label='Test loss')
plt.legend()
```

```
plt.subplot(1,2,2)
plt.plot(accuracy_tracker, label='Test accuracy')
plt.show()
```



CIFAR-10 Image Dataset

- The CIFAR-10 dataset (<https://www.cs.toronto.edu/~kriz/cifar.html>)
- 60000 32x32 colour images
- 10 classes, with 6000 images per class.
- 50000 training images and 10000 test images

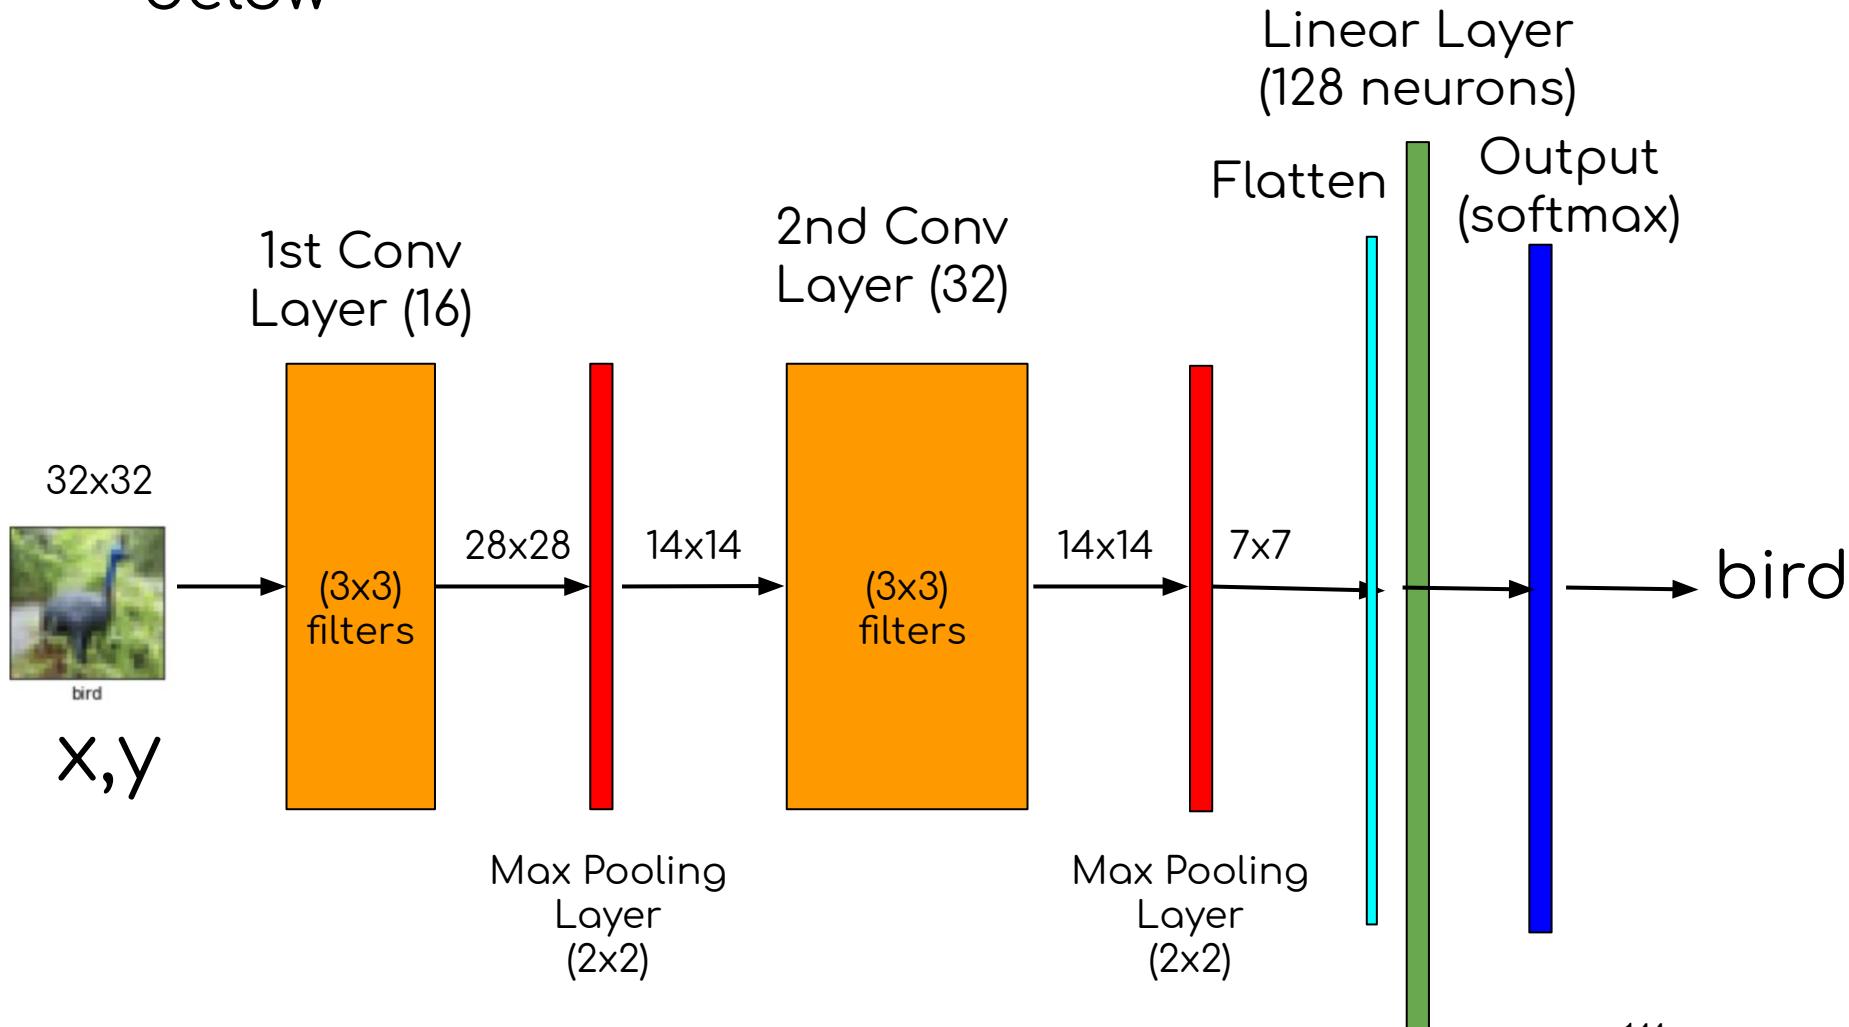


CIFAR10 Classes

airplane	[1 0 0 0 0 0 0 0 0 0]
automobile	[0 1 0 0 0 0 0 0 0 0]
bird	[0 0 1 0 0 0 0 0 0 0]
cat	[0 0 0 1 0 0 0 0 0 0]
deer	[0 0 0 0 1 0 0 0 0 0]
dog	[0 0 0 0 0 1 0 0 0 0]
frog	[0 0 0 0 0 0 1 0 0 0]
horse	[0 0 0 0 0 0 0 1 0 0]
ship	[0 0 0 0 0 0 0 0 1 0]
truck	[0 0 0 0 0 0 0 0 0 1]

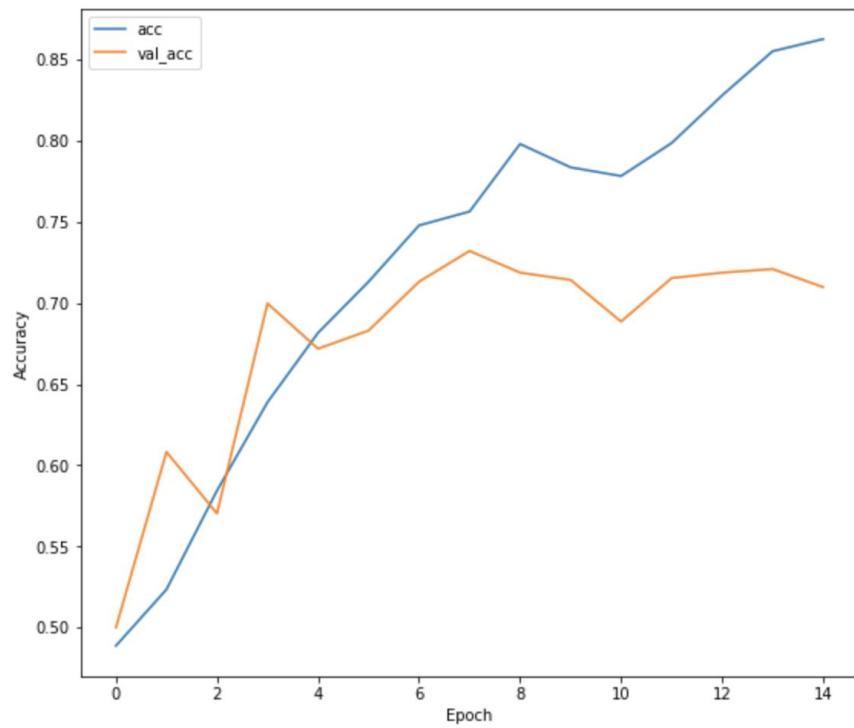
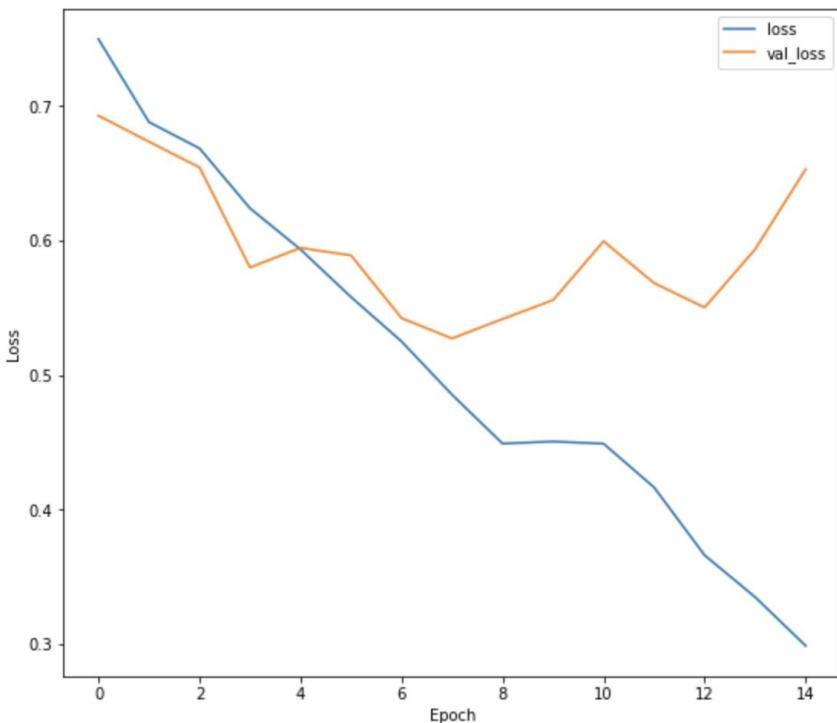
Activity: CNN on CIFAR10 Dataset

- Create a CNN model for CIFAR10 dataset as shown below



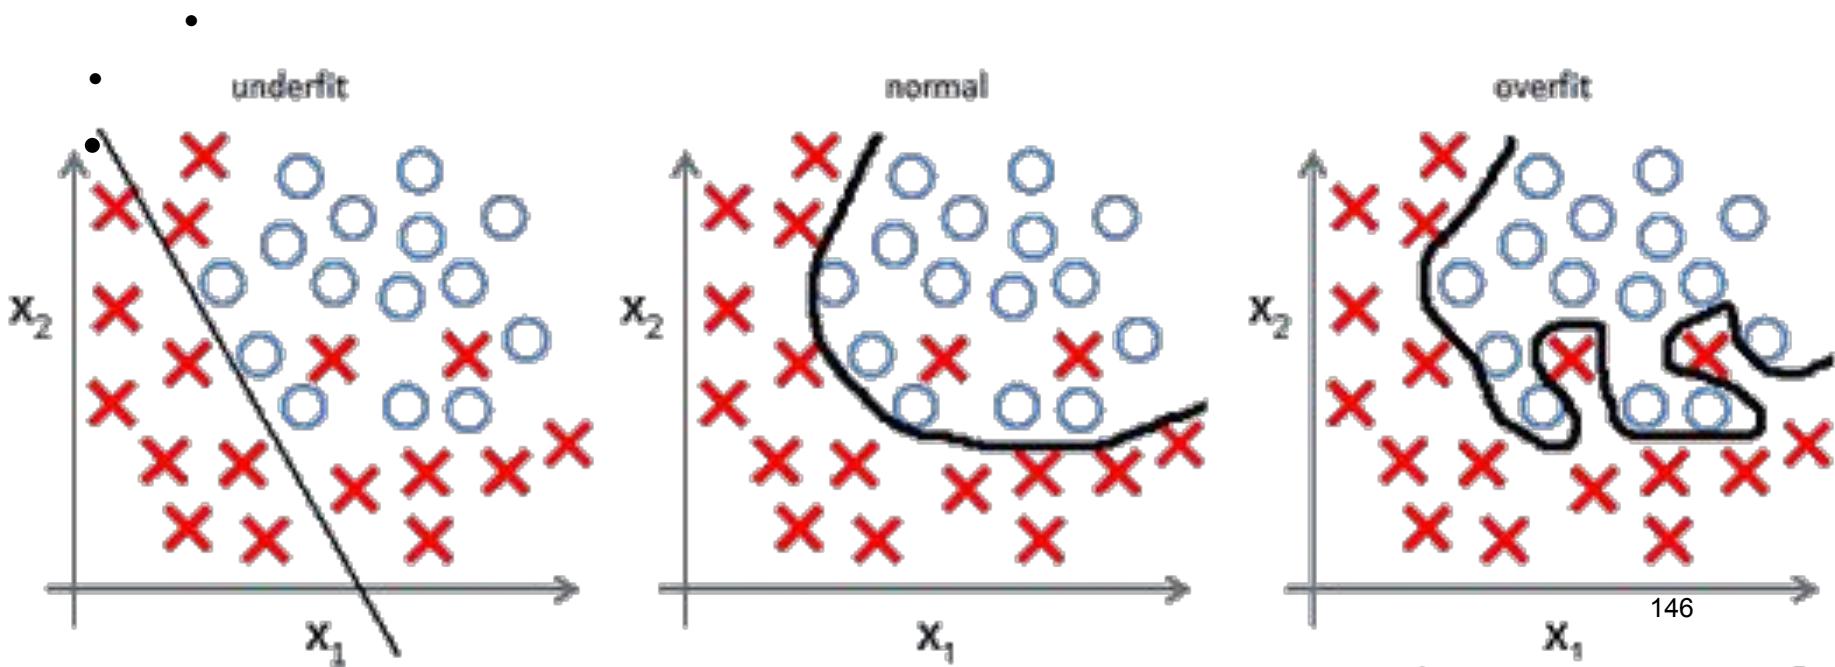
Issues with Small Dataset

For small dataset, a general issue is overfitting. While the training loss reduce with training epoch, the testing loss has U-turned and increased with training epoch.



Overfitting and Underfitting

- Overfitting occurs when a statistical model or machine learning algorithm captures the noise of the data. Intuitively, overfitting occurs when the model or the algorithm fits the data too well. Specifically, overfitting occurs if the model or algorithm shows low bias but high variance.
- Underfitting occurs when a statistical model or machine learning algorithm cannot capture the underlying trend of the data. Intuitively, underfitting occurs when the model or the algorithm does not fit the data well enough. Specifically, underfitting occurs if the model or algorithm shows low variance but high bias



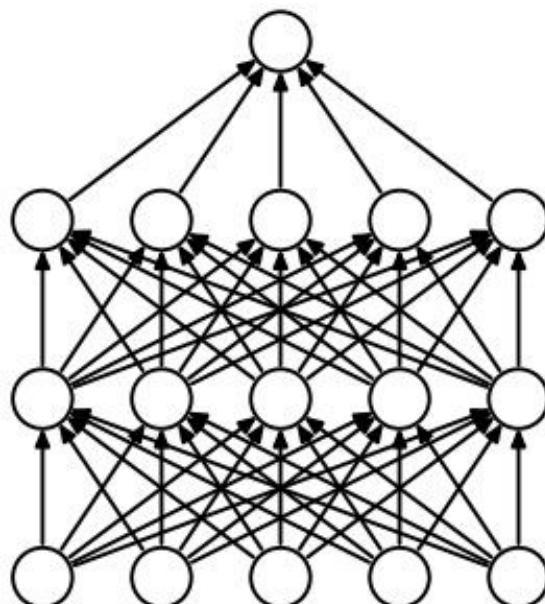
Methods to Solve Overfitting

Overfitting issues are commonly solved by the following techniques:

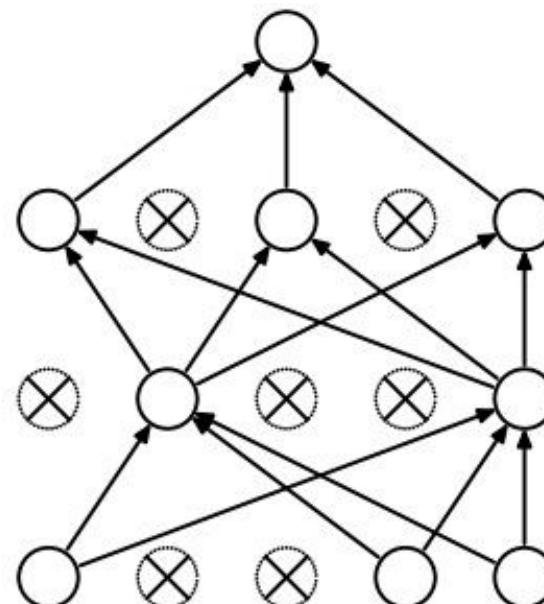
- Increase training data (**expensive**)
- Reduce model capacity (**reduce learning**)
- L1 or L2 weight regularization
- Early stopping (**reduce learning**)
- Data augmentation
- Dropout

Dropout

- Dropout is a way regularize the parameters, thus reduce over-fitting and general fit better.
- During training, we randomly select activation functions and make the output equal to zero.
- It is generally apply to the fully connected layer.



(a) Standard Neural Net



(b) After applying dropout.

Load Small Datasets

- Download the hymenoptera_data dataset from Google classroom and upload to your Google Drive or local computer.
- That data has about 120 training images each for ants and bees.
- There are 75 validation images for each class. Usually, for a very small dataset, the model is prompt to over-hitting.

```
data_dir = '/content/drive/MyDrive/dataset/hymenoptera_data'
```

```
trainset = datasets.ImageFolder(os.path.join(data_dir, 'train'),  
transforms['train'])  
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,  
shuffle=True)
```

```
testset = datasets.ImageFolder(os.path.join(data_dir, 'val'),  
transforms['val'])  
testloader = torch.utils.data.DataLoader(testset, batch_size=4,  
shuffle=True)
```

Mount Google Drive

- If you use Google Drive for storing data, you need to mount the Google Drive to access your data

```
from google.colab import drive  
drive.mount('/content/drive')
```

Resize the Images

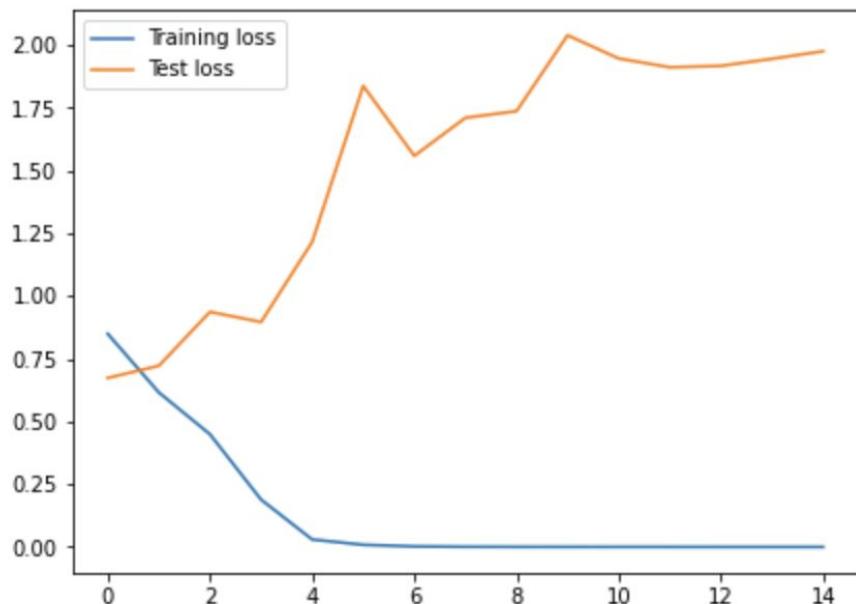
```
mean = [0.5, 0.5, 0.5]
```

```
std = [0.5, 0.5, 0.5]
```

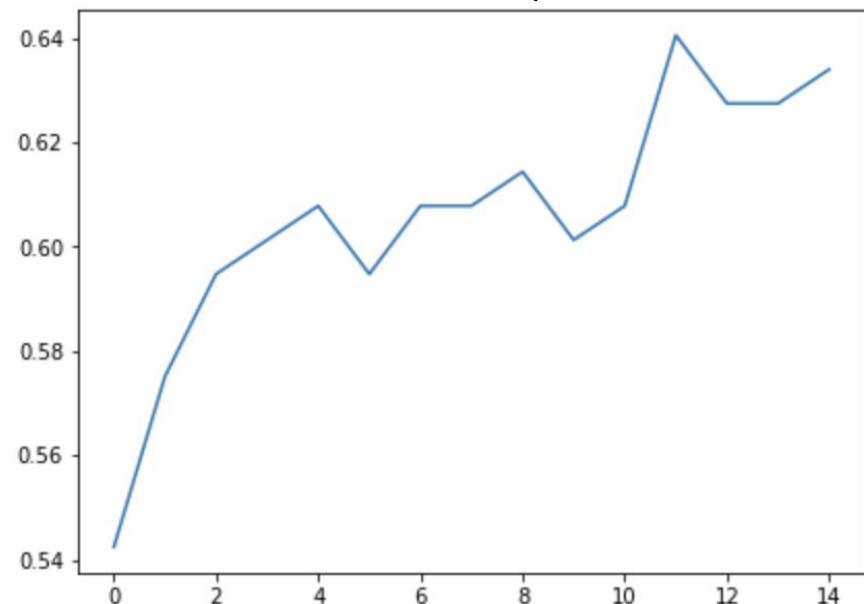
```
transforms = {
    'train': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize(mean=mean, std=std)
    ]),
    'val': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize(mean=mean, std=std)
    ]),
}
```

Result (Baase line)

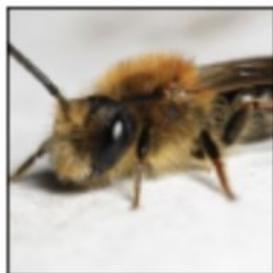
Loss



Accuracy



bee/ant



ant/bee



ant/bee



ant/ant



Apply Data Augmentation

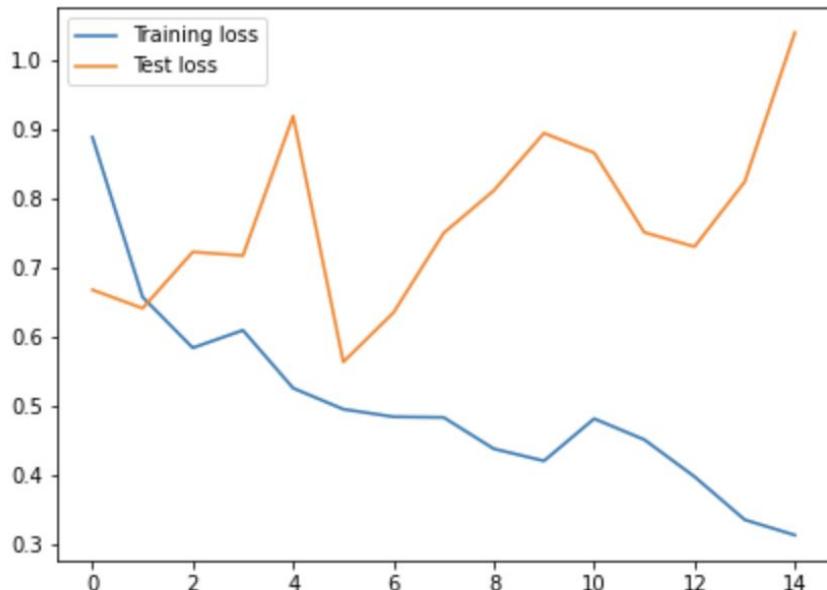
```
mean = [0.5, 0.5, 0.5]
```

```
std = [0.5, 0.5, 0.5]
```

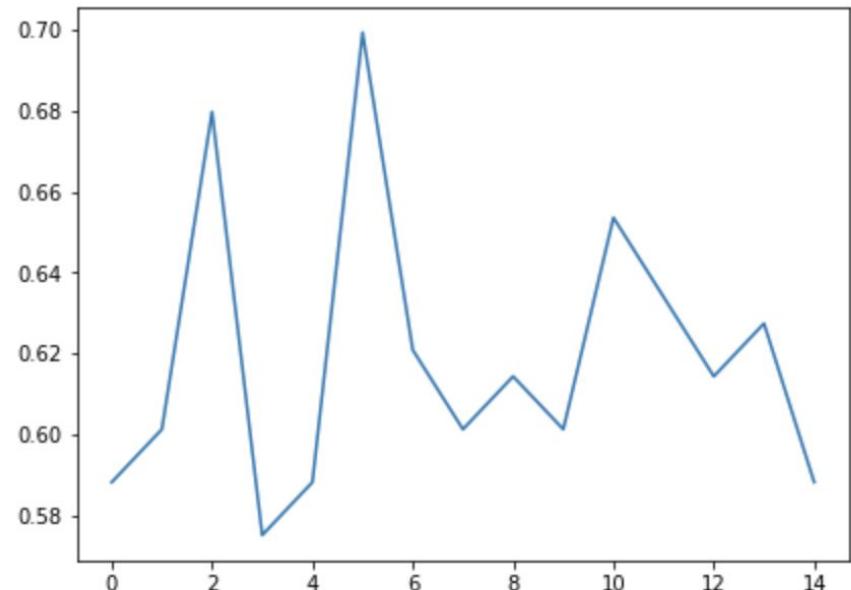
```
transforms = {
    'train': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.RandomRotation(20),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize(mean=mean, std=std)
    ]),
    'val': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize(mean=mean, std=std)
    ]),
}
```

Result (Data Augmentation)

Loss



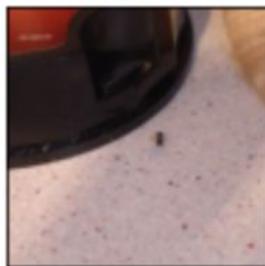
Accuracy



bee/ant



ant/ant



ant/ant



bee/ant



Apply DA and Dropout

```
class Model(nn.Module):
    def __init__(self):
        super(Model,self).__init__()
        self.conv1 = nn.Conv2d(3,L1,3,1,1)
        self.d1 = nn.Dropout(0.2)
        self.conv2 = nn.Conv2d(L1,L2,3,1,1)
        self.fc1 = nn.Linear(L2*56*56,L3)
        self.fc2 = nn.Linear(L3, 2)

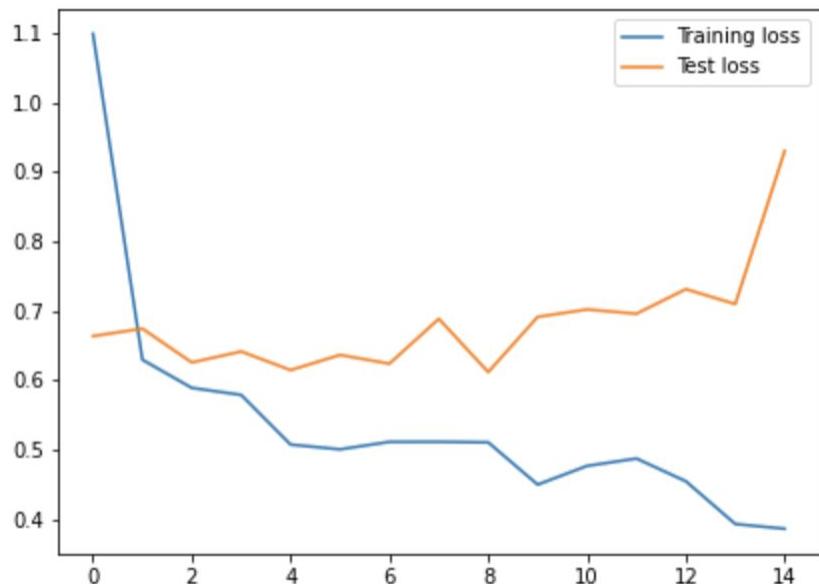
    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)
        x = self.d1(x)
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view( x.size(0),-1)

        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x

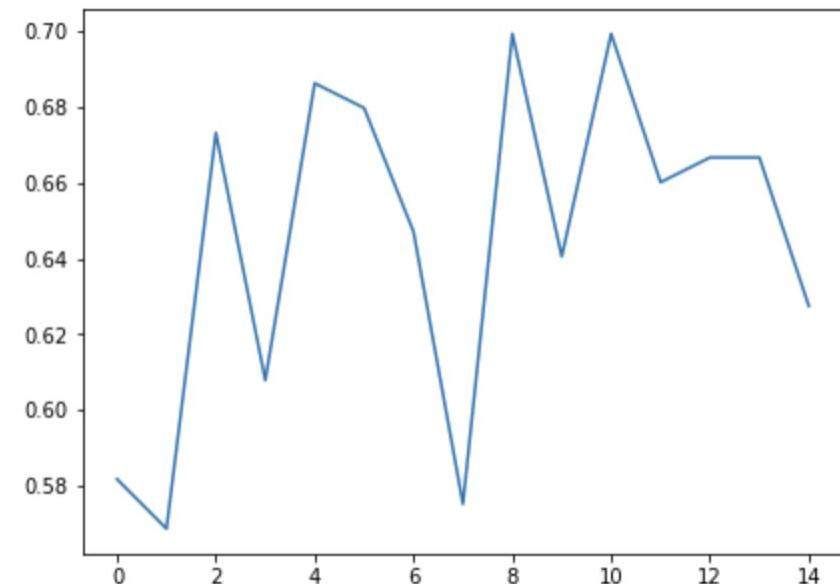
model = Model()
```

Result (DA+Dropout)

Loss



Accuracy



bee/ant



ant/ant



ant/ant



bee/ant



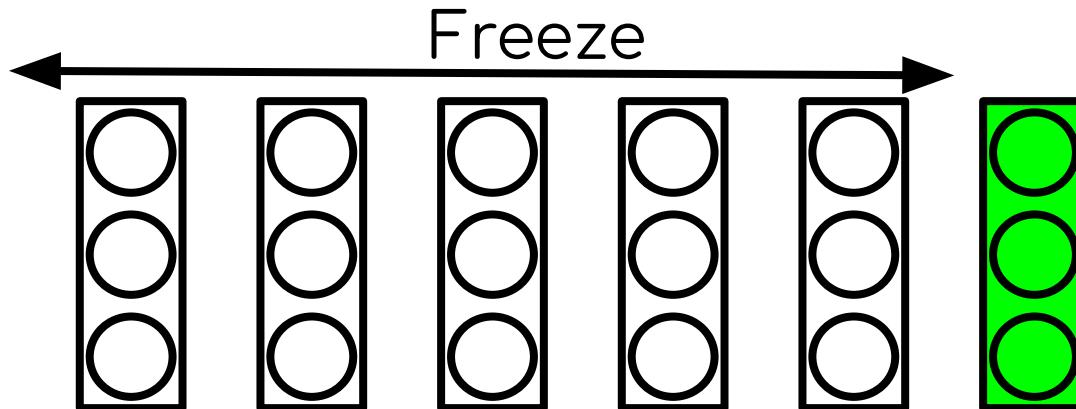
Topic 5

Transfer Learning

Transfer Learning

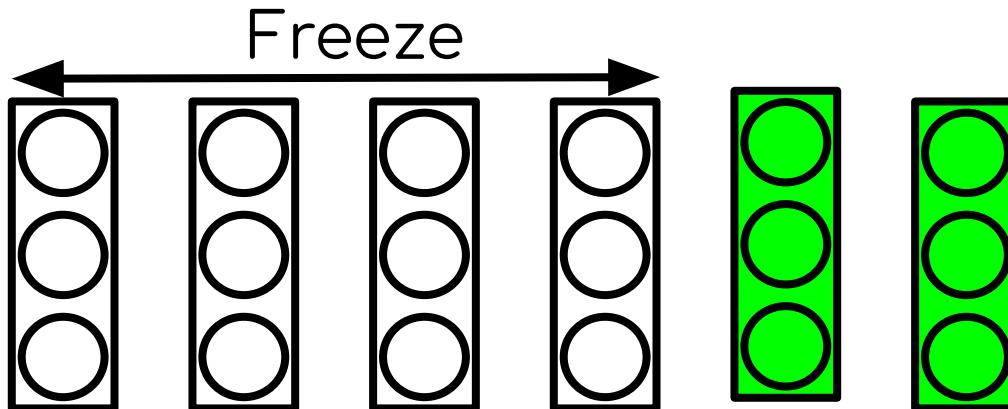
- A **pre-trained model** is a saved network that was previously trained on a large dataset, typically on a large-scale image-classification task.
- You either use the pretrained model as it is, or use transfer learning to customize this model to a given task.
- The intuition behind transfer learning is that if a model trained on a large and general enough dataset, this model will effectively serve as a generic model of the visual world.
- You can then take advantage of these learned feature maps without having to start from scratch training a large model on a large dataset.

Transfer Learning Approaches



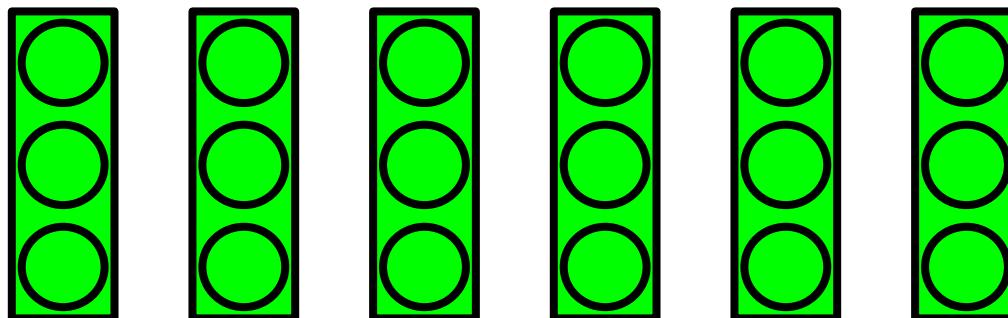
Feature Extraction:

- Freeze feature extraction layers
- Replace the classifier layers
- Train the classifier layers



Fine Tuning:

- Freeze the bottom layers.
- Train the top few layers and classifier layers



Initialization

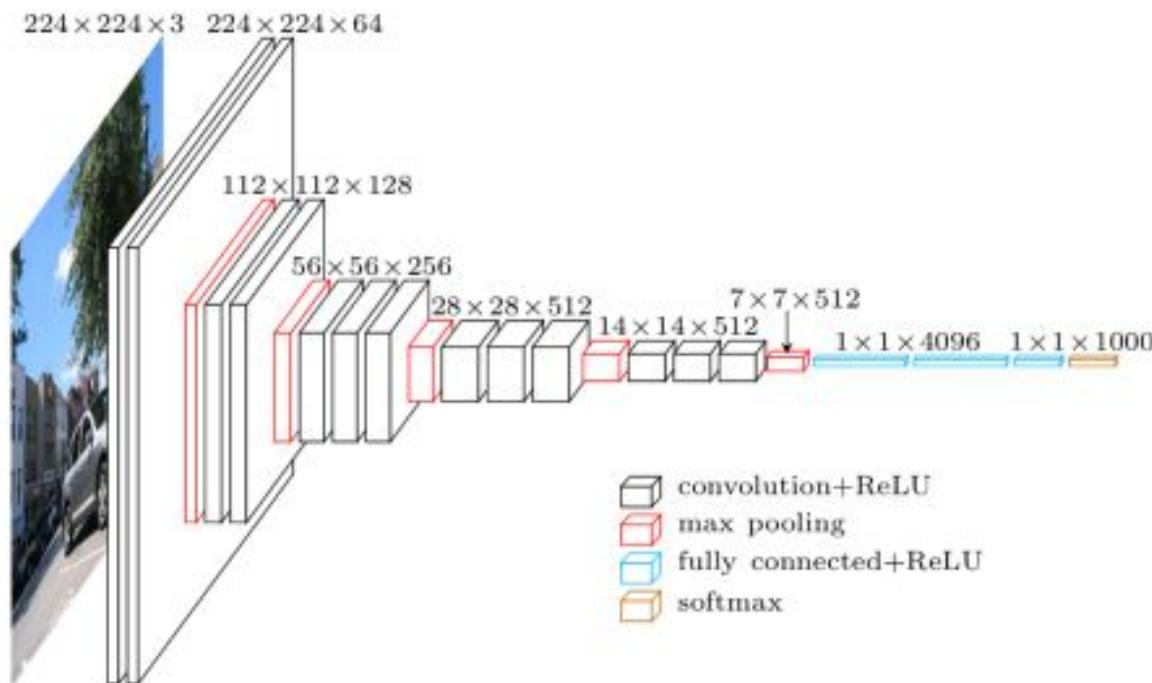
- Retrain the entire layers with preloaded weights as initialization

Pre-Trained TorchVision Models

- AlexNet
- VGG
- ResNet
- SqueezeNet
- DenseNet
- Inception v3
- GoogLeNet
- ShuffleNet v2
- MobileNet v2
- ResNeXt
- Wide ResNet
- MNASNet

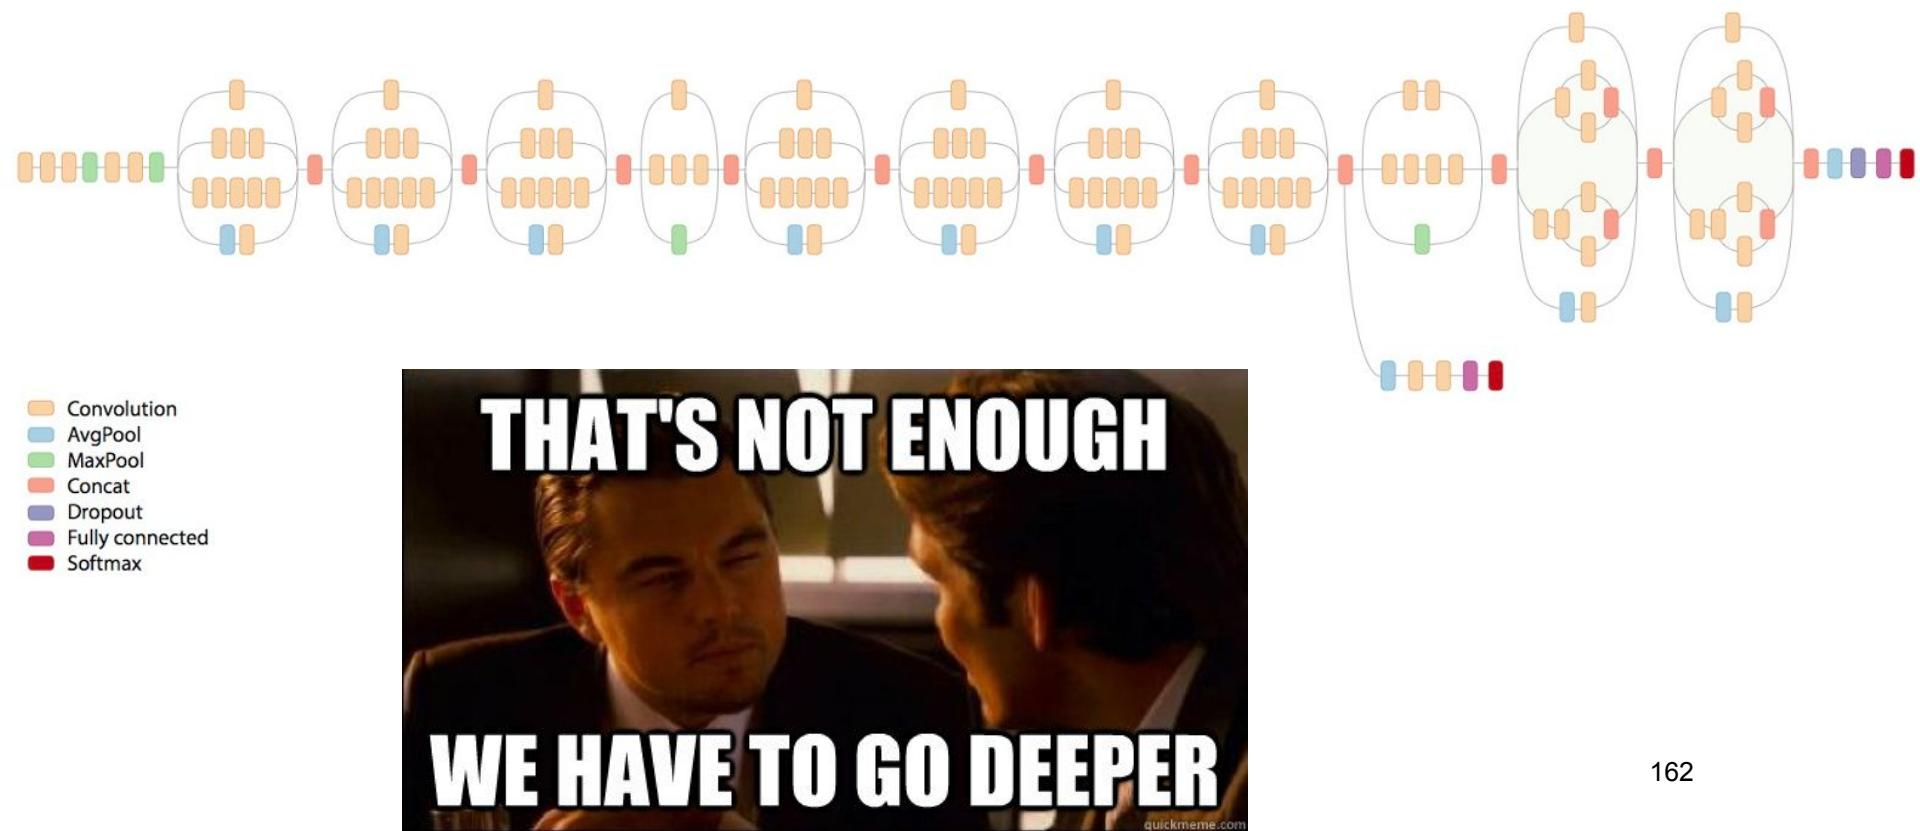
VGG16 Model

- Runner of ILSVRC 2014
- Pre-trained on ImageNet. Input image size is 224x224px
- Simple architecture - 3x3 filters, 2x2 max pooling, double filter numbers after each max pooling



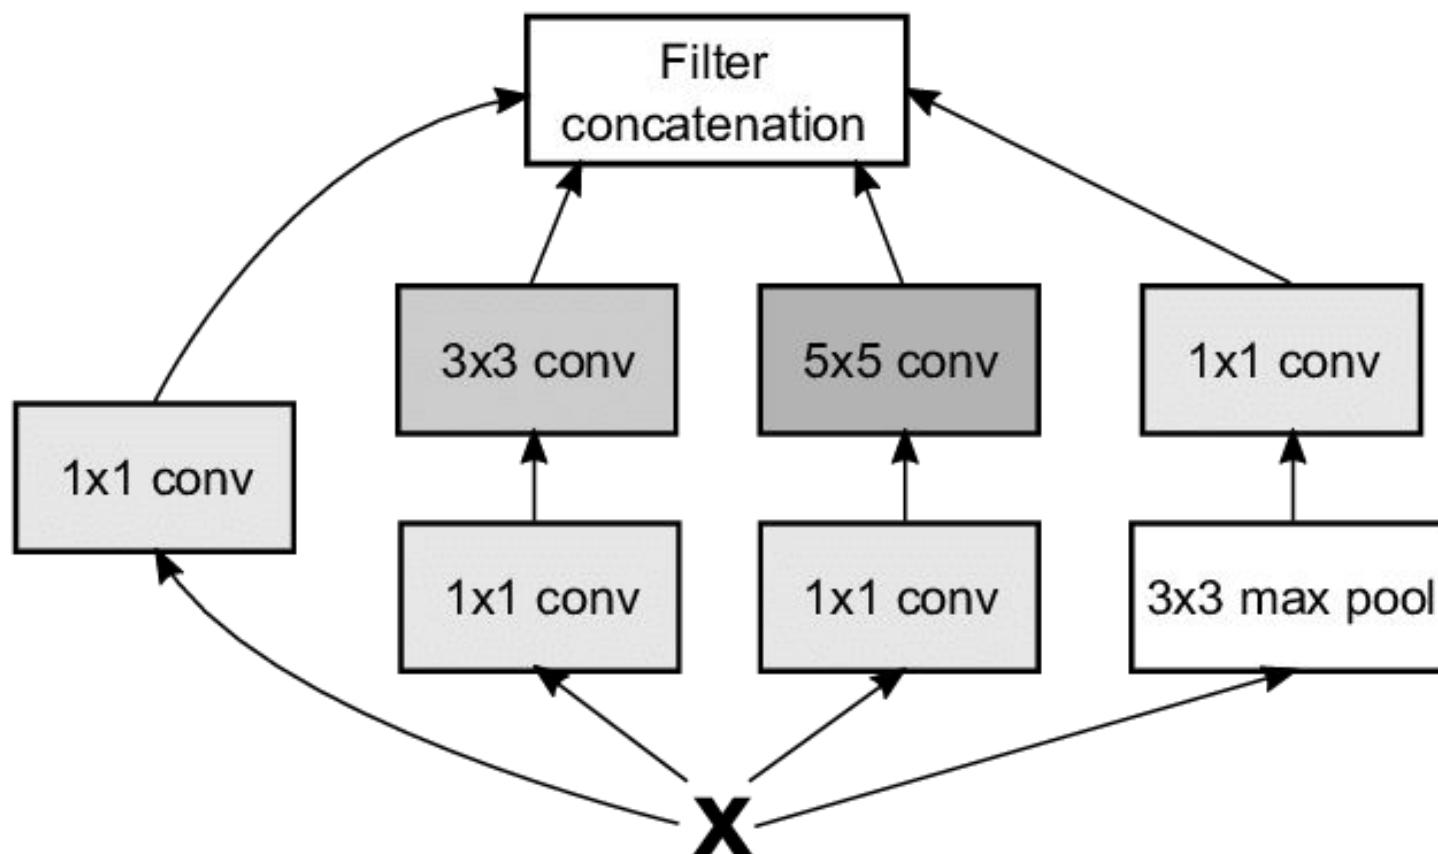
Inception Model

- Winner of ILSVRC 2014, also known as GoogleNet
- Pre-trained on ImageNet. Input Image size is 299x299px
- Consist of many inception modules



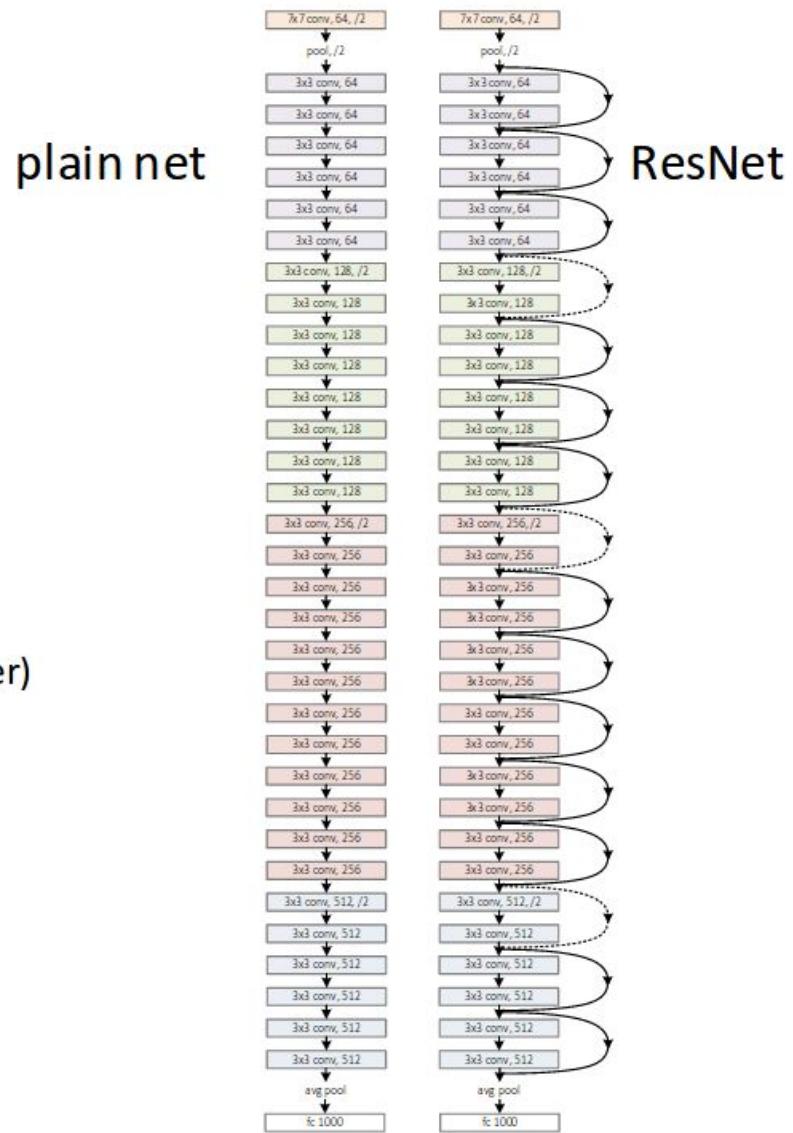
Single Inception Module

The basic idea of inception is to apply different filter sizes eg 1x1, 3x3 and 5x5



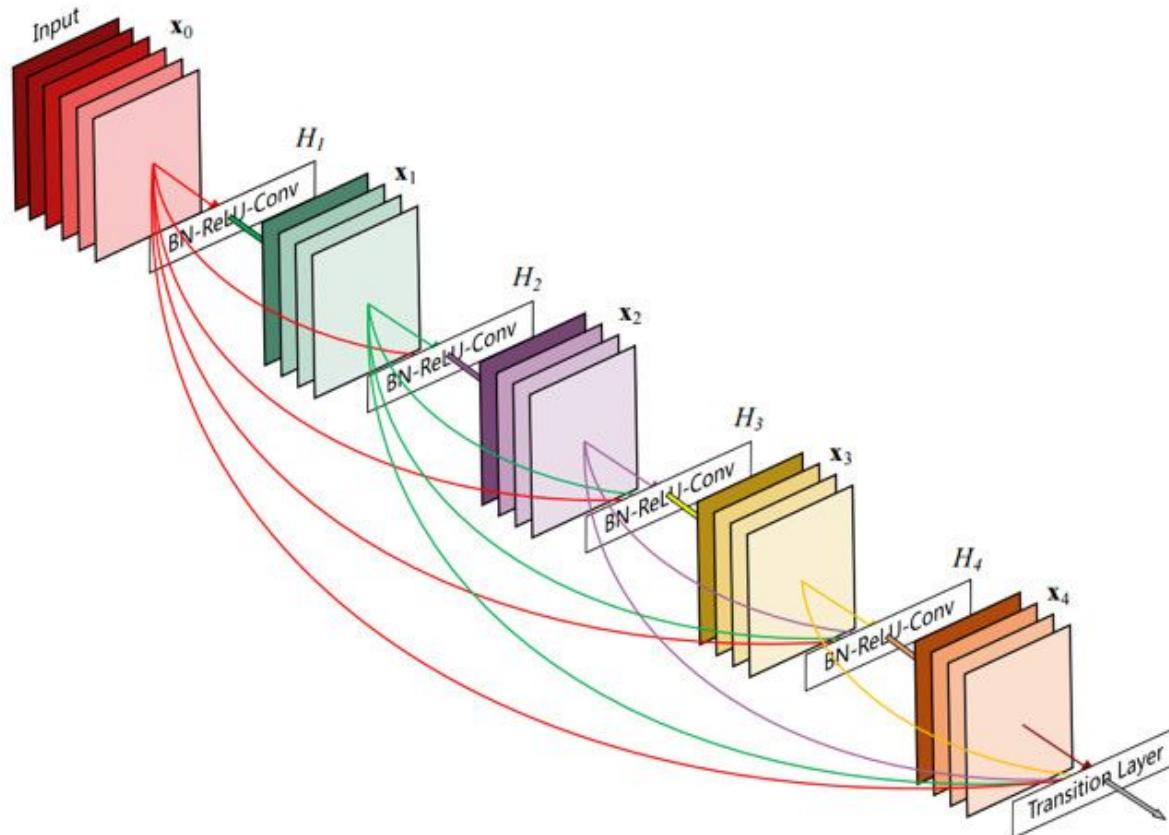
ResNet

- Residual network - winner of ILSVRC 2015 from Microsoft
- Pre-trained on ImageNet. Input image size is 224x224px
- With more layers, the NN performance might degrade due to vanishing gradient issue.
- ResNet allow NN to go deeper by skipping the bottleneck layers



DenseNet

Dense Convolutional Network (DenseNet) is an extension of ResNET, which connects each layer to every other layer in a feed-forward fashion.



Import Pretrained Model

```
import torchvision.models as models
```

```
model = models.vgg16(pretrained=True)
```

Model Features

model.features

```
Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): ReLU(inplace=True)  
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (3): ReLU(inplace=True)  
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (6): ReLU(inplace=True)  
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (8): ReLU(inplace=True)  
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (11): ReLU(inplace=True)  
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (13): ReLU(inplace=True)  
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (15): ReLU(inplace=True)  
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (18): ReLU(inplace=True)  
    (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (20): ReLU(inplace=True)  
    (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (22): ReLU(inplace=True)  
    (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (25): ReLU(inplace=True)  
    (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (27): ReLU(inplace=True)  
    (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (29): ReLU(inplace=True)  
    (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
)
```

Import and Preprocess Data

```
from PIL import Image
```

```
image_path = '/content/cat.jpg'
```

```
preprocess = transforms.Compose([
    transforms.Resize(224),
    transforms.ToTensor(),
    transforms.Normalize((0.485, 0.456, 0.406), (0.229,
0.224, 0.225))
])
```

```
img = Image.open(image_path)
img_tensor = preprocess(img)
img_tensor.unsqueeze_(0)
predict = model(img_tensor)
```

Download ImageNet Labels

```
import requests
```

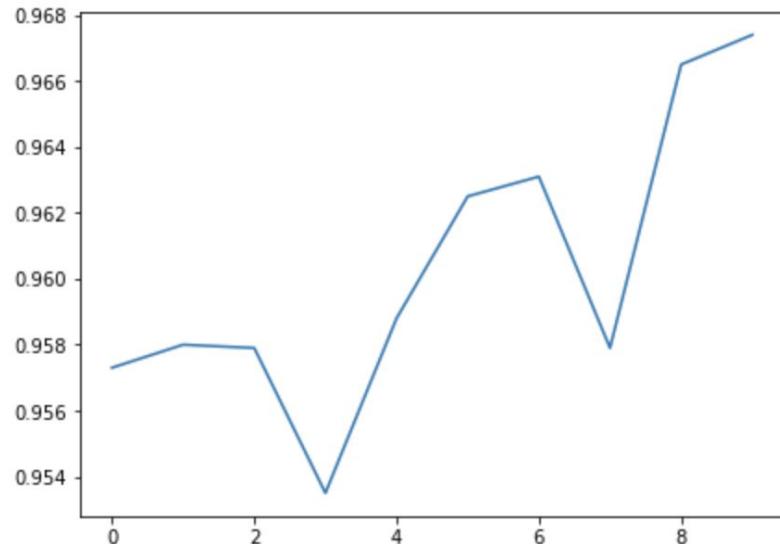
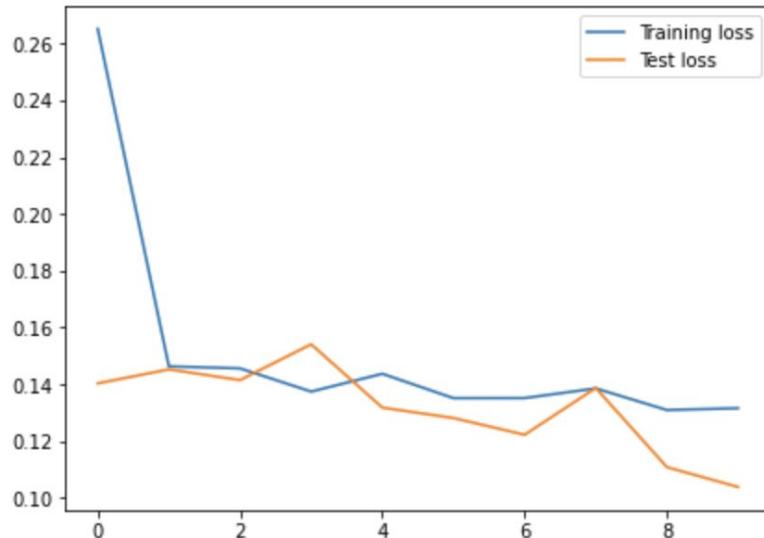
```
LABELS_URL =  
'https://s3.amazonaws.com/deep-learning-models/image-models/ima  
genet_class_index.json'  
labels = {int(key):value for (key, value) in  
requests.get(LABELS_URL).json().items()}
```

Evaluate the Model

```
import matplotlib.pyplot as plt
```

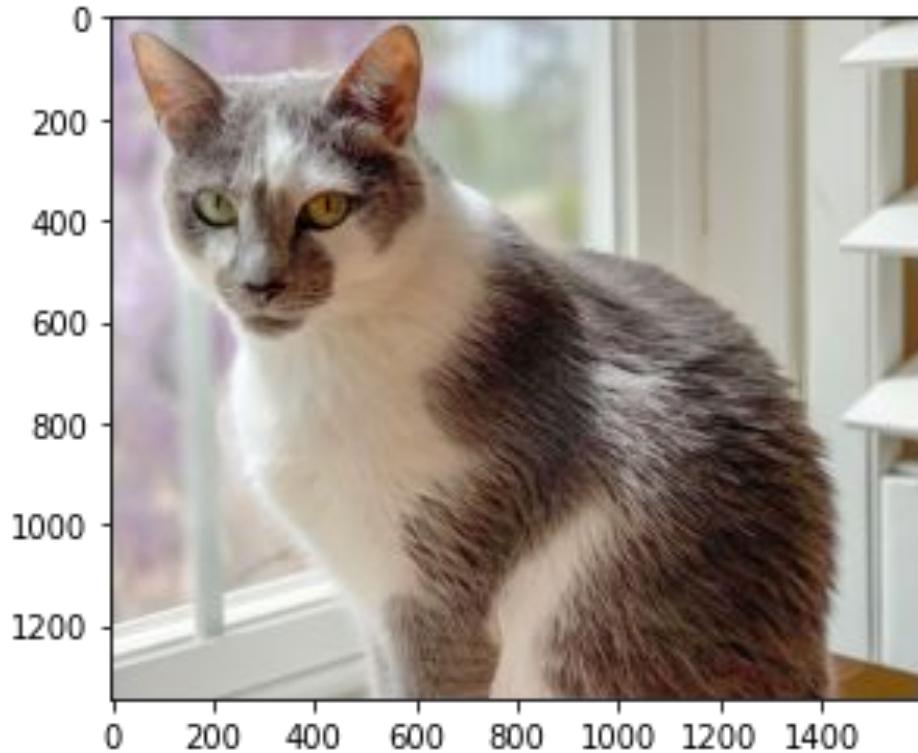
```
plt.figure(figsize=(15,5))
plt.subplot(1,2,1)
plt.plot(train_tracker, label='Training loss')
plt.plot(test_tracker, label='Test loss')
plt.legend()
```

```
plt.subplot(1,2,2)
plt.plot(accuracy_tracker, label='Test accuracy')
plt.show()
```



Prediction Using Pre-trained Model

```
import matplotlib.pyplot as plt  
  
predict = model(img_tensor)  
  
print(labels[predict.data.numpy().argmax()])  
plt.imshow(img)
```



[n02124075, 'Egyptian_cat']

Activity: Pre-trained Model

- Load the RESNET pre-trained model
- Test out an image downloaded from internet.



Feature Extraction- Import Pretrained Model

```
import torchvision.models as models
```

```
model = models.vgg16(pretrained=True)
```

Replace Classifier

```
for param in model.parameters():
    param.requires_grad = False

model.classifier[-1] = nn.Sequential(
    nn.Linear(in_features=4096,
out_features=10)

Sequential(
    (0): Linear(in_features=25088, out_features=4096, bias=True)
    (1): ReLU(inplace=True)
    (2): Dropout(p=0.5, inplace=False)
    (3): Linear(in_features=4096, out_features=4096, bias=True)
    (4): ReLU(inplace=True)
    (5): Dropout(p=0.5, inplace=False)
    (6): Sequential(
        (0): Linear(in_features=4096, out_features=10, bias=True)
    )
)
```

Loss Function & Optimizer

```
from torch.optim import Adam
```

```
criterion = nn.CrossEntropyLoss()
```

```
optimizer = torch.optim.Adam(model.parameters(),  
lr=0.001)
```

Retrain the Model

```
from tqdm.notebook import trange,tqdm
num_epochs = 2

for i in trange(num_epochs):
    train_loss = 0

    for (X, y) in tqdm(trainloader):
        X = X.to(device)
        y = y.to(device)

        yhat = model(X)

        optimizer.zero_grad()
        loss = criterion(yhat, y)
        loss.backward()
        optimizer.step()

        train_loss += loss.item()

    print(f"Epoch({i+1}/{num_epochs}) | Training loss: {train_loss/len(trainloader)} | ",end=)
```

Evaluate the Model

```
from tqdm.notebook import trange,tqdm
```

```
model.to('cpu')
model.eval()
with torch.no_grad():
    num_correct = 0
    total = 0

    for (X,y) in tqdm(testloader):
        yhat = model(X)
        _, pred = torch.max(yhat.data, 1)
        total += y.size(0)
        num_correct += (pred == y.data).sum()
    print(f'Accuracy of the model on {total} test images: {num_correct * 100 / total}% ')
```

100%

39/39 [01:08<00:00, 1.75s/it]

Accuracy of the model on 153 test images: 92.1568603515625%

Model.eval()

- `model.eval()` is a kind of switch for some specific layers/parts of the model that behave differently during training and inference (evaluating) time. For example, Dropouts Layers, BatchNorm Layers etc.
- You need to turn off them during model evaluation, and `.eval()` will do it for you.
- In addition, the common practice for evaluating/validation is using `torch.no_grad()` in pair with `model.eval()` to turn off gradients computation.

```
# evaluate model:  
model.eval()  
with torch.no_grad():  
    ...  
    out_data = model(data)
```

Test the Model

```
import matplotlib.pyplot as plt  
images,labels = next(iter(testloader))  
output = model(images)  
_,pred = torch.max(output,1)  
  
fig = plt.figure(figsize=(10,10))  
for i in range(4):  
    plt.subplot(5,5,i+1)  
    show_img(images[i])  
    plt.title(get_class(labels[i])+"/"+get_class(pred[i]))  
    plt.xticks([])  
    plt.yticks([])
```

ant/ant



bee/bee



ant/ant



bee/bee



Fine Tuning - Load Pre-trained Model

```
import torchvision.models as models  
  
model = models.vgg16(pretrained=True)
```

UnFreeze Layers and Replace Classifier

```
for param in model.parameters():
    param.requires_grad = False
```

```
for i in range(24,31):
    model.features[i].requires_grad = True
```

```
model.classifier[-1] = nn.Sequential(
    nn.Linear(in_features=4096,
out_features=10)
)
```

Loss Function & Optimizer

```
from torch.optim import Adam
```

```
loss_func = nn.CrossEntropyLoss()
```

```
model = model.to(device)
```

```
optimizer = Adam(filter(lambda p: p.requires_grad, model.parameters()))
```

Re-train the Model

```
from tqdm.notebook import trange,tqdm

num_epochs = 2

for i in trange(num_epochs):
    train_loss = 0

    for (X, y) in tqdm(trainloader):
        X = X.to(device)
        y = y.to(device)

        yhat = model(X)

        optimizer.zero_grad()
        loss = criterion(yhat, y)
        loss.backward()
        optimizer.step()

        train_loss += loss.item()

    print(f"Epoch({i+1}/{num_epochs}) | Training loss: {train_loss/len(trainloader)} | ",end="")
```

Evaluate the Model

```
from tqdm.notebook import trange,tqdm
```

```
model.to('cpu')
model.eval()
with torch.no_grad():
    num_correct = 0
    total = 0
```

```
#set_trace()
for (X,y) in tqdm(testloader):

    yhat = model(X)
    _, pred = torch.max(yhat.data, 1)
    total += y.size(0)
    num_correct += (pred == y.data).sum()
```

```
print(f'Accuracy of the model on {total} test images: {num_correct * 100 / total}%
```

100%

39/39 [01:08<00:00, 1.75s/it]

Accuracy of the model on 153 test images: 93.46405029296875%

Test the Model

```
import matplotlib.pyplot as plt
images,labels = next(iter(testloader))
output = model(images)
_,pred = torch.max(output,1)
fig = plt.figure(figsize=(10,10))
for i in range(4):
    plt.subplot(5,5,i+1)
    show_img(images[i])
    plt.title(get_class(labels[i])+"/"+get_class(pred[i]))
    plt.xticks([])
    plt.yticks([])
```



Topic 6

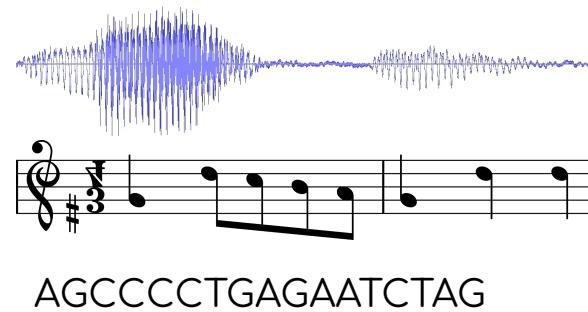
Recurrent

Neural Networks (RNN)

Sequential Data

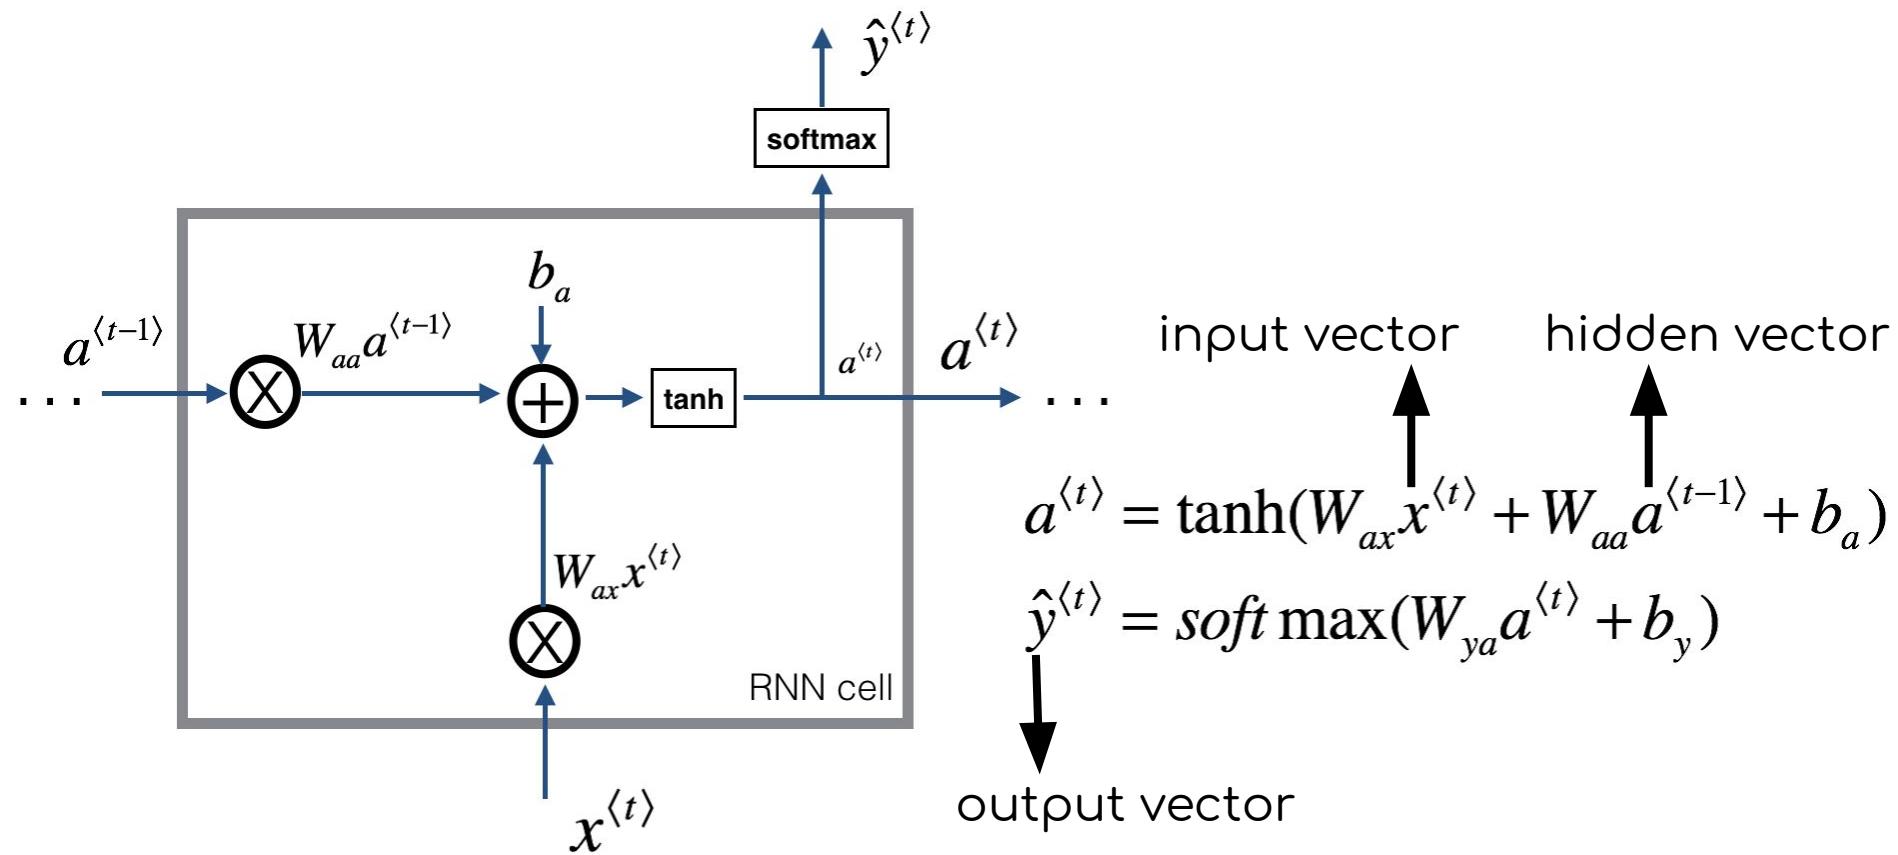
Conventional NN and CNN are not able to model sequential data such as

- Speech
- Music
- DNA sequence
- Video
- Machine translation
- Name entity detection



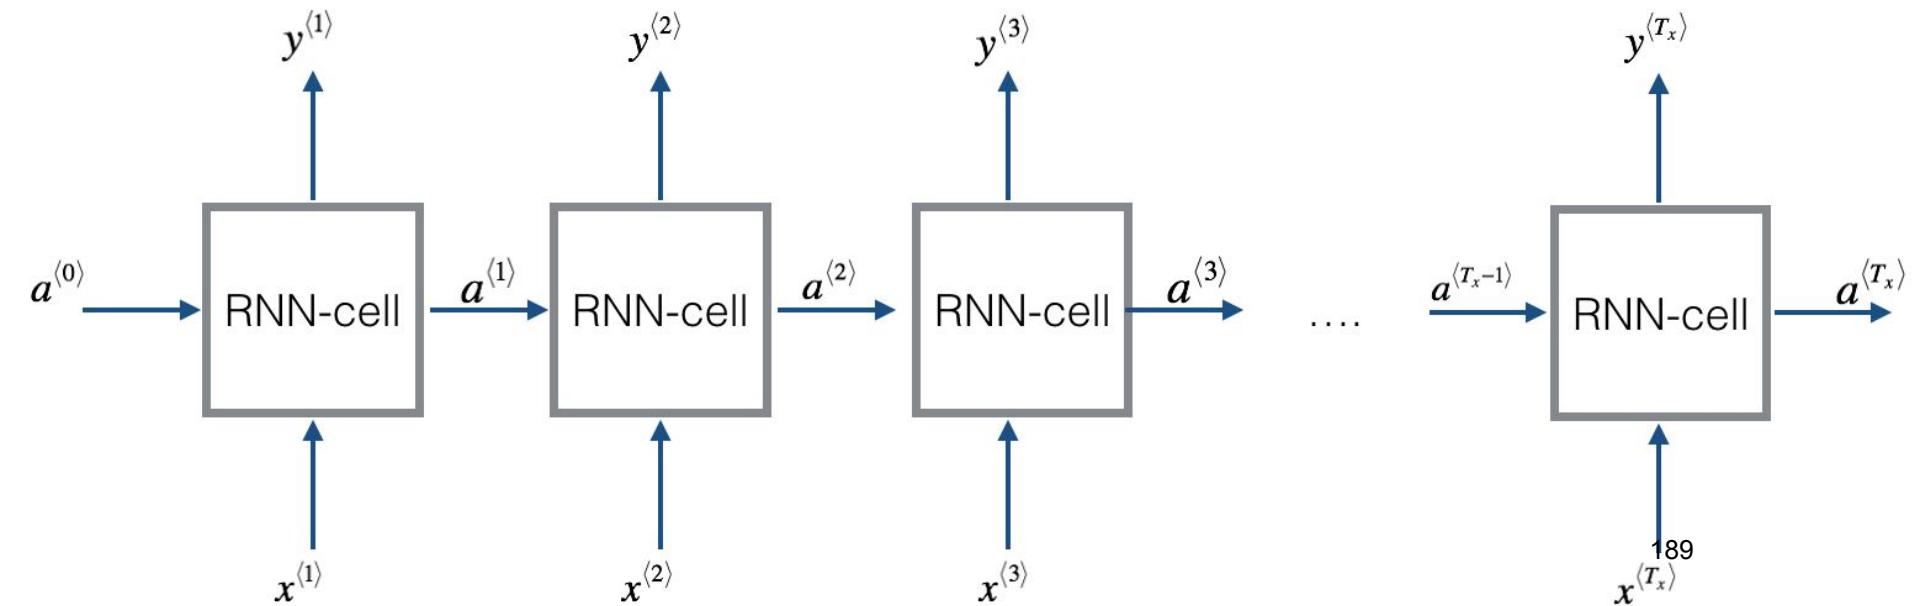
What is RNN

A RNN cell can be seen as a simple one hidden layer NN (a) with input layer (a+x) and output layer (y)



Rollout RNN vs Time

- To input the sequential data to the RNN, RNN will roll out with time.
- RNN will read inputs $x^{(t)}$ one at a time, and remember some info through the hidden layer activations that get passed from one time-step to the next.
- This allows a uni-directional RNN to take info from the past to process later inputs



Long Term Dependencies Issue

"*Michel C. was born in Paris, France.....His mother tongue is ??????*"

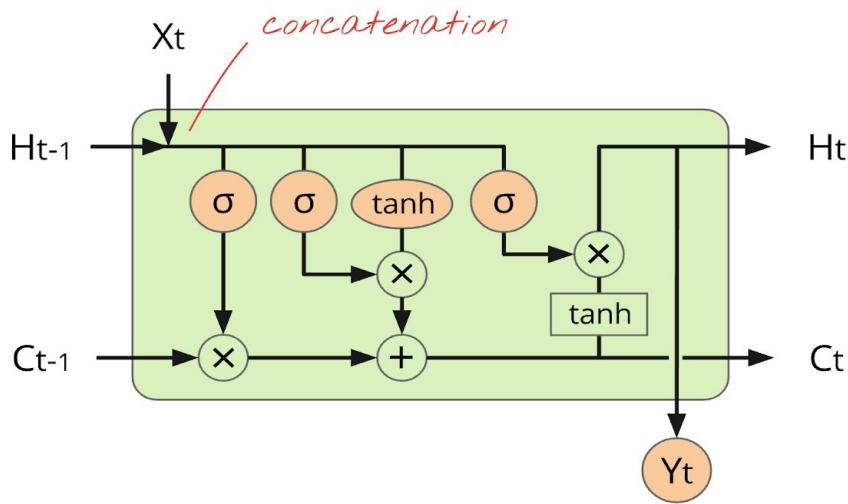
"*The cat, which already ate.....was*"

"*The cats, which already atewere*"

It is difficult to predict "French" or was/were if the sentences were very long due to vanishing gradients

LSTM Cell

- There are two internal states , C_t and H_t .
- C_t is the cell state (memory) to retain via forget and update gates
- H_t is the hidden state fed to the softmax via the result gate.

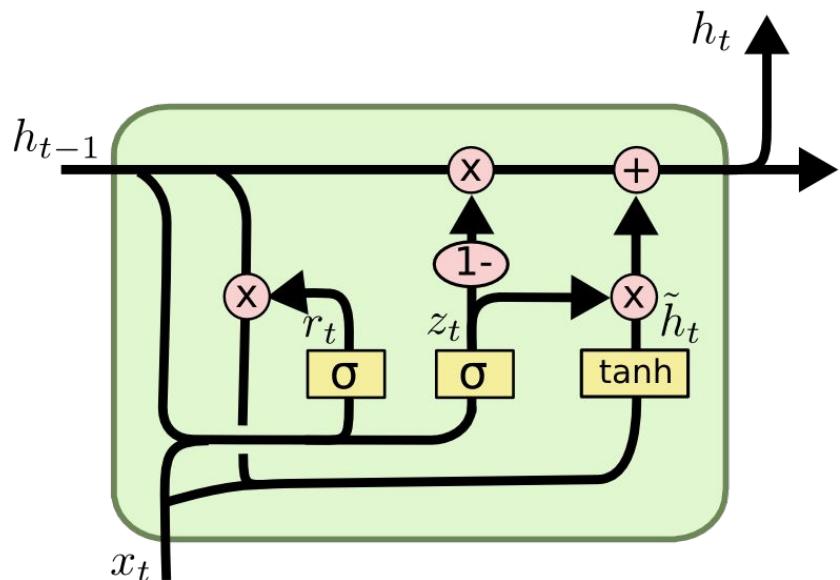


$$\begin{aligned} f &= \sigma(X \cdot W_f + b_f) \\ u &= \sigma(X \cdot W_u + b_u) \\ r &= \sigma(X \cdot W_r + b_r) \\ X' &= \tanh(X \cdot W_c + b_c) \\ C_t &= f * C_{t-1} + u * X' \\ H_t &= r * \tanh(C_t) \end{aligned}$$

GRU Cell

GRU: Gated Recurrent Unit

Only two gates instead of 3 gates



$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

Types of RNN Architecture

RNN can have a variety of architecture

Image Classification

Sentiment Analysis

Video Classification

one to one

one to many

many to one

many to many

many to many

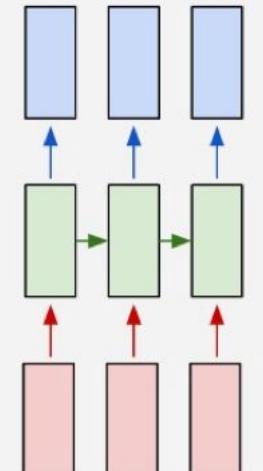
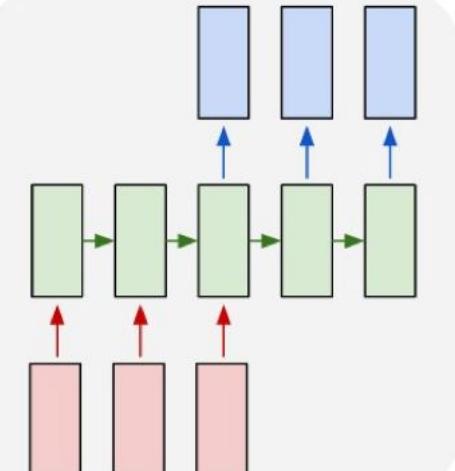
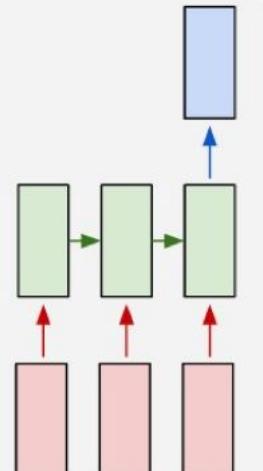
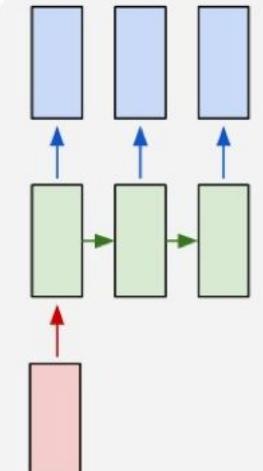
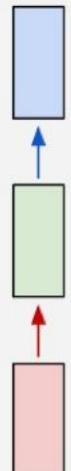
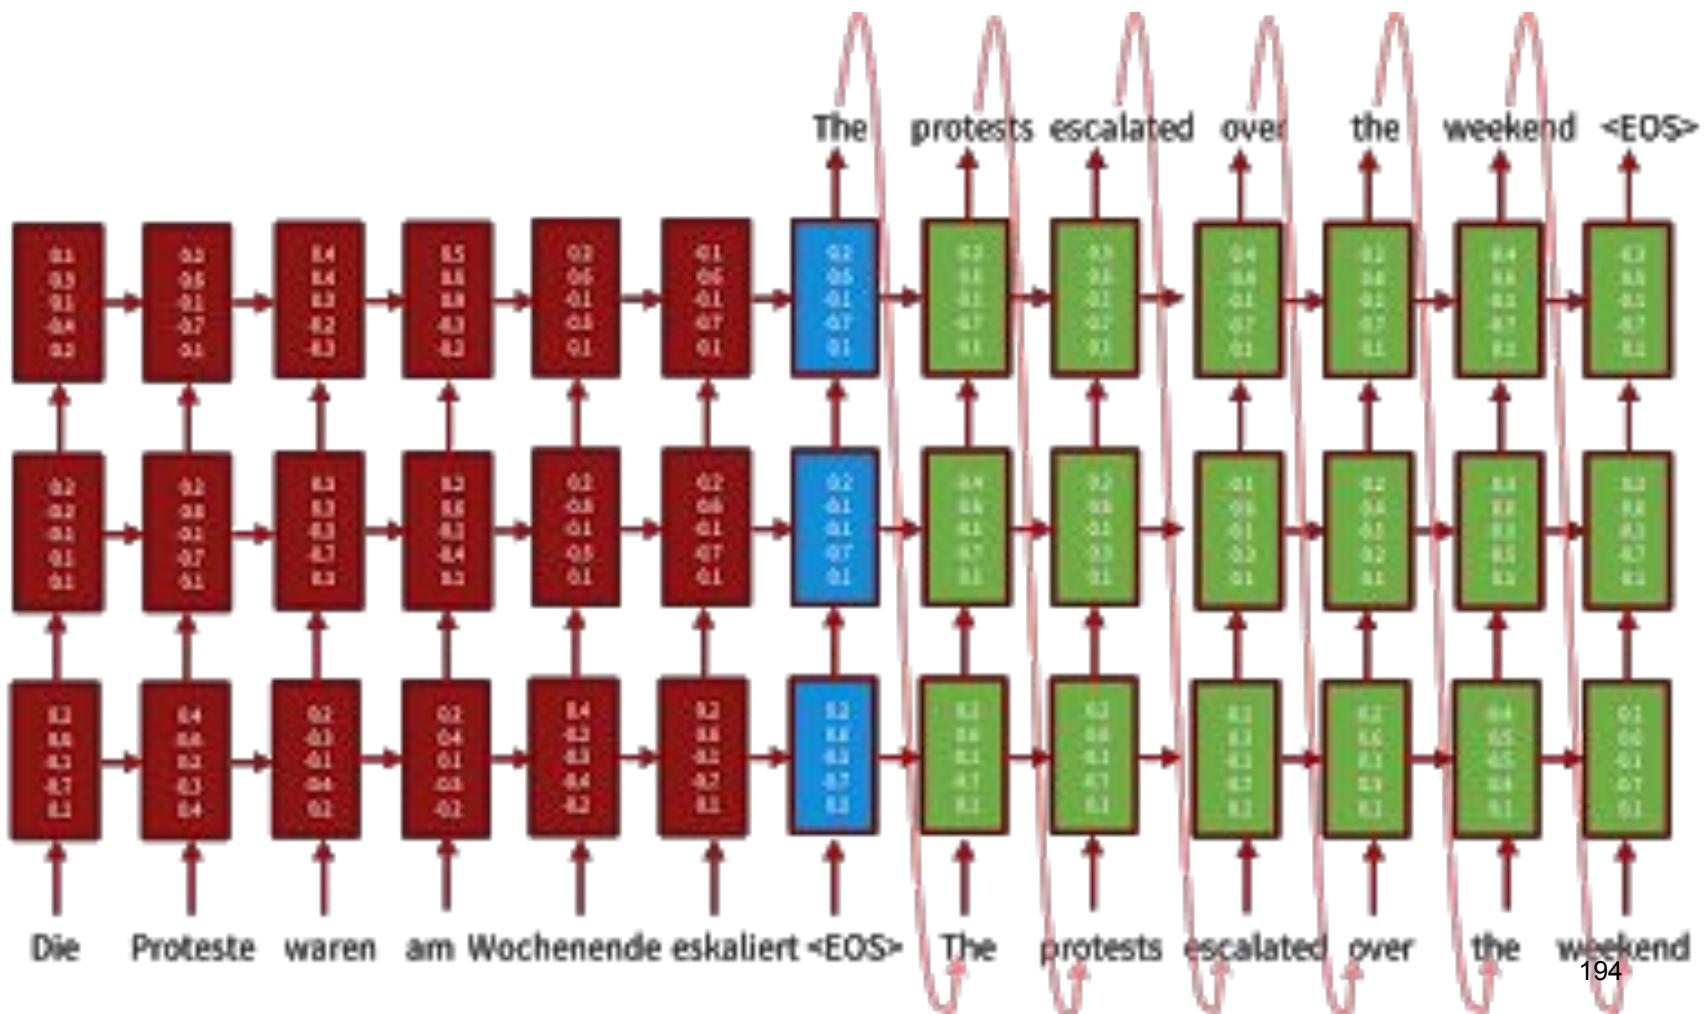


Image Captioning
Music Generation
Text Generation

Sequence to Sequence
Language Translation
Chatbot

Sequence to Sequence RNN

Sequence to Sequence RNN have been used successfully for chatbot or machine translation..



Brief History of RNN

- 1982 - Hopfield Network
- 1997 - Hochreiter invented LSTM
- 2009 - Alex Graves demo cursive writing recognition with LSTM
- 2014 - Baidu demo speech recognition with LSTM
- 2016 - Google deployed seq2seq for machine translation

RNN Parameters

`torch.nn.RNN(...)` - output, h_n

Parameters

- **input_size** – The number of expected features in the input x
- **hidden_size** – The number of features in the hidden state h
- **num_layers** – Number of recurrent layers. E.g., setting `num_layers=2` would mean stacking two RNNs together to form a *stacked RNN*, with the second RNN taking in outputs of the first RNN and computing the final results. Default: 1
- **nonlinearity** – The non-linearity to use. Can be either `'tanh'` or `'relu'`. Default: `'tanh'`
- **bias** – If `False`, then the layer does not use bias weights b_{ih} and b_{hh} . Default: `True`
- **batch_first** – If `True`, then the input and output tensors are provided as $(batch, seq, feature)$ instead of $(seq, batch, feature)$. Note that this does not apply to hidden or cell states. See the Inputs/Outputs sections below for details. Default: `False`
- **dropout** – If non-zero, introduces a *Dropout* layer on the outputs of each RNN layer except the last layer, with dropout probability equal to `dropout`. Default: 0
- **bidirectional** – If `True`, becomes a bidirectional RNN. Default: `False`

LSTM Parameters

`torch.nn.LSTM(...)` - output, h_n,c_n

Parameters

- **input_size** – The number of expected features in the input x
- **hidden_size** – The number of features in the hidden state h
- **num_layers** – Number of recurrent layers. E.g., setting `num_layers=2` would mean stacking two RNNs together to form a *stacked RNN*, with the second RNN taking in outputs of the first RNN and computing the final results. Default: 1
- **nonlinearity** – The non-linearity to use. Can be either `'tanh'` or `'relu'`. Default: `'tanh'`
- **bias** – If `False`, then the layer does not use bias weights b_{ih} and b_{hh} . Default: `True`
- **batch_first** – If `True`, then the input and output tensors are provided as $(batch, seq, feature)$ instead of $(seq, batch, feature)$. Note that this does not apply to hidden or cell states. See the Inputs/Outputs sections below for details. Default: `False`
- **dropout** – If non-zero, introduces a *Dropout* layer on the outputs of each RNN layer except the last layer, with dropout probability equal to `dropout`. Default: 0
- **bidirectional** – If `True`, becomes a bidirectional RNN. Default: `False`

LSTM Input and Output Sizes

source: <https://pytorch.org/docs/stable/generated/torch.nn.LSTM.html>

- **input:** tensor of shape (L, N, H_{in}) when `batch_first=False` or (N, L, H_{in}) when `batch_first=True` containing the features of the input sequence. The input can also be a packed variable length sequence. See `torch.nn.utils.rnn.pack_padded_sequence()` or `torch.nn.utils.rnn.pack_sequence()` for details.
- **h_0:** tensor of shape $(D * \text{num_layers}, N, H_{out})$ containing the initial hidden state for each element in the batch. Defaults to zeros if (h_0, c_0) is not provided.
- **c_0:** tensor of shape $(D * \text{num_layers}, N, H_{cell})$ containing the initial cell state for each element in the batch. Defaults to zeros if (h_0, c_0) is not provided.

where:

N = batch size

L = sequence length

D = 2 if `bidirectional=True` otherwise 1

H_{in} = input_size

H_{cell} = hidden_size

H_{out} = `proj_size` if `proj_size > 0` otherwise `hidden_size`

GRU Parameters

`torch.nn.GRU(...)` - output, h_n

Parameters

- **input_size** – The number of expected features in the input x
- **hidden_size** – The number of features in the hidden state h
- **num_layers** – Number of recurrent layers. E.g., setting `num_layers=2` would mean stacking two GRUs together to form a *stacked GRU*, with the second GRU taking in outputs of the first GRU and computing the final results. Default: 1
- **bias** – If `False`, then the layer does not use bias weights b_{ih} and b_{hh} . Default: `True`
- **batch_first** – If `True`, then the input and output tensors are provided as $(batch, seq, feature)$ instead of $(seq, batch, feature)$. Note that this does not apply to hidden or cell states. See the Inputs/Outputs sections below for details. Default: `False`
- **dropout** – If non-zero, introduces a *Dropout* layer on the outputs of each GRU layer except the last layer, with dropout probability equal to `dropout`. Default: 0
- **bidirectional** – If `True`, becomes a bidirectional GRU. Default: `False`

Parameters Description Examples

Consider the case of text classification:

- Each batch consist of a sequence of sentences
- Each sentence consist of a sequence of word tokens
- Each token is represented by an embedding vector

For example:

batch_size = 5 sentences or 100 days of stock price

sequence_length = 80 word tokens or 3 days of stock price

input_size = word embedding dimension or 1 for stock price

Batch-First (True)

```
import torch
import torch.nn as nn

number_layers = 1
batch_size = 1
input_size = 10
hidden_size = 20
sequence_length = 3

rnn = nn.LSTM(input_size=input_size, hidden_size=hidden_size,
num_layers=number_layers, batch_first=True)
input = torch.randn(batch_size, sequence_length, input_size)
h0 = torch.randn(number_layers, batch_size, hidden_size)
c0 = torch.randn(number_layers, batch_size, hidden_size)
_,(hn, cn) = rnn(input, (h0, c0))
```

Batch First (False)

```
import torch
import torch.nn as nn

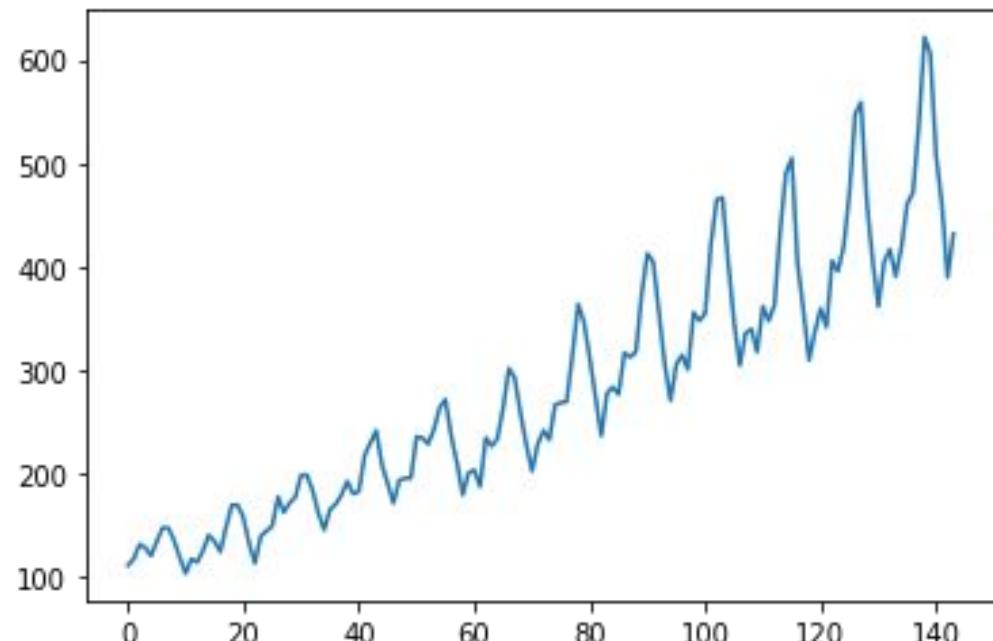
number_layers = 1
batch_size = 1
input_size = 10
hidden_size = 20
sequence_length = 3

rnn = nn.LSTM(input_size=input_size, hidden_size=hidden_size,
num_layers=number_layers, batch_first=False)
input = torch.randn(sequence_length, batch_size, input_size)
h0 = torch.randn(number_layers, batch_size, hidden_size)
c0 = torch.randn(number_layers, batch_size, hidden_size)
_,(hn, cn) = rnn(input, (h0, c0))
```

Load the Data

```
training_set = pd.read_csv('airline-passengers.csv')  
training_set.head(10)
```

	Month	Passengers
0	1949-01	112
1	1949-02	118
2	1949-03	132
3	1949-04	129
4	1949-05	121
5	1949-06	135
6	1949-07	148
7	1949-08	148
8	1949-09	136
9	1949-10	119



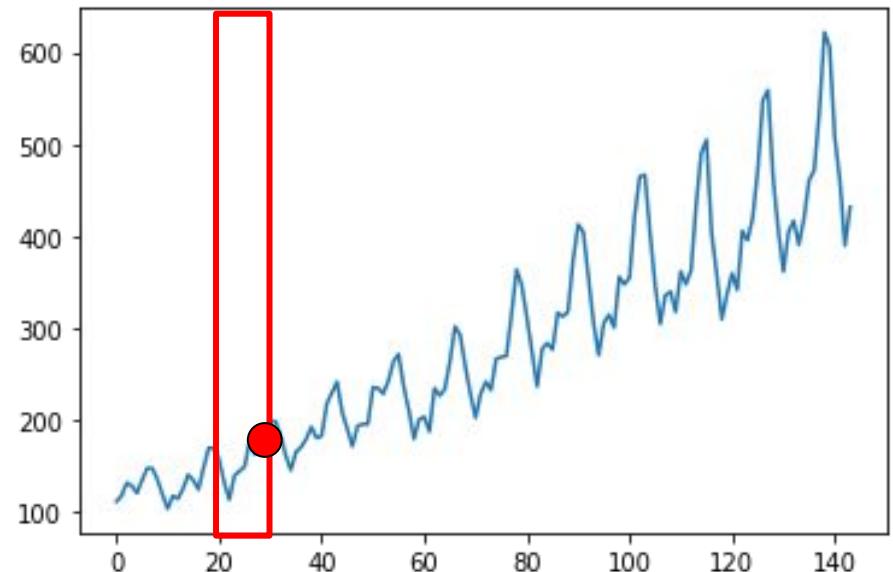
Sliding Window

```
def sliding_window(data, seq_length):
    x = []
    y = []

    for i in range(len(data)-seq_length-1):
        _x = data[i:(i+seq_length)]
        _y = data[i+seq_length]
        x.append(_x)
        y.append(_y)

    return np.array(x),np.array(y)
```

```
seq_length = 4
x, y = scaling_window(training_data,
seq_length)
```



Scale the Data

```
sc = MinMaxScaler()  
training_data = sc.fit_transform(training_set)
```

Split the Data

```
train_size = int(len(y) * 0.67)
```

```
test_size = int(len(y)) - train_size
```

```
dataX = torch.Tensor(np.array(x))
```

```
dataY = torch.Tensor(np.array(y))
```

```
trainX = torch.Tensor(np.array(x[0:train_size]))
```

```
trainY = torch.Tensor(np.array(y[0:train_size]))
```

```
testX = torch.Tensor(np.array(x[train_size:len(x)]))
```

```
testY = torch.Tensor(np.array(y[train_size:len(y)]))
```

Define the LSTM Model

```
class LSTM(nn.Module):  
  
    def __init__(self, num_classes, input_size, hidden_size, layer1, layer2):  
        super(LSTM, self).__init__()  
  
        self.num_classes = num_classes  
        self.num_layers = num_layers  
        self.input_size = input_size  
        self.hidden_size = hidden_size  
  
        self.lstm = nn.LSTM(input_size=input_size, hidden_size=hidden_size,  
                           num_layers=num_layers, batch_first=True)  
  
        self.fc = nn.Linear(hidden_size, num_classes)  
  
    def forward(self, x):  
        h_0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size)  
        c_0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size)  
  
        ula, (h_out, _) = self.lstm(x, (h_0, c_0))  
        h_out = h_out.view(-1, self.hidden_size)  
        out = self.fc(h_out)  
  
        return out
```

the batch size is
 $x.size(0)$

Loss Function and Optimizer

```
learning_rate = 0.01
```

```
criterion = torch.nn.MSELoss()  
optimizer = torch.optim.Adam(lstm.parameters(),  
lr=learning_rate)
```

Train the LSTM Model

```
num_epochs = 2000
```

```
for epoch in range(num_epochs):
```

```
    outputs = lstm(trainX)
```

```
    optimizer.zero_grad()
```

```
    loss = criterion(outputs, trainY)
```

```
    loss.backward()
```

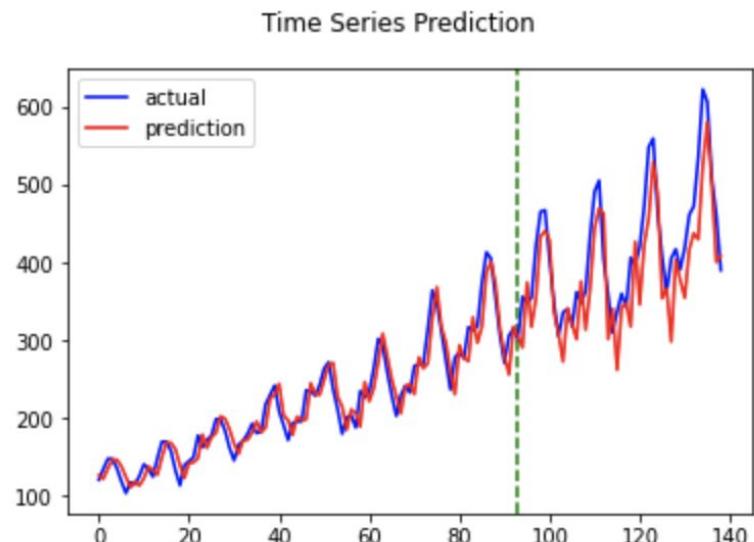
```
    optimizer.step()
```

```
if epoch % 100 == 0:
```

```
    print("Epoch: %d, loss: %1.5f" % (epoch, loss.item()))
```

Evaluate the LSTM Model

```
lstm.eval()  
train_predict = lstm(dataX)  
  
data_predict = train_predict.data.numpy()  
dataY_plot = dataY.data.numpy()  
  
data_predict = sc.inverse_transform(data_predict)  
dataY_plot = sc.inverse_transform(dataY_plot)  
  
plt.axvline(x=train_size, c='g', linestyle='--')  
  
plt.plot(dataY_plot,'b',label='actual')  
plt.plot(data_predict,'r',label='prediction')  
plt.suptitle('Time Series Prediction')  
plt.legend()  
plt.show()
```



Define the GRU Model

```
class GRU(nn.Module):  
  
    def __init__(self, num_classes, input_size, hidden_size, num_layers):  
        super(GRU, self).__init__()  
  
        self.num_classes = num_classes  
        self.num_layers = num_layers  
        self.input_size = input_size  
        self.hidden_size = hidden_size  
        self.seq_length = seq_length  
  
        self.gru = nn.GRU(input_size=input_size, hidden_size=hidden_size,  
                         num_layers=num_layers, batch_first=True)  
  
        self.fc = nn.Linear(hidden_size, num_classes)  
  
    def forward(self, x):  
        h_0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size)  
  
        # Propagate input through LSTM  
        _, h_out = self.gru(x, h_0)  
        h_out = h_out.view(-1, self.hidden_size)  
        out = self.fc(h_out)  
  
    return out
```

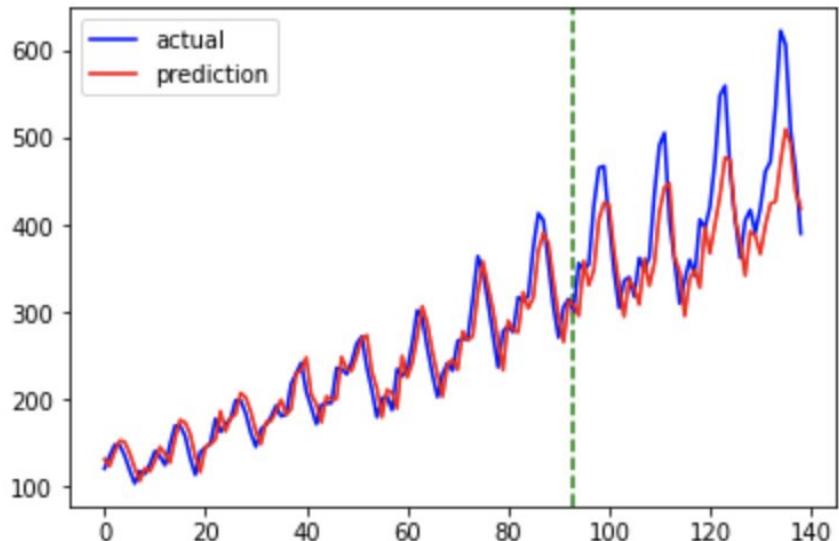
Define the RNN Model

```
class RNN(nn.Module):  
  
    def __init__(self, num_classes, input_size, hidden_size, num_layers):  
        super(RNN, self).__init__()  
  
        self.num_classes = num_classes  
        self.num_layers = num_layers  
        self.input_size = input_size  
        self.hidden_size = hidden_size  
        self.seq_length = seq_length  
  
        self.rnn = nn.RNN(input_size=input_size, hidden_size=hidden_size,  
                          num_layers=num_layers, batch_first=True)  
  
        self.fc = nn.Linear(hidden_size, num_classes)  
  
    def forward(self, x):  
        h_0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size)  
  
        # Propagate input through LSTM  
        _, h_out = self.rnn(x, h_0)  
        h_out = h_out.view(-1, self.hidden_size)  
        out = self.fc(h_out)  
  
        return out
```

Comparison LSTM, GRU & RNN

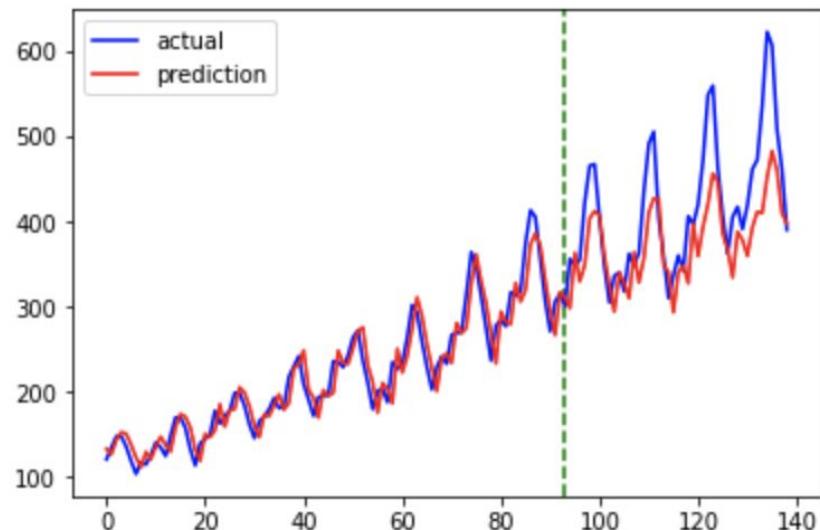
GRU

Time Series Prediction



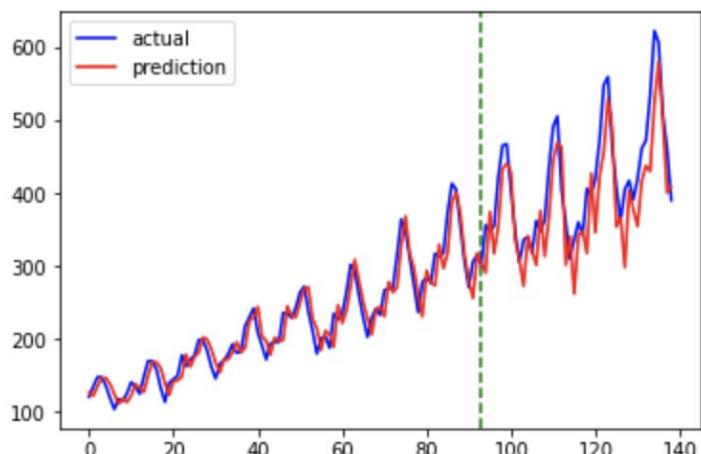
RNN

Time Series Prediction



LSTM

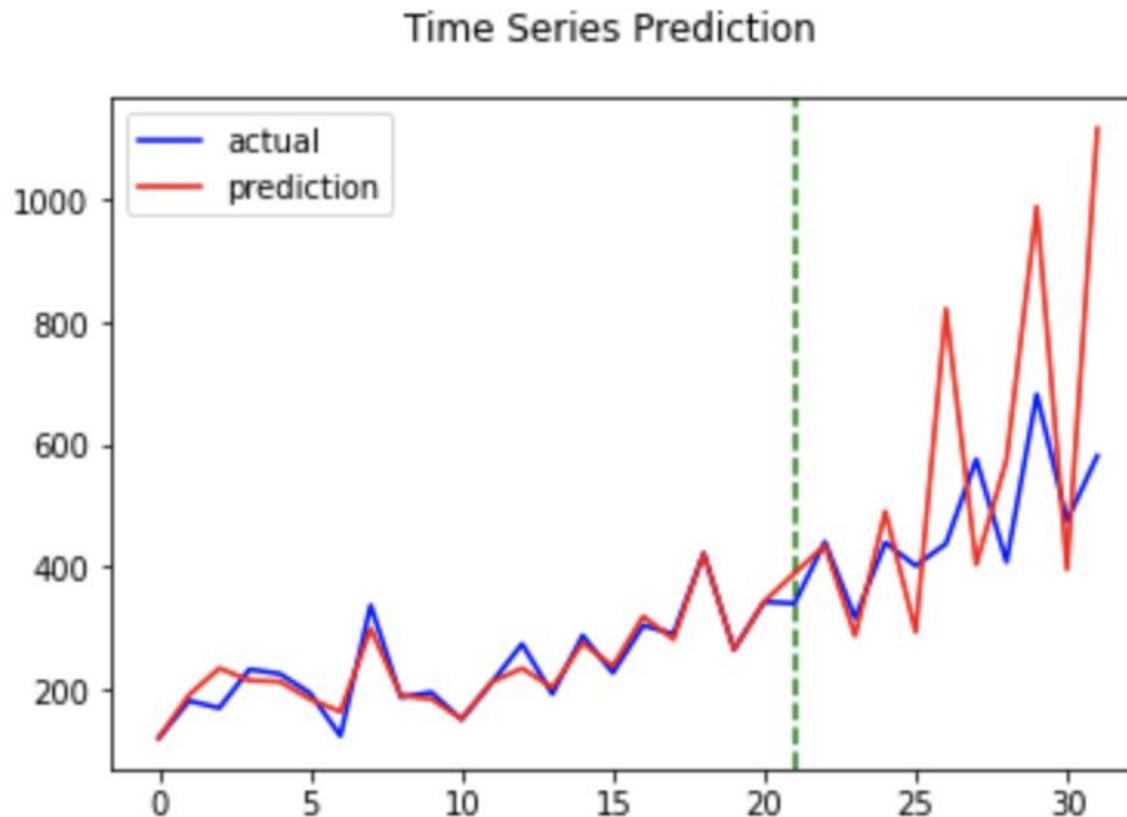
Time Series Prediction



- LSTM shows slightly better result than GRU and RNN.

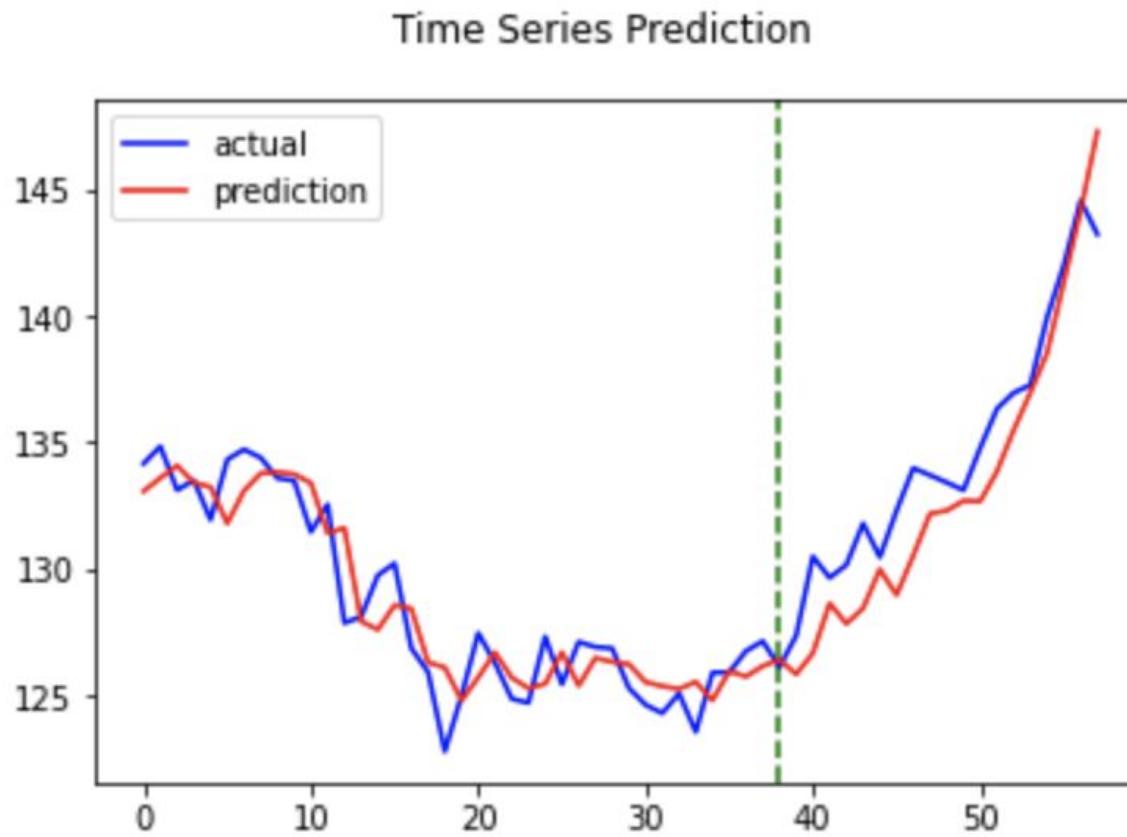
Activity: Sales Prediction

- Load the shampoo.csv dataset
- Create a LSTM model to predict the sales based on historical data.



Activity: Stock Price Prediction

- Goto Yahoo Finance to download at least 3 months of a stock price for a particular stock
- Create a LSTM model to predict the stock price based on historical data.



Summary

Q&A



Course Feedback

<https://goo.gl/R2eumq>





CERTIFICATE *of ACCOMPLISHMENT*

You will receive a digital certificate in your email
after the completion of the class

If you did not receive the digital certificate,
please send your request to
enquiry@tertiaryinfotech.com

Thank You!

Dr. Alfred Ang
Tel: 96983731

Email: angch@tertiaryinfotech.com

Facebook:

<https://www.facebook.com/angchewhoe>

Linkedin:

<https://www.linkedin.com/in/angchewhoe/>